Interpretable Allocation and Model-Predictive Rebalancing of Equity Portfolios, Tested on Real and Synthetic Markets

Jeffrey D. Varner*
Robert Frederick Smith School of Chemical and Biomolecular Engineering
Cornell University, Ithaca, NY 14850
jdv27@cornell.edu

Abstract

A live equity portfolio poses three problems: which stocks to hold, when to trade, and whether the strategy holds up beyond the observed market. We built one system for all three and tested it on 373 US equities, calibrated on 2014–2024 and traded January 2025 to September 2026. A selection rule chose stocks by score, by a bandit or at random, and one of seven weighting rules, four consumer-theory utilities and three standard portfolios, set share counts under limits on budget, single-stock weight, portfolio volatility and turnover, with each stock's preference weight a bounded score of its trend, market beta and the market's recent direction. A model-predictive controller forecast portfolio value with a multi-asset geometric Brownian motion and re-planned at its horizon, when the value left its forecast band, at a drawdown stop, or when a scheduled pick changed the stocks. To probe deeper declines than the real window held, we also ran the system on simulated markets from per-asset hidden Markov generators with jumps, linked by market and sector factors. On real prices, the stocks chosen mattered more to terminal wealth than the weighting rule, and the early re-plans shrank the worst drawdown in most rule combinations against a fixed calendar without consistently raising wealth. In the deepest simulated declines, every combination, random picks included, ended above simulated SPY and the equal-weight universe in the median, even without early re-plans. The equal-weight universe also ended above simulated SPY, so part of the margin over SPY came from the simulated market and the rest from shared parts of the system, not from stock selection or re-plan timing. A market hedge that at each re-plan capped portfolio beta more tightly when the market was weaker cost wealth on real prices but paid off in the most sustained simulated declines, and the generator's persistent trends favored score-based selection more than the real market did, which suggests a synthetic test judges stock selection fairly only if its generator does not build in that persistence.

Keywords: portfolio allocation; utility maximization; model-predictive control; rebalancing; backtesting; synthetic financial data; hidden Markov model; stress testing.

*Corresponding author.

Preprint.

1 Introduction

Managing a live equity portfolio raises three linked questions. An allocation rule answers what to hold, given forecasts and limits on risk, concentration and trading; a rebalancing rule answers when the holdings should be recomputed; and an evaluation asks whether the two rules together would have worked on markets other than the one that was observed. Mean-variance optimization [1] and its continuous-time extensions [2] give a principled answer to the first question, but optimized weights are sensitive to errors in the estimated means and covariances [3, 4], and an equal-weight portfolio often matches optimized portfolios out of sample [5]. Transaction costs turn the second question into a control problem. With proportional costs, the optimal policy leaves the holdings alone inside a no-trade region and trades only at its boundary [6, 7], and multi-period formulations solve a sequence of forecast-based optimizations over a receding horizon [8, 9]. The third question is usually answered with a backtest on the historical record, which is a single path of the market. Choosing among rules on one path inflates their measured performance [10, 11], and a short window may contain few of the market conditions a portfolio must survive.

These lines of work mostly treat the three questions one at a time, although a live portfolio answers them together and the answers interact. The allocation rule decides which positions the rebalancing rule must track, and the market's path between rebalances decides when the next trade comes and from which holdings it starts. An allocation that is easy to interpret helps here, because a change in a holding can then be traced to a change in one stated number. Utility functions from consumer theory offer such a rule when each stock's preference weight is a score with a plain meaning [12–15]. Testing the combined system beyond the historical path requires synthetic markets that reproduce the stylized facts of daily growth rates, in particular heavy tails and volatility clustering [16, 17], while keeping the dependence between stocks and the market [18]. Per-asset hidden Markov generators with a jump mechanism, combined through a common market factor with each asset's own draw rescaled so that the market term does not inflate the asset's total variance, reproduce the heavy tails of individual stocks within a multi-asset market [19]. Such a market makes it possible to run a complete system, with its calibration repeated on each simulated history, on futures that did not occur, including futures with deeper declines than the historical window contains.

In this study, we built a portfolio management system that answers the three questions together and tested it on real and synthetic markets. A selection rule chose the stocks a portfolio could hold, by ranking a score across the whole universe or within each sector, by a bandit within each sector, or at random as a control, and two short lists chosen in advance by the author served as further selections. Each stock's score divided its expected growth rate from a single-index fit to the market by its absolute beta raised to a power that was positive when the market had recently been falling and negative when it had been rising, so that, among stocks with positive expected growth, low-beta stocks scored higher after declines and high-beta stocks after rises, and a hyperbolic tangent kept the score bounded. Utility functions from consumer theory, with the scores as preference weights, and three standard portfolios (minimum variance, tangency and long-short) set the holdings under shared limits on the budget, each stock's weight, portfolio volatility and trading cost. An optional market hedge capped the portfolio's beta, each time the allocation was re-solved, more tightly when the market was weaker and added up to five stocks with zero or negative beta to those the portfolio could hold. A model-predictive controller forecast the value of the holdings with a multi-asset geometric Brownian motion and re-solved the allocation from the current holdings when its horizon had passed, when the value left a band around the forecast, when a drawdown limit was reached, or when a scheduled pick changed the stocks. We calibrated the system on a decade of United States equity prices and ran it on the following twenty months, and we ran it again on simulated price paths of the same period (synthetic futures) built from per-asset hidden Markov generators with jumps, including futures, selected from many simulated paths, with deeper market declines than the historical window contained. On real prices, which stocks a portfolio held mattered more for its terminal wealth than which utility function weighted them, and the forecast band and the drawdown limit brought many re-plans forward of the horizon. Compared with re-planning on the calendar alone, these early re-plans made the worst drawdown smaller in most combinations of rules, on real prices and in the deepest simulated declines, but did not raise terminal wealth consistently. In the deepest simulated declines, every combination of selection and weighting rule ended ahead of both the simulated market index and an equal-weight portfolio of all stocks, random selections included, and did so with the early re-plans switched off as well. The equal-weight portfolio also ended ahead of the index, so part of the margin came from the simulated market and the rest from parts of the system that all

combinations shared, such as the volatility cap, whose separate effect we did not test, rather than from the choice of stocks or the re-planning triggers. The market hedge cost wealth while markets rose and helped when declines persisted, and the synthetic futures favored score-based selection partly because the generator preserved each stock's historical trend. More broadly, the results suggest that portfolio systems are best evaluated as a whole, on synthetic as well as historical markets, and that a synthetic test measures the value of stock selection fairly only when its generator does not build in the trend persistence that selection rules exploit.

2 Related Work

Portfolio allocation is usually posed as a trade-off between expected growth and risk, following mean-variance analysis [1] and its dynamic extensions [2]. Optimized weights amplify errors in the estimated inputs [3]. Practical work therefore shrinks the covariance estimate [4] or adds limits on weights and trading, and judges the result against simple rules such as equal weighting, which is difficult to beat out of sample [5]. The single-index model [18] reduces the inputs to each stock's intercept and market loading, and factor models add further common factors [20]. The allocation used here keeps the single-index inputs for its risk limits and its scores, but replaces the mean-variance objective with utility functions from consumer theory: the Cobb–Douglas form [12], the constant elasticity of substitution form [13], applied once to share counts and once to portfolio weights, and the Stone–Geary form with a minimum holding of each stock [14, 15], four utilities in all. In each utility, a stock's preference weight is its score, a number between -1 and 1 equal to the hyperbolic tangent of the ratio of its expected growth rate from the single-index fit to a penalty on its market loading whose strength depends on the recent direction of the market. Each preference weight is thus a single number that states why a stock is held. Stocks with a negative score receive no preference weight in the long-only utilities. The expected growth rate rewards a stock's trend, and trend-following signals in prices have a long empirical record [21, 22]. The penalty on large loadings, which is harsher when the market is falling, follows the evidence that portfolios tilted away from high-beta stocks have earned higher risk-adjusted performance than portfolios tilted toward them [23]. Three benchmarks run under the same limits: the minimum-variance portfolio, the tangency portfolio and a $130/30$ long-short portfolio, which holds 130% of wealth long and 30% short and so relaxes the long-only constraint of the utility portfolios [24]. Choosing which stocks to hold before weighting them is a separate decision, and one of the selection rules studied here ranks the stocks of each sector with a multi-armed bandit [25], whose arms are the sector's stocks and whose rewards are their growth in excess of the market over past 21-day windows.

With proportional trading costs, the optimal rebalancing policy leaves the holdings alone inside a no-trade region and trades only at its boundary [6, 7]. When expected growth rates are predictable, the optimal policy instead trades a fraction of the gap between the current portfolio and the portfolio the investor would hold without costs [26]. Model-predictive control solves a finite-horizon optimization from the measured state, applies the first decision, and solves again from the next measured state [27]. Its portfolio applications include stochastic formulations with forecast models of asset prices [8], linear control policies for multi-period allocation [28], and convex multi-period trading with costs and limits [9], all of which re-solve at every period. Drawdown control has its own literature, in which exposure is adjusted continuously to keep wealth above a fixed fraction of its running maximum [29]. The controller used here solves a single-period allocation at each re-plan, forecasts the value of the resulting holdings over a 21-day horizon, and keeps the holdings unchanged until a trigger fires. Its default schedule is the horizon itself; a re-plan comes earlier when the measured portfolio value leaves a band around the forecast, so that a forecast miss brings trading forward, or when the drawdown reaches a limit. A re-plan re-solves the allocation from the current holdings and starts a new forecast; it does not by itself reduce exposure.

Evaluating trading rules on the historical record carries a selection risk: when many rules or settings are tried on the same path, the best measured result overstates the result to expect [10, 11]. Synthetic data can test rules on paths that were not observed [30, 31], and the test is informative to the extent that the generator reproduces the stylized facts of daily growth rates, in particular heavy tails and volatility clustering [16, 17]. Block bootstrap resampling keeps short-range dependence [32]. Conditional variance models such as GARCH reproduce volatility clustering [33, 34]. Regime-switching models, including hidden Markov models, represent shifts between market states [35], and neural generators learn the joint distribution directly [36]. Reviews of time-series synthesis note

that a realistic distribution of daily values does not imply realistic behavior over time [37]. Earlier work built the generator used here, a hidden Markov model for each asset whose hidden states are bins of the daily growth rate, with heavy-tailed draws in each state and occasional episodes in which the chain is held in its most extreme states for a random number of days, producing bursts of large moves [19]. The same work combined the draws of many assets through a common market factor, rescaling each asset's own draw so that adding the market term does not inflate the asset's total variance. The present study adds drifting market loadings and sector factors to that combination, and uses the result only as a market for the complete allocation and rebalancing system to trade in; the controller's own forecast remains a geometric Brownian motion. The generator reproduces heavy tails more closely than volatility clustering, and the evaluation carries that limitation.

3 Methods

We tested a rebalancing strategy built from three parts, a score that ranks stocks by expected growth, a constrained utility maximization that sets how many shares of the selected stocks to hold, and a model-predictive controller that decides when to re-plan, in a backtest on recorded prices and in a stress test on synthetic prices. The recorded prices were the daily opens and closes (USD/share) and volumes (shares/day) of the constituents of the S&P 500 index and of the exchange-traded fund SPY, which served as the market and was never held. They came from Polygon.io in three files covering 2014-01-03 to 2024-12-31, calendar 2025, and 2026-01-02 to 2026-09-04, and the trading days were those of SPY. Each name took one of the 11 sectors of the Global Industry Classification Standard (GICS) from a snapshot of the index constituent list. The universe kept every ticker that appeared in all three files with a sector and had a positive open and close on every trading day from 2014-01-03 to 2026-09-04, which left $N = 373$ names with between 15 and 60 per sector (Appendix S2), and we write $s(i)$ for the sector of name i. Because the snapshot lists current members, the universe holds only names that were index members at the snapshot date and traded on every day of 2014–2026, each with its current sector. The prices were split-adjusted by the vendor and not adjusted for dividends, and we used them as given, so neither the portfolios nor SPY received dividends. Each file was split-adjusted as of its own creation date, so for the seven names that split between the creation dates of two files we put the earlier rows on the newest file's basis (Appendix S2).

The calibration window ran from 2014-01-03 to 2024-12-31 (2,767 trading days) and the trading window from 2025-01-02 to 2026-09-04 (420 trading days). Quantities fitted on the calibration window (the prior of the online single-index fit, the growth moments of SPY, the volatility cap, the gain of the market state and the covariance shrinkage) were fitted once and then held fixed. In the trading window every decision was made after the close of day t from data dated on or before day t, and its orders filled at the open of day $t + 1$. We measured performance with continuously compounded growth rates. For a price or wealth series S observed on trading days, the growth rate over the k trading days that end on day t is

$$g(t;k) \;\equiv\; \frac{1}{k\,\Delta t}\,\ln\!\left(\frac{S(t)}{S(t-k)}\right), \tag{1}$$

where $\Delta t = 1/252$ yr is the length of one trading day, so that g has units of yr^{-1}. We write $g_i(t) = g_i(t;1)$ for the daily growth rate of name i and $g_M(t)$ for that of SPY, and the realized growth rate of a run is Eq. (1) applied to its wealth over the $k = 419$ daily steps of the trading window.

The risk-free growth rate came from the 3-month Treasury bill rate DTB3 [38], a bank discount rate d (a fraction, on a 360-day year), converted to the continuously compounded growth rate of a 91-day bill,

$$g_f(t) \;=\; -\,\frac{365}{91}\,\ln\!\left(1-\frac{91\,d}{360}\right), \tag{2}$$

in yr^{-1}. Day t used the last observation dated before t, because the rate for a day is published on the following day, and g_f ranged from 0.036 to 0.044 yr^{-1} over the trading window. It set the growth of cash in the backtest and in the controller's forecast, the excess growth of the tangency benchmark, and the excess growth behind the Sharpe ratio and the regression on SPY.

Each run started with $W_0 = 100{,}000$ USD in cash at the close of 2025-01-02, and its wealth at the close of day t was $W(t) = c(t) + \sum_i n_i(t)\,S_i(t)$, with cash c (USD), fractional shares n_i (negative for a short position) and close prices S_i (USD/share). Positive cash grew by the factor $e^{g_f(t)\,\Delta t}$ each

day. Orders filled at the next open, sells before buys, with all buys scaled down by one common factor when the cash could not cover them, so cash never went negative, and a target position worth at most 0.1% of wealth was set to zero shares. A fill of q shares (signed) of name i at the open price S_i^{o} (USD/share) paid

$$C(q) \;=\; \bar{c}\, S_i^{\mathrm{o}}\, |q| \;+\; \kappa\, \frac{|q|}{A_i}\, S_i^{\mathrm{o}}\, |q| \tag{3}$$

in USD, the half spread $\bar{c} = 5$ bps of the traded value plus a linear impact cost [39] with coefficient $\kappa = 0.001$ on the participation $|q|/A_i$, where A_i (shares/day) is the mean daily volume of name i over the 63 trading days that end on the decision day. There was no commission, and the same half spread prices the turnover limit of the allocation problem. Short positions, which only the 130/30 benchmark took, paid a borrow fee at the rate $b = 0.0075\ \mathrm{yr}^{-1}$ on their value at each close, charged daily, and their sale proceeds were held as cash.

3.1 The score and the selection rules

The score combines a single-index fit of each name to the market with the market state. The daily growth rates $g_i(k)$ of name i and $g_M(k)$ of SPY (Eq. 1, yr^{-1}) follow the single-index model $g_i(k) = \alpha_i + \beta_i\, g_M(k) + \varepsilon_i(k)$ [18], with intercept α_i (yr^{-1}), market loading β_i (dimensionless) and residual $\varepsilon_i(k)$ (yr^{-1}). We fit the model by ordinary least squares on 2014–2024 and, during the trading window, updated (α_i, β_i) every trading day by exponentially weighted least squares (EWLS),

$$(\alpha_i, \beta_i)_t = \arg\min_{a,\,b} \Big[\sum_{k=t_0+1}^{t} \delta^{\,t-k} \big(g_i(k) - a - b\, g_M(k)\big)^2 + \delta^{\,t-t_0} N_0\, \mathrm{E}_0\big[(g_i - a - b\, g_M)^2\big] \Big], \tag{4}$$

where t_0 is the last trading day before the window and $\delta = 2^{-1/63}$ is the daily decay, a half-life of 63 trading days. The prior term counts as $N_0 = 63$ observations, with E_0 the expectation under the 2014–2024 fit and SPY's 2014–2024 mean ($0.106\ \mathrm{yr}^{-1}$) and variance ($7.51\ \mathrm{yr}^{-2}$) of daily growth; its minimizer is the 2014–2024 fit, and its weight decays with the data's. The estimates used on day t include the growth into day t and nothing later. Two market inputs come from SPY's closes. The market growth average $\tilde{g}_M$ (yr^{-1}) is the exponential moving average of SPY's daily growth rate with a span of 252 trading days (smoothing factor $2/253$). The market state is

$$\xi_t = -G \left(\frac{\mathrm{EMA}_{21}(t)}{\mathrm{EMA}_{63}(t)} - 1 \right), \tag{5}$$

where EMA_{21} and EMA_{63} are exponential moving averages of SPY's close with spans of 21 and 63 trading days, and the gain $G = 21.23$ is the reciprocal of the 99th percentile of $|\mathrm{EMA}_{21}/\mathrm{EMA}_{63} - 1|$ over 2014–2024, so $|\xi_t| \le 1$ on 99% of those days. The state ξ_t is positive exactly when the 21-day average lies below the 63-day average, as in a falling market, and it was positive on 75 of the 420 trading days of the window (Fig. 1). The score of name i is

$$\gamma_i = \tanh\!\left(\frac{\alpha_i + \beta_i\, \tilde{g}_M}{\max(|\beta_i|,\, 0.1)^{\xi_t}} \right), \tag{6}$$

with growth rates in yr^{-1} entering the hyperbolic tangent as plain numbers, so $\gamma_i \in (-1, 1)$. The numerator is the name's expected growth rate $\mu_{g,i}$ from the single-index fit. It keeps the sign of β_i, so a name with a negative loading is charged rather than credited for market growth when $\tilde{g}_M > 0$. The denominator, the beta penalty, is positive, so γ_i always has the sign of the expected growth rate and the market state changes only its size. When $\xi_t > 0$ the penalty divides the expected growth of a name with $|\beta_i| > 1$ by a number above one and that of a name with $|\beta_i| < 1$ by a number below one, so the score favors names that move less with the market; when $\xi_t < 0$ the effect reverses and the score favors names that move more; at $\xi_t = 0$ the score is the hyperbolic tangent of expected growth. The floor of 0.1 keeps the penalty finite for names with a loading near zero, and at the window's largest ξ_t of 0.90 it limits their amplification to a factor of $0.1^{-0.90} \approx 7.9$. The same expected growth rate enters the tangency rule and the controller's forecast.

A selection rule chose the names of each decision from the 373 names of the universe in 11 sectors, using data dated on or before the day of the pick only, and the study compared ten rules. Three picked 3 names in each sector, 33 in all: γ *per sector* kept the 3 highest γ_i in each sector, *bandit per sector* kept the 3 names a per-sector bandit ranked highest, and *random per sector* drew 3 names per

sector uniformly at random. Five picked K names from the whole universe with no sector quota, with $K = 13$ in the main comparison and other sizes in the size comparison (Fig. S2). *Global* γ kept the K highest γ_i. *Correlation-aware* γ walked down the same ranking and accepted a name only if the Pearson correlation of its daily growth rates with those of every name already accepted, over the trailing 252 trading days, was at most 0.7, and filled any shortfall from the ranking. *Liquid then* γ kept the K highest γ_i among the 100 names with the largest median daily dollar volume (close times volume) over the trailing 63 trading days, *random liquid 100* drew K names uniformly from that shortlist, and *random whole universe* drew K names uniformly from the universe. The random rules drew afresh at every pick, from a stream seeded by the run seed and the date. The last two rules were fixed lists: the hand-picked 13 (AAPL, MSFT, INTC, MU, AMD, NVDA, GS, BAC, WFC, C, F, GM, JNJ), which the author chose in advance from sectors the author follows and without 2025 data, and the same list with WMT and PG added. Each is one list chosen by one person, and a pick never changed it, although the hedge pool (the stocks with zero or negative beta that join each decision when the market hedge is on) could change.

The bandit rule ran one ϵ-greedy bandit [25] per sector, retrained from scratch at each pick, with each of the sector's A names (between 15 and 60) as one arm. A pull of arm j drew a window start d uniformly from the 252 trading days whose 21-day window ends within the trailing 252 trading days, so every window was complete by the day of the pick, and earned the excess growth rate $x_j(d) = [\ln(S_j(d+21)/S_j(d)) - \ln(S_M(d+21)/S_M(d))]/(21\,\Delta t)$ (yr^{-1}) of name j over SPY, with S_M SPY's close. The bandit made $\min(\max(50A, 500), 5000)$ pulls. The first pull chose an arm uniformly at random; at pull $k > 1$ the bandit chose a uniformly random arm with probability $\epsilon_k = \max\{0.05,\ \min[1,\ k^{-1/3}(A \ln k)^{1/3}]\}$, and otherwise the pulled arm with the highest running-mean reward, ties going to a random rank drawn at the start. With these settings ϵ_k stayed above 0.51 at every pull in every sector, so the floor of 0.05 never bound and each pull explored with probability above one half. The rule kept the 3 arms with the highest mean reward, unpulled arms last (Algorithm S1). Because every arm can be scored on every window, the bandit approximates a ranking of each sector's names by their mean 21-day growth rate in excess of SPY over the trailing year, a sector-neutral momentum screen [40].

3.2 Allocation by constrained utility maximization

At each decision, made after the close of trading day t, the allocator chose a continuous number of shares $n_i \geq 0$ for every name i in the decision's name set $\mathcal{N}$, which held the names picked by the selection rule and, with the market hedge on, a pool of hedge names. With S_i the closing price of name i (USD/share), W the wealth at the close (USD) and $w_i = n_i S_i / W$ the weight of name i (dimensionless), the allocator solved

$$\begin{aligned} \max_{n \geq 0} \quad & U(n;\gamma) \\ \text{subject to} \quad & \textstyle\sum_{i\in\mathcal{N}} n_i S_i \leq W, \qquad w_i \leq \bar{w}_i \ \ (i \in \mathcal{N}), \qquad \sqrt{w^\top \Sigma_g\, w} \leq \sigma_{\mathrm{SPY}}, \\ & \bar{c}\,\bar{S} \textstyle\sum_{i\in\mathcal{N}} |n_i - n_i^{\mathrm{prev}}| \leq B, \qquad \sum_{i\in\mathcal{N}} \beta_i w_i \leq \bar{\beta} \ \ \text{(hedge on)}, \end{aligned} \tag{7}$$

where $U(n;\gamma)$ is the weighting rule's utility and γ_i the score of name i (Eq. 6); wealth not invested is held as cash, which has a market loading of zero. The single-name cap was $\bar{w}_i = \max(0.10,\ 1.5/K)$ for a name picked by the selection rule, with K the number of picks, and 0.10 for hedge-pool names and 130/30 short names, so a pick could hold up to 10% of wealth in the 33-name baskets, 11.5% in the 13-name baskets and 30% at $K = 5$. In the volatility cap, Σ_g (yr^{-2}) is the covariance of the names' daily growth rates, the single-index market term plus the sample residual covariance over the trailing 126 trading days (Eq. 9 with its calibrated shrinkage of zero), and $\sigma_{\mathrm{SPY}} = 2.741\ \mathrm{yr}^{-1}$ is the sample standard deviation of SPY's daily growth rate over 2014–2024, an annualized volatility of 17.3%; the solver imposed the cap as a second-order cone on $L^\top w$ with $LL^\top = \Sigma_g + 10^{-8} I$. The turnover budget caps the half-spread cost of the share changes from the holdings n^{prev} before the decision, priced at the half spread $\bar{c}$ of the cost model (Eq. 3) times the mean closing price $\bar{S}$ (USD/share) of the names in $\mathcal{N}$, with $\bar{c} = 5 \times 10^{-4}$ and $B = 0.10\,W$ (USD). For a long-only book this budget can bind only when $\bar{S}$ exceeds 100 times the cheapest name's price. The sum in the budget ran over the names whose shares the solver chose; the sale of a name with $\gamma_i \leq 0$ down to its floor and the exit of a held name outside $\mathcal{N}$ were not counted. The study set no sector limit. With the hedge on, the portfolio's market loading $\sum_i \beta_i w_i$, computed with each name's signed online loading (Eq. 4), was capped at $\bar{\beta} = \max(0,\ 1 - \xi_t)$, where ξ_t is the market state (Eq. 5). The cap equals

SPY's loading of one at $\xi_t = 0$, tightens as ξ_t rises in a falling market, reaches zero at $\xi_t \geq 1$ and rises above one when $\xi_t < 0$; over the trading window ξ_t ranged from -0.80 to 0.90, so $\bar{\beta}$ ranged from 1.80 down to 0.10 (Fig. 1). The cap applied to the targets of each re-plan. Between re-plans the shares were held, so the book's loading moved with prices and the online loadings and was not held under the cap.

The hedge pool was refreshed at every pick of the basket. A universe name was eligible when its online market loading was at most zero and its down-day beta was at most zero, where the down-day beta is the ordinary least squares slope, with intercept, of the name's daily growth rate on SPY's over the days of the trailing 252 trading days on which SPY's growth rate was negative. The pool held up to $H = 5$ eligible names, most negative down-day beta first, and joined every decision's name set on top of the picks, a name both picked and pooled counting once as a hedge name. In the rules that use γ, a pool name with $\gamma_i > 0$ entered the utility like any other name, while a pool name with $\gamma_i \leq 0$ had no utility term and was a free variable above its floor, so the allocator bought it only when the β cap or the volatility cap made it worth its share of the budget. Because the objective is flat along such a name's weight whenever no limit binds, it subtracted a tie-break cost of 10^{-6} per unit of hedge weight, and a second, lexicographic solve then fixed every other name's shares at their first-stage values and minimized the total weight of the free hedge names under the same limits. The utility depends only on the fixed names, so the second solve left it unchanged and held a hedge name above its floor only as far as a binding limit required; if the second solve failed, the first-stage shares stood. With the hedge off, a decision had neither the β cap nor the pool and was otherwise identical. We added the β cap and the pool to the design after inspecting how the online loadings behaved in the February–April 2025 sell-off; every decision used trailing data only, and every result is reported with the hedge on and off.

We compared four utilities built on the score. With $\mathcal{P} = \{i \in \mathcal{N} : \gamma_i > 0\}$ the names with a positive score, they are

$$\begin{aligned} U_{\mathrm{CD}} &= \sum_{i \in \mathcal{P}} \gamma_i \ln n_i, & U_{\mathrm{CES}} &= \Big(\sum_{i \in \mathcal{P}} \gamma_i \, n_i^{\rho} \Big)^{1/\rho}, \\ U_{\mathrm{SG}} &= \sum_{i \in \mathcal{P}} \gamma_i \ln \left(n_i - b_i \right), & U_{\mathrm{CESw}} &= \Big(\sum_{i \in \mathcal{P}} \gamma_i \, w_i^{\rho} \Big)^{1/\rho}, \end{aligned} \tag{8}$$

the Cobb–Douglas (CD) [12], Stone–Geary (SG) [15, 14] and constant elasticity of substitution (CES) [13] utilities, the last on share counts (CES) and on weights (CESw). Both CES rules used the elasticity $\eta = 0.5$, so $\rho = 1 - 1/\eta = -1$, and SG used the baseline $b_i = 0.01\, W/S_i$ shares, a core position worth 1% of wealth that SG holds in every name of $\mathcal{N}$ ($n_i \geq b_i$). Under the budget alone each utility has a closed-form optimum (Table 1; derivations in Appendix S1). CD splits the invested wealth in proportion to γ_i. CES sets $n_i \propto (\gamma_i / S_i)^{\eta}$, so its weights $w_i \propto \gamma_i^{\eta} S_i^{1-\eta}$ depend on each name's share price when $\eta \neq 1$, whereas CESw sets $w_i \propto \gamma_i^{\eta}$ and does not; both equal CD at $\eta = 1$, approach equal share counts (CES) or equal weights (CESw) [5] as $\eta \to 0$, and concentrate in the name with the largest γ_i / S_i (CES) or γ_i (CESw) as $\eta \to \infty$. SG first buys every baseline and then splits the remaining wealth in proportion to γ_i. With the limits of Eq. 7 there is no closed form, and each utility becomes a convex cone program [41]: CD and SG use exponential cones, the two CES rules use power cones, the volatility cap is a second-order cone and the other limits are linear. We built each program in JuMP $1.32.1$ [42] and solved it with the interior-point solver Clarabel $0.11.1$ [43] (absolute and relative duality-gap tolerances 10^{-12}, feasibility tolerance 10^{-9}, at most 500 iterations), accepting its optimal and almost-optimal statuses; when Clarabel reached neither, the same program went to SCS $2.7.0$ [44] with its default settings. To check the solver against the closed forms, we allocated a new book of $100{,}000$ USD on 2 January 2025 across the hand-picked 13-name list, with the hedge off and every limit except the budget removed (the turnover budget stayed and did not bind). Over 17 solves (CD, and CES and CESw at eight values of η from 0.05 to 20), the largest gap between a solver weight and its closed form was 1.0×10^{-5}, below the tolerance of 10^{-3} we required.

The rules treat names with $\gamma_i \leq 0$ differently. CD, CES and CESw leave them out of the utility and hold them at a floor of $\varepsilon = 10^{-3}$ shares, which disappears when targets worth at most 0.1% of wealth are set to zero before trading, so the three rules held only names with $\gamma_i > 0$ and any hedge names the limits required; when no name in $\mathcal{N}$ had $\gamma_i > 0$, they skipped the solve and moved the book to cash. SG holds a name with $\gamma_i \leq 0$ at its baseline b_i; with no positive score it held the baselines, adding the least hedge weight that met the β cap, and it reported the problem infeasible when the

baselines exceeded the budget or broke another limit. The baselines use 1% of wealth per name, so SG cannot hold more than 100 names. In the large-book runs every SG decision at $K = 200$ was infeasible and the book stayed in cash, and at $K = 100$ most SG decisions were infeasible and kept the previous book.

Three benchmarks ran under the same limits. Minimum variance (MinVar) and tangency (Tangent) [1, 45] do not use γ and could hold any name of $\mathcal{N}$, with $n_i \geq 0$ and no floors. Each followed a two-fund form: it found the largest invested fraction $s = \min(1, \max \sum_i w_i)$ that the budget, single-name, β and turnover limits allowed without the volatility cap, chose the mix at $\sum_i w_i = s$ that minimized $w^\top \Sigma_g w$ (MinVar) or maximized $(\mu_g - g_f \mathbf{1})^\top w / \sqrt{w^\top \Sigma_g w}$ (Tangent), and scaled the mix by $\min(1, \sigma_{\mathrm{SPY}}/\sigma_{\mathrm{mix}})$ onto the volatility cap, holding the rest in cash. Here $\mu_{g,i} = \alpha_i + \beta_i \tilde{g}_M$ (yr^{-1}) is the expected growth rate of name i from the score's inputs and g_f (yr^{-1}) the risk-free growth rate (Eq. 2). Tangent solved its ratio as a quadratic program after a Charnes–Cooper change of variables [46] and fell back to MinVar when no feasible mix had positive excess growth. The $130/30$ long-short Cobb–Douglas rule (LS) [24] maximized $\sum_{\gamma_i > 0} \gamma_i \ln n_i^+ + \sum_{\gamma_i < 0} |\gamma_i| \ln n_i^-$ over long shares n_i^+ and short shares n_i^- ($n_i = n_i^+ - n_i^-$), so it shorted names with $\gamma_i < 0$ weighted by $|\gamma_i|$. Long weights summed to at most 1.30, short weights to at most 0.30 in absolute value and the net weight to at most 1, so cash stayed nonnegative; the single-name cap applied to $|w_i|$, the volatility and β caps to the signed weights and the turnover budget to the signed share changes. Its long side held the picks with $\gamma_i > 0$ and the hedge pool, and its short side held the short candidates, up to 3 names per sector with the most negative $\gamma_i < 0$ outside the pool, chosen at each pick of the basket and kept only while $\gamma_i < 0$. A hedge name was never shorted, since shorting a name with a negative loading raises the portfolio's loading. With no short candidate LS ran as long-only CD, and with no positive score it moved to cash.

Algorithm 1 lists one allocation decision. A decision whose program a solver certified infeasible, that neither solver solved, or that produced non-finite shares kept every holding and placed no orders. Apart from SG with 100 or more names, this happened in under 1% of the decisions of every grid of real-data runs.

3.3 Model-predictive rebalancing

The controller decides when to re-plan the book and leaves it alone in between. It follows model-predictive control in the receding-horizon sense [47, 27]. At each re-plan it measures the state of the account (the shares n_i, the cash c, the wealth $W(t)$, the close prices $S_i(t)$ and the online estimates of that day), solves the constrained utility maximization of Algorithm 1 for target shares $n^\star$, and predicts the distribution of the wealth of holding $n^\star$ over the next $T = 21$ trading days with a multi-asset geometric Brownian motion (GBM). It then compares the measured wealth with the prediction every day and re-plans from the newly measured state when the horizon elapses, when the wealth leaves the predicted band, or when the book falls a set fraction below its peak. The plan over the horizon is to hold $n^\star$ without further trades, and the GBM is the controller's only model of prices. The plant is the price process that the plan meets, the recorded prices in the backtest and synthetic prices from per-asset JumpHMM generators [19] in the stress test, where the controller still forecasts with the GBM and never sees the generator. Multi-period trading by convex optimization [9] also re-solves from the measured state, but it plans a sequence of trades over the horizon and executes the first of them every period. Our controller instead solves a single-period allocation, uses the forecast as a reference for the realized wealth, and trades only when a trigger fires, so the band around the forecast acts as a no-trade region [7] placed on the portfolio value relative to its forecast rather than on the weights.

At a re-plan on day t the drift of name i is its expected growth rate $\mu_{g,i} = \alpha_i(t) + \beta_i(t)\, \tilde{g}_M(t)$ (yr^{-1}). Here $\alpha_i(t)$ (yr^{-1}) and $\beta_i(t)$ (dimensionless) are the online single-index estimates of name i against SPY [18], fitted by exponentially weighted least squares with a half-life of 63 trading days from a prior worth 63 pseudo-observations at the calibration fit (Appendix S2), and $\tilde{g}_M(t)$ (yr^{-1}) is the exponential moving average of the daily growth rate of SPY with a span of 252 trading days. The covariance Σ_g (yr^{-2}) of the daily growth rates of the names in the allocation problem is

$$\Sigma_g \;=\; \sigma_M^2(t)\, \beta\beta^\top \;+\; (1-\lambda)\, \hat{\Omega}_\varepsilon \;+\; \lambda\, \mathrm{diag}\big(\hat{\Omega}_\varepsilon\big), \tag{9}$$

where β is the vector of the loadings $\beta_i(t)$, $\sigma_M(t)$ (yr^{-1}) is the exponentially weighted standard deviation of the daily growth rate of SPY with the same 63-day half-life, and $\hat{\Omega}_\varepsilon$ (yr^{-2}) is the sample

Algorithm 1 One allocation decision after the close of day t.

Require: Closes through day t; held shares n^{prev}; wealth W (USD); weighting rule U; selection rule; hedge on or off; whether a pick of the basket is due (Algorithm 2).
Ensure: Target shares $n^\star$ and orders for the open of day $t+1$, or the kept book and no orders.

$\gamma_i \leftarrow$ Eq. 6 for every universe name, from the day's (α_i, β_i) (Eq. 4), $\tilde{g}_M$ and ξ_t (Eq. 5)
if a pick is due **then**
  $\mathcal{K} \leftarrow$ picks of the selection rule; $\mathcal{H} \leftarrow$ hedge pool if the hedge is on, else $\emptyset$
  $\mathcal{S}^- \leftarrow$ up to 3 names per sector with the most negative $\gamma_i < 0$, outside $\mathcal{H}$ {LS only; else $\emptyset$}
end if
$\mathcal{N} \leftarrow \mathcal{K} \cup \mathcal{H}$; for LS, $\mathcal{N} \leftarrow \{i \in \mathcal{K} : \gamma_i > 0\} \cup \mathcal{H} \cup \{i \in \mathcal{S}^- : \gamma_i < 0\}$
Build $\mu_{g,i} = \alpha_i + \beta_i \tilde{g}_M$, Σ_g (Eq. 9), $\bar{w}_i$, $\bar{c}\,\bar{S}$, $B = 0.10\,W$, and $\bar{\beta} = \max(0, 1 - \xi_t)$ (hedge on) or $\bar{\beta} = \infty$
if U uses γ and no $i \in \mathcal{N}$ has $\gamma_i > 0$ **then**
  $n^\star \leftarrow$ floors: ε (CD, CES, CESw), 0 (LS), or b_i plus the least hedge weight that meets $\bar{\beta}$ (SG) {no utility solve}
else
  Solve Eq. 7 for U with Clarabel {MinVar, Tangent: mix, then scale}
  if Clarabel is neither optimal nor almost optimal **then**
    Solve the same program with SCS
  end if
  if a hedge name has $\gamma_i \le 0$ and U uses γ **then**
    Fix the other names' shares; minimize $\sum_{i \in \mathcal{H}} w_i$ under the same limits
  end if
end if
if the program is infeasible, both solvers failed, or some $n_i^\star$ is not finite **then**
  $n^\star \leftarrow n^{\text{prev}}$; **return** no orders {keep the whole book}
end if
$n_i^\star \leftarrow 0$ where $|n_i^\star S_i| \le 0.001\,W$, and for held names outside $\mathcal{N}$
return orders $n^\star - n^{\text{prev}}$, filled at the open of day $t+1$

covariance of the residuals $\varepsilon_i(s) = g_i(s) - \alpha_i(t) - \beta_i(t)\, g_M(s)$ over the $L_\Omega = 126$ daily growth rates that end on day t. The weight $\lambda \in [0, 1]$ shrinks the residual covariance toward its diagonal, from the full sample covariance at $\lambda = 0$ to the covariance of the single-index model, the target of Ledoit and Wolf [48], at $\lambda = 1$. We chose it on the calibration window by the QLIKE loss [49] of 21-day variance forecasts for random portfolios, which selected $\lambda = 0$ (Appendix S2). The allocation problem uses the same μ_g and Σ_g in its volatility limit and in its minimum-variance and tangency benchmarks, so each book is forecast with the model that sized it.

From the target shares $n^\star$ and the close prices $S_i(t)$, the controller simulates $N_p = 1{,}000$ paths over the next $T = 21$ trading days. On path p and day $\tau = 1, \ldots, T$ after the re-plan, the log price of each name advances by Δt times a vector of growth rates drawn from a multivariate normal distribution with mean μ_g and covariance Σ_g, independently across days and paths,

$$\ln S_i^{(p)}(\tau) \;=\; \ln S_i^{(p)}(\tau - 1) \;+\; g_i^{(p)}(\tau)\,\Delta t, \qquad g^{(p)}(\tau) \sim \mathcal{N}\big(\mu_g, \Sigma_g\big), \qquad S_i^{(p)}(0) = S_i(t), \tag{10}$$

a multi-asset GBM [50] written in growth rates. Because μ_g is the expected growth rate of the log price, the drift needs no Itô correction. The wealth of the plan on path p is

$$W^{(p)}(\tau) \;=\; \sum_i n_i^\star\, S_i^{(p)}(\tau) \;+\; c^\star\, e^{g_f(t)\,\tau\,\Delta t}, \qquad c^\star \;=\; W(t) - \sum_i n_i^\star\, S_i(t), \tag{11}$$

where $c^\star$ (USD) is the cash left after paying for the targets at the close of the decision day. The forecast is the sample mean $\mu_W(\tau)$ and standard deviation $\sigma_W(\tau)$ (USD) of $W^{(p)}(\tau)$ across the paths. It omits the trading costs, the overnight price change before the fill and, for a 130/30 book, the borrow fee, and its paths are seeded with the run seed and the decision date, so a run reproduces its forecasts exactly.

The forecast defines a band $\mu_W(\tau) \pm z\,\sigma_W(\tau)$ with $z = 1.96$, which would contain the wealth on about 95% of days if the forecast wealth were normal and the model were correct. With τ the number

of trading days since the last re-plan, the controller marks the wealth $W(t)$ at each close, after that day's fills, costs, cash growth and fees, updates its running peak $W^{\mathrm{pk}}(t)$, and checks three triggers in turn, acting on the first that holds. The drawdown trigger fires when $1 - W(t)/W^{\mathrm{pk}}(t) > D_{\max}$ with $D_{\max} = 0.08$, the horizon trigger when $\tau \geq T$, and the band trigger when $W(t)$ lies outside $\mu_W(\tau) \pm z\,\sigma_W(\tau)$, so the wealth measured τ days after a re-plan is compared with the forecast for τ days ahead on days $\tau = 1, \ldots, T-1$. A drawdown re-plan sets the peak back to the current wealth, so the trigger fires once for each fresh loss of $D_{\max}$ rather than on every day that the book stays below an old peak. We adopted this reset after an earlier pass of the Tier 1 grid (the three per-sector selection rules crossed with the seven weighting rules) without it, in which 88.5% of the decisions were drawdown re-plans. A re-plan that leaves no position worth more than 0.1% of wealth sets no band, and the next re-plan then comes from the drawdown trigger, from a basket change or, after 21 trading days, from a cash revisit (when no name had a positive score) or the horizon.

The selection rule picks the basket on the first day, every 21 trading days after its previous pick, and at every band exit. With the hedge on, each pick also refreshes the hedge pool, and for the 130/30 benchmark the short candidates. When a timer pick changes the set of names on a day when no trigger fired, the controller re-plans as well. Because band exits depend on each run's own wealth path, the dates of the picks, and with them the baskets, could differ across weighting rules within a seed. At a re-plan it solves the allocation problem from the shares actually held, so the turnover limit counts the changes from that book, queues orders to $n^\star$ for the next open, forecasts the new book from the measured wealth $W(t)$ and resets τ to 0. Each forecast therefore starts from the measured wealth rather than from the previous forecast, and the horizon recedes with every re-plan. If the solve fails, the controller keeps every holding, queues no orders and forecasts the kept book. One episode of a real run, three consecutive forecasts and the re-plans that ended them, shows the mechanism (Fig. 2), and Algorithm 2 gives the daily loop of a run.

3.4 Backtest design

A backtest run fixes a selection rule, a weighting rule, the hedge setting and a seed, and replays Algorithm 2 over the 420 trading days of the trading window. We used the seeds 5660 to 5689, 30 runs per cell of an experiment grid. The seed enters in three places: the random selection rules draw each basket from a stream seeded with the run seed and the pick date, each sector's bandit seeds its exploration with the run seed, the pick date and the sector, and the forecast paths of every re-plan are seeded with the run seed and the decision date. For the deterministic selection rules (the γ-based rules and the hand-picked lists) the seed changes only the forecast paths, and through them the band exits and so the dates of the re-plans and picks. Each run with the hedge on has a twin with the hedge off and the same seed, identical except that the beta cap and the hedge pool are removed.

We ran nine grids over the full trading window, and all 18,238 distinct runs finished (Table S3). The Tier 1 grid crossed the seven weighting rules with three per-sector selection rules (3 names per sector, 33 names) and the hedge on and off (854 runs). The fixed-schedule grid repeated its hedge-on half with basket picks only on the 21-day timer (427 runs), while band exits and drawdown stops still triggered re-plans, unlike the calendar-only grid. Because the random and bandit streams depend only on the seed and the pick date, and the γ ranking and the hedge pool only on market data, every weighting rule in that grid held the same baskets within a seed. The universe-wide grid applied five selection rules that pick K names from the whole universe, at $K = 5, 8, 13, 21$ and 34 (4,410 runs), and the hand-picked grid ran the two hand-picked lists (28 runs). A seed rerun gave each deterministic rule all 30 seeds (4,200 runs, 140 of which repeat earlier runs), and the large-book grid extended global γ and random picks from the whole universe to $K = 100$ and 200 (1,680 runs). A calendar-only grid repeated the 4,200 runs of the main comparison, described next, with the band exit and the drawdown stop switched off ($z = \infty$ and $D_{\max} = \infty$), so that a run re-planned only when the 21-day horizon elapsed, and paired each run with the triggered run of the same selection rule, weighting rule, hedge setting and seed (4,200 runs). Two sensitivity grids varied the CES elasticity η, one on the Tier 1 grid with the single-name cap on and off (2,562 runs) and one on the hand-picked 13-name list under the budget alone (17 runs). This paper does not report results from the fixed-schedule grid or from the elasticity grid on the Tier 1 rules, and the budget-only grid served as the solver check against the closed forms.

The main comparison crosses ten selection rules with the seven weighting rules, with the hedge on and off. The ten are γ, bandit and random per sector (33 names each), global γ, correlation-aware

Algorithm 2 Daily loop of one backtest run. A run fixes a selection rule, a weighting rule, the hedge setting and a seed. Every decision uses data dated on or before the close of its day.

Require: trading days $t_1, \ldots, t_{420}$; calibration fit on 2014–2024; W_0; T, z, $D_{\max}$, N_p
Ensure: daily wealth $W(t)$ and the log of re-plans

$c \leftarrow W_0$, $n \leftarrow 0$, $W^{\mathrm{pk}} \leftarrow W_0$, $\tau \leftarrow 0$; no band, no pending orders
for each trading day $t = t_1, \ldots, t_{420}$ **do**
  if orders are pending **then**
    fill them at the open: sells first, then buys scaled by one factor in $[0, 1]$ so that $c \geq 0$; pay the costs of Eq. (3)
  end if
  if $t > t_1$ **then**
    $c \leftarrow c\, e^{g_f(t)\Delta t}$ if $c > 0$; charge the borrow fee on short positions; $\tau \leftarrow \tau + 1$
  end if
  mark $W(t) \leftarrow c + \sum_i n_i S_i(t)$ at the close; $W^{\mathrm{pk}} \leftarrow \max(W^{\mathrm{pk}}, W(t))$
  update $\alpha_i, \beta_i, \sigma_M, \tilde{g}_M, \xi$ and γ_i with the growth rates into day t
  reason $\leftarrow$ the first that holds of: first day; cash revisit (last re-plan went to cash and $\tau \geq 21$); drawdown ($1 - W(t)/W^{\mathrm{pk}} > D_{\max}$); horizon ($\tau \geq T$); band exit ($|W(t) - \mu_W(\tau)| > z\,\sigma_W(\tau)$); otherwise none
  if first day, or 21 trading days since the last pick, or reason is band exit (except on the fixed schedule) **then**
    pick the basket with the selection rule; refresh the hedge pool and the short candidates
    if reason is none and the set of names changed **then**
      reason $\leftarrow$ basket change
    end if
  end if
  if reason is not none **then**
    build μ_g and Σ_g (Eq. (9)) for the names of the allocation problem
    solve the allocation problem (Algorithm 1) from n, $W(t)$, $S(t)$: targets $n^\star$
    if the solve failed **then**
      $n^\star \leftarrow n$; queue no orders
    else if $t < t_{420}$ **then**
      queue orders to $n^\star$ for the open of the next day, each target worth at most $0.001\,W(t)$ set to zero shares
    end if
    simulate N_p paths of Eqs. (10) and (11) for the book $n^\star$; store $\mu_W(\tau)$, $\sigma_W(\tau)$, $\tau = 1, \ldots, T$
    set the band from the forecast (no band if the book holds no position above $0.001\,W(t)$); $\tau \leftarrow 0$
    if reason is drawdown **then**
      $W^{\mathrm{pk}} \leftarrow W(t)$
    end if
  end if
end for

γ, liquid then γ, random from the 100 most liquid names and random from the whole universe ($K = 13$ each), and the two hand-picked lists. The 13-name list (AAPL, MSFT, INTC, MU, AMD, NVDA, GS, BAC, WFC, C, F, GM, JNJ) and its 15-name extension with WMT and PG were chosen in advance by the author, from sectors the author follows and without 2025 data. Each is one list chosen by a person. Each of the 140 cells holds 30 runs, and we report the median over the runs of each cell, never a median pooled across selection rules. For each run we computed the terminal wealth ratio W_T/W_0, with W_T the wealth at the close of 2026-09-04, the realized growth rate $\ln(W_T/W_0)/(419\,\Delta t)$, the maximum drawdown $\max_t \left(1 - W(t)/\max_{s \leq t} W(s)\right)$, and the annualized Sharpe ratio [51] $\bar{y}/(\sqrt{\Delta t}\, s_y)$, where $\bar{y}$ and s_y are the mean and standard deviation of the daily excess growth $y(t) = g_W(t) - g_f(t)$, with g_W the daily growth rate of the run's wealth. We also recorded the intercept α_W (yr^{-1}) and slope β_W of an ordinary least-squares regression of $y(t)$ on the daily excess growth of SPY, and the number of re-plans by trigger. Two passive benchmarks received the same metrics. SPY buy-and-hold puts W_0 into SPY at the close of 2025-01-02, and

the equal-weight universe puts $W_0/373$ into each name at that close and never rebalances. Both are valued at the closes and pay no trading costs. The universe-wide selection rules, their thresholds and the grid of K were designed after the Tier 1 results were known, so their comparison is exploratory. The 30-seed reruns and the large-book runs were added afterwards to measure how much the forecast seed moves a deterministic rule and whether the comparison of global γ with random picks holds for larger books. The calendar-only grids, on real prices and on the synthetic stress sets, were added last, after all other results were known, to measure what the two early triggers changed.

3.5 Synthetic markets from per-asset JumpHMM generators

The backtest covers one real history of 420 trading days with one sell-off, so it cannot show how the method behaves in conditions that did not occur in that window. We therefore built synthetic markets for the whole universe and ran the complete method on each. Every synthetic series came from JumpHMM, the hidden Markov generator with jumps of Alswaidan and Varner [19] (Fig. 4). A JumpHMM is a hidden Markov model [52, 35] whose states are intervals of the daily growth rate bounded by equal-probability quantiles of a fitted Laplace distribution, whose transition matrix is estimated by counting, and whose draw for each day is Student-t with $\nu = 5$ degrees of freedom, located and scaled by the training days in the selected state. Because the counted chain mixes within a few days, JumpHMM adds jump episodes to reproduce the slow decay of the autocorrelation of absolute growth rates [16, 17, 53]. On a day with no episode under way, an episode starts with probability ϵ and lasts a Poisson number of days with mean λ, and on each of its days the state is drawn uniformly from the N_{tail} most negative states with probability p_{neg} or from the N_{tail} most positive states otherwise, which makes an episode a sustained burst of large moves of both signs rather than a single price jump [54, 55]. The pair (ϵ, λ) minimizes on a grid, over simulated paths that contain an episode, the squared error in the absolute-growth autocorrelation at lags 1 to 25 plus 0.20 times the squared error in excess kurtosis. We write ϵ for the jump probability and ε for a growth-rate draw.

Every JumpHMM here used 50 states, $\nu = 5$, $p_{\text{neg}} = 0.52$ and $N_{\text{tail}} = 5$, with jumps enabled (Table S4). Two settings differ from the SPY generator of the JumpHMM paper. That paper used 100 states, so its tail sets held the most extreme 5% of states on each side, whereas with 50 states they hold 10%. Its multi-asset comparisons either disabled jumps or gave every asset the SPY values of (ϵ, λ), whereas each sector and name model here was tuned on its own series. As for SPY in that paper, the tuned ϵ sat at the lower edge of its grid, 10^{-4} per day, for all 11 sector models and 370 of the 373 name models, while λ spread over its whole grid of 10 to 120 trading days. We used JumpHMM.jl version 0.1.0 at a pinned revision (commit `aaed969`) without modification and called only its fitting, tuning and simulation functions. Its multi-asset models either reorder each asset's draw to the ranks of a copula sample or add bootstrapped or Gaussian regression residuals to a market factor, and neither keeps each name's own JumpHMM path in order, so we wrote the composition in the study code.

The synthetic market combined three kinds of JumpHMM (Fig. 5). The market model, fitted in earlier work to SPY closing prices over 2014–2024 (the study's series), has $\epsilon = 10^{-4}$ and $\lambda = 118.9$ trading days. Its draw is the growth rate $g_m(t)$ (yr^{-1}) of synthetic SPY on day t, with a flag $J(t)$ that marks the days of a jump episode. For the sector models, we took the residual $e_i(t) = g_i(t) - \hat{\alpha}_i - \hat{\beta}_i\, g_M(t)$ (yr^{-1}) of each name from its ordinary least-squares (OLS) single-index fit to SPY over 2014–2024 [18], where $g_i(t)$ and $g_M(t)$ are the real growth rates of name i and SPY and $\hat{\alpha}_i$ (yr^{-1}) and $\hat{\beta}_i$ are the fitted intercept and slope. The sector factor $f_s(t)$ (yr^{-1}) is the equal-weight mean residual of the names in sector s, and because JumpHMM.jl fits price series, each sector model was fitted to the compounded index $\exp(\Delta t \sum_{u \le t} f_s(u))$. As in the JumpHMM paper, g_i and f_s denote real series when we estimate and synthetic ones when we compose. Finally, one JumpHMM was fitted to each name's 2014–2024 closing prices. We refitted all 373 name models on the study's data rather than reuse models stored from earlier work, whose training file held nine unadjusted stock splits, and each model's draws had a standard deviation within 0.5 to 2 times the real one (median ratio 1.037).

For name i in sector $s = s(i)$ on day t, let α_i (yr^{-1}) be its synthetic intercept, $\beta_i(t)$ its market loading, δ_i its sector loading, $f_s(t)$ (yr^{-1}) the draw of its sector's model, $\tilde{\varepsilon}_i(t)$ (yr^{-1}) the centered draw of its own model and s_i a scale on that draw. The synthetic growth rate is

$$g_i(t) = \alpha_i + \beta_i(t)\, g_m(t) + \delta_i\, f_s(t) + s_i\, \tilde{\varepsilon}_i(t), \qquad t = 1, \ldots, 3{,}186, \tag{12}$$

over every trading day of the real calendar from 2014-01-03 to 2026-09-04, so one continuous path covers both the calibration and the trading window and every date-based setting of the method applies unchanged. Closes start at each name's real close on 2014-01-03 and compound as $S_i(t) = S_i(t-1)\exp(g_i(t)\,\Delta t)$, and synthetic SPY compounds from its real close with $g_m(t)$ (Algorithm 3).

The intercept is the name's real 2014–2024 OLS intercept pulled toward the universe mean $\bar{\alpha}$, $\alpha_i = \bar{\alpha} + \phi(\hat{\alpha}_i - \bar{\alpha})$ with $\phi = 0.6$, constant in time, so names keep their 2014–2024 order at 60% of their real distance from the mean. The pull reflects how little of the real spread persists. In 2014–2024 the spread of $\hat{\alpha}_i$ across names was 0.076 yr^{-1}, against a root-mean-square standard error of 0.075 yr^{-1}, and the intercepts of 2014–2019 and 2020–2024 correlated 0.175 across names. We fitted ϕ so that OLS intercepts estimated on synthetic paths reproduced both statistics (Appendix S3). We dropped a planned drift in α_i, of the same form as that of $\beta_i(t)$, because its fitted size was zero and real intercepts changed between the halves of 2014–2024 by no more than their estimation noise. We also set aside the raw estimates, which overstate the real spread and persistence, and a common intercept for every name, which removes them. Because each name keeps a persistent intercept, the synthetic markets favor selection by γ more than real 2025–2026 did, whose intercepts correlated -0.12 across names with real 2014–2024 intercepts, against 0.10 on synthetic paths (Table S6).

The market loading drifts as $\beta_i(t) = \hat{\beta}_i + x_i(t)$, where $x_i(t)$ is an exponential moving average (EMA) of independent standard normal shocks $z_i(t)$ with weight a, half-life h_β (trading days) and scale k_i:

$$x_i(t) = (1-a)\,x_i(t-1) + a\,k_i\,z_i(t), \qquad a = 1 - 2^{-1/h_\beta}. \tag{13}$$

The first value is drawn from the stationary distribution, whose standard deviation is $k_i\sqrt{a/(2-a)}$, so the synthetic calibration window already contains drift. We calibrated the drift against the method's own online estimate $\hat{\beta}_i(t)$ (exponentially weighted least squares with a half-life of 63 trading days) rather than an assumed truth. Each k_i reproduces on synthetic 2014–2024 the standard deviation of the name's real $\hat{\beta}_i(t)$, and the common half-life reproduces the median lag-252 autocorrelation of $\hat{\beta}_i(t)$, 0.206 (Appendix S3). The fit chose $h_\beta = 126$ trading days, and the stationary standard deviation of the drift had a median of 0.26 across names (interquartile range 0.20 to 0.34). The synthetic $\hat{\beta}_i(t)$ thus carries the estimation noise of the real one, and a name whose $\beta_i(t)$ drifts below zero becomes a hedge candidate on that path.

The sector loading δ_i is the OLS slope of the name's real residual $e_i(t)$ on the mean residual of the other names in its sector, which keeps its own noise out of the regressor in sectors with as few as 15 names, times one factor κ_s per sector. The slope on this noisy proxy for the common factor is biased toward zero, yet it multiplies a factor drawn with the variance of the full sector mean, and the unscaled loadings undercounted the shared co-movement in every sector (in Communication Services, a mean residual correlation of 0.052 against a real 0.094). We chose κ_s so that $\delta_i\delta_j\sigma_s^2$, averaged over the pairs $i \neq j$ of the sector, equals the mean pairwise covariance of the real residuals, where σ_s^2 (yr^{-2}) is the variance of the real sector factor. The factors ran from 1.02 in Financials to 1.34 in Communication Services (Table S5), and the sector models were drawn independently of each other and of the market.

The own draw is centered on its model's long-run mean, $\tilde{\varepsilon}_i(t) = \varepsilon_i(t) - \bar{\varepsilon}_i$, where $\varepsilon_i(t)$ is a path of the name's JumpHMM and $\bar{\varepsilon}_i$ (yr^{-1}) the mean of one 10^6-day draw with a fixed seed. Let σ_m^2, σ_s^2 and σ_i^2 (yr^{-2}) be the variances of the real 2014–2024 growth rates of SPY, the sector factor and the name, held fixed rather than measured on each path. The factor share of the name's variance, averaged over the drift in $\beta_i(t)$, is

$$\bar{\rho}_i = \frac{\left(\hat{\beta}_i^2 + k_i^2\,a/(2-a)\right)\sigma_m^2 + \delta_i^2\,\sigma_s^2}{\sigma_i^2}, \tag{14}$$

where the term in parentheses is the long-run mean of $\beta_i(t)^2$. With the JumpHMM paper's floor $s_{\min}^2 = 0.10$ on the own draw's share, we set $s_i^2 = 1 - \bar{\rho}_i$ when $\bar{\rho}_i \le 1 - s_{\min}^2$. Otherwise we multiplied $\beta_i(t)$ and δ_i by $\theta_i = \sqrt{(1-s_{\min}^2)/\bar{\rho}_i}$ for the whole path, which keeps their signs, and set $s_i^2 = s_{\min}^2$. This happened for two Utilities names ($\bar{\rho}_i = 0.917$ and 0.918), or 0.54% of name-days on every path, and the paper's branch that keeps the real R^2 for names with $R^2 \ge 0.80$ was never used (largest $R^2 = 0.65$). Because the shares are real and each JumpHMM draws slightly more variance than its training series, the composed variance sat at the level of the name's own draw, with median ratios of 0.994 to the draw's variance and 1.066 to the real variance (Table S6).

We did not reorder the residuals with a copula. The rank reorder that the JumpHMM paper offers as an option imposes a fitted cross-asset dependence [56, 57], but it permutes each residual path over days and so destroys the name's volatility clustering and the timing of its jump episodes. Each name here keeps its own draw in the order it was produced, and cross-sectional dependence enters only through the shared terms of Eq. (12). Names in one sector share $g_m(t)$ and $f_s(t)$, names in different sectors share only $g_m(t)$, and all dependence is contemporaneous. The mean residual correlation across sectors was therefore 0.0002 on synthetic paths against 0.031 in real 2014–2024, and correlations between names rise in a sell-off only through the market term, whereas real equity correlations rise further in down markets [58, 59].

The composition extends the variance-corrected single-index composition of the JumpHMM paper [19] and departs from it in six ways. First, the paper centered each draw on its realized mean, which forces the residuals to sum to zero over the path. On a path that spans both windows, that sum ties the noise of the trading window to minus the noise of the calibration window, and on 20 test paths the noise parts of each name's intercepts in the two windows correlated -0.48 across names under path centering against 0.003 under centering on the model mean. Second, the paper's variance target was each draw's own sample variance, whereas the variances in Eq. (14) are fixed, so a market jump burst raises a name's variance through its market term instead of shrinking its own noise and loadings over the whole path, calibration years included. Third, the paper composed with the market alone, and Eq. (12) adds the sector term. Fourth, the paper's loading was constant, whereas $\beta_i(t)$ drifts here and the variance is matched on average over the drift with one s_i per path. Matching it day by day, as an earlier version did, scaled down 15 to 24% of name-days in Energy, Utilities, Financials and Real Estate, weakened the sector co-movement, and scaled 8.8% of all name-days rather than 0.54%. Fifth, when the floor binds, the paper reduced the market loading alone, whereas here the market and sector loadings shrink by one factor θ_i. Sixth, the paper kept the fitted intercept, and here it is pulled toward the universe mean. The paper also compared its centered and rescaled full-growth draws with comparators fitted to regression residuals. Our name terms follow the full-growth approach, whereas our sector terms are residual-fit models, since each sector JumpHMM is fitted to a mean of regression residuals.

Before running the method on any synthetic path, we compared the generator with the real market on 22 gated checks of SPY's growth, the gain from daily rebalancing (a guard against a market that rewards trading more than the real one), each name's variance, market loading and clustering, the co-movement within sectors, the spread of $\hat{\beta}_i(t)$, reproducibility from seeds and the OLS intercepts (Table S6). Twenty-one checks passed. The cross-name standard deviation of OLS intercepts failed, at 0.089 yr^{-1} against 0.076 yr^{-1} with a tolerance of 0.01 yr^{-1}. Estimation noise in synthetic intercepts adds spread, and a smaller ϕ would push the correlation between the halves of 2014–2024 further below its target, so we reported the failure rather than tune the generator to pass. Two reported statistics limit what the synthetic market can show. The composed names clustered as their own draws did but less than real names, with absolute-growth autocorrelations below the real ones by a median of 0.10, 0.17 and 0.10 at lags 1, 5 and 20. A median of 7 names per path had $\hat{\beta}_i(t) \leq 0$ at some point in the trading window, near the median of 5 over real windows of that length in 2014–2024 but against 120 in real 2025–2026, so the hedge had fewer candidates than in the backtest. The study's test suite, including its look-ahead tests, passed with the synthetic code in place.

3.6 Evaluation on synthetic futures

We drew $10{,}000$ market paths over the full calendar, seeding path p with p, and ranked them by the maximum drawdown of synthetic SPY within the trading window, 2025-01-02 to 2026-09-04, with the peak reset on the window's first day (Algorithm 4). We tagged a window as containing a jump burst when the market model's flag $J(t)$ was set on any of its growth days, which happened in 523 of the $10{,}000$ windows (5.2%), as the market model's jump settings imply (Appendix S3), and in 30 of the 50 deepest drawdowns. The typical set holds 600 paths drawn at random from all $10{,}000$, the burst-stress set holds the 200 deepest drawdowns among windows with a burst, and the bear-stress set holds the 200 deepest among windows without one. Since 22 typical paths also fell in a stress set, the sets hold 978 distinct paths. Synthetic SPY's median window drawdown was 0.176, 0.364 and 0.406 in the typical, burst-stress and bear-stress sets, and its median growth rate over the window was 0.102, -0.087 and -0.199 yr^{-1}, so a burst window holds months of large moves of both signs with less net decline than the sustained falls of the bear-stress set. Because the sector, drift and name draws are

Algorithm 3 One synthetic future of the 373-name market over the full calendar, 2014-01-03 to 2026-09-04.

Require: Path number p; JumpHMMs of the market, of each sector $s = 1, \dots, 11$ and of each name $i = 1, \dots, N$; real 2014–2024 estimates $\hat{\alpha}_i$, $\hat{\beta}_i$, R_i^2 and loadings δ_i; real variances σ_m^2, σ_s^2, σ_i^2; long-run means $\bar{\varepsilon}_i$; pull ϕ; drift half-life h_β and scales k_i; floor $s_{\min}^2$; real closes $S_{\mathrm{SPY}}(0)$, $S_i(0)$ on 2014-01-03.

Ensure: Closes $S_{\mathrm{SPY}}(t)$ and $S_i(t)$ and jump flags $J(t)$ for $t = 1, \dots, 3{,}186$.

$a \leftarrow 1 - 2^{-1/h_\beta}$; $\quad \bar{\alpha} \leftarrow N^{-1} \sum_i \hat{\alpha}_i$

Draw $g_m(t)$ and $J(t)$ from the market JumpHMM with seed p {JumpHMM simulation [19]}

for $s = 1, \dots, 11$ **do**

  Draw $f_s(t)$ from the JumpHMM of sector s with a seed set by (p, s)

end for

for $i = 1, \dots, N$ **do**

  Draw $\varepsilon_i(t)$ from the JumpHMM of name i with a seed set by (p, i); $\quad$ draw $z_i(t) \sim \mathcal{N}(0, 1)$

  $\alpha_i \leftarrow \bar{\alpha} + \phi(\hat{\alpha}_i - \bar{\alpha})$

  $x_i(1) \leftarrow k_i \sqrt{a/(2-a)}\, z_i(1)$; $\quad x_i(t) \leftarrow (1-a)\, x_i(t-1) + a\, k_i\, z_i(t)$ for $t \geq 2$; $\quad \beta_i(t) \leftarrow \hat{\beta}_i + x_i(t)$ {Eq. (13)}

  $\bar{\rho}_i \leftarrow \left[(\hat{\beta}_i^2 + k_i^2 a/(2-a))\, \sigma_m^2 + \delta_i^2\, \sigma_{s(i)}^2 \right] / \sigma_i^2$ {Eq. (14)}

  if $R_i^2 \geq 0.80$ **then**

    $\theta_i \leftarrow 1$; $\quad s_i^2 \leftarrow \bar{\rho}_i (1 - R_i^2)/R_i^2$ {keeps the real R^2; no name qualified}

  else if $\bar{\rho}_i \leq 1 - s_{\min}^2$ **then**

    $\theta_i \leftarrow 1$; $\quad s_i^2 \leftarrow 1 - \bar{\rho}_i$

  else

    $\theta_i \leftarrow \sqrt{(1 - s_{\min}^2)/\bar{\rho}_i}$; $\quad s_i^2 \leftarrow s_{\min}^2$ {one scale for the whole path}

  end if

  $g_i(t) \leftarrow \alpha_i + \theta_i \beta_i(t)\, g_m(t) + \theta_i \delta_i\, f_{s(i)}(t) + s_i \big(\varepsilon_i(t) - \bar{\varepsilon}_i \big)$ for every t {Eq. (12)}

end for

$S_{\mathrm{SPY}}(t) \leftarrow S_{\mathrm{SPY}}(0) \exp\left(\Delta t \sum_{u \leq t} g_m(u) \right)$; $\quad S_i(t) \leftarrow S_i(0) \exp\left(\Delta t \sum_{u \leq t} g_i(u) \right)$

return S_{SPY}, S_i for every name, and J

independent of the market draw, we drew them only for the chosen paths, from seeds derived from the path number, and any path can be rebuilt bit for bit.

On each of the 978 paths, we treated the synthetic history as if it were real and ran the whole method on it. We first recalibrated every estimated input on synthetic 2014-01-03 to 2024-12-31 with the code of the real run, namely the OLS single-index fit of each name to synthetic SPY, SPY's growth moments and the volatility cap σ_{SPY}, the gain G of the market state ξ, and the shrinkage of the residual correlation toward its diagonal, which the calibration chose anew on every path (zero on 664 of the 978 paths, at most 0.2). The method thus estimated α_i and $\beta_i(t)$ from synthetic prices without seeing the generator's values. We then traded the trading window with the controller and settings of the real backtest. The controller still forecast the portfolio value with a multi-asset geometric Brownian motion (GBM), now calibrated on the synthetic history, and the plant it acted on was the synthetic closes, so JumpHMM played the market and not the forecast model. The synthetic data kept the real calendar, sectors and 3-month Treasury bill rate, held each name's daily share volume at its real 2014–2024 median, so that the dollar volume ranked by the liquidity-based rule moves with the synthetic price, and set opens equal to closes. Orders filled at the next open, so a synthetic order filled at the next day's close, a full trading day after the decision rather than overnight as in the real run, and its fill absorbed the whole day's move.

Each path received 532 runs, a subset of the real grid sized to the compute budget. All seven weighting rules, with the hedge off and on, were run under the three per-sector selection rules (γ, bandit and random, each picking three names in each of the 11 sectors), under five universe-wide rules at $K = 13$ (global γ, correlation-aware γ, liquid then γ, random from the whole universe and random from the 100 most liquid names), under global γ and random selection from the whole universe at $K = 5$ and $K = 34$, and on the two lists of 13 and 15 names that the author chose in advance. The randomized rules, bandit and random, used 5 seeds (5660 to 5664) per path instead of the 30 of the real backtest, and the other rules ran once. Each path also recorded synthetic SPY bought and held and the equal-weight universe bought once and held. The 978 paths gave 520,296 runs, all of which finished without error. On each of the 400 stress paths, the 364 runs of the ten main selection rules were repeated with the band exit and the drawdown stop switched off, as in the calendar-only grid on real prices, and each was paired with the triggered run of the same selection

rule, weighting rule, hedge setting and seed on the same path (145,600 runs, all of which finished without error).

For each set and hedge setting, we summarized the terminal wealth ratio W_T/W_0, the realized growth rate, the maximum drawdown and the Sharpe ratio in selection rule by weighting rule grids with the two benchmark rows, never pooling across selection rules (Fig. 6). We took the median over seeds within a path and then, across paths, the median, the 10th and 90th percentiles, and the share of paths on which the cell's terminal wealth ratio exceeded that of the path's own synthetic SPY. That share is a paired comparison, so a cell can beat SPY on most paths while its median lies below that of the SPY row. We also formed paired differences across paths between global γ and random selection at $K = 5$, 13 and 34, between γ per sector and random per sector, and between hedge on and hedge off, and we located each real backtest result among the 600 typical futures as the share of typical paths that ended at or below it (Fig. S1). These summaries describe this generator combined with this method and are not forecasts.

Algorithm 4 Evaluation of the method on synthetic futures.

Require: Generator of Algorithm 3; $P = 10{,}000$ market paths; calibration window $\mathcal{C}$ (2014-01-03 to 2024-12-31); trading window $\mathcal{W}$ (2025-01-02 to 2026-09-04); run grid $\mathcal{R}$ of 532 runs; seeds 5660 to 5664 for the randomized selection rules.

Ensure: Selection rule by weighting rule summaries for the typical, burst-stress and bear-stress sets.

for $p = 1, \ldots, P$ **do**
  Draw $g_m(t)$ and $J(t)$ with seed p and compound synthetic SPY closes
  $d_p \leftarrow$ maximum drawdown of synthetic SPY over $\mathcal{W}$, peak reset at the start of $\mathcal{W}$
  $b_p \leftarrow$ true if $J(t)$ is set on any growth day of $\mathcal{W}$
end for
Burst stress $\leftarrow$ the 200 paths with the largest d_p among those with b_p true
Bear stress $\leftarrow$ the 200 paths with the largest d_p among those with b_p false
Typical $\leftarrow$ the first 600 paths of a seeded random permutation of $1, \ldots, P$
for each path p in the union of the three sets **do**
  Build the synthetic market of path p {Algorithm 3}
  Recalibrate on synthetic $\mathcal{C}$: single-index fits, SPY moments and volatility cap, gain G of ξ, residual shrinkage
  for each run in $\mathcal{R}$ (selection rule, weighting rule, hedge setting, seed) **do**
    Trade $\mathcal{W}$ with the rebalancing controller, forecasting with GBM calibrated on synthetic data and acting on the synthetic closes
    Record W_T/W_0, growth rate, maximum drawdown and Sharpe ratio
  end for
  Record synthetic SPY and the equal-weight universe, each bought and held over $\mathcal{W}$
end for
For each set, hedge setting and cell: median over seeds within a path, then the median, the 10th and 90th percentiles, and the share of paths above the path's own synthetic SPY

4 Results

We ran the same allocation and rebalancing system on real prices and on synthetic futures and compared the median terminal wealth of every cell with two passive benchmarks (Tables 2 and 3). A cell was one selection rule, one weighting rule, and one hedge setting. On real prices, the ten selection rules (γ per sector, bandit per sector, random per sector, global γ, correlation-aware γ, liquid then γ, random liquid 100, random whole universe, and the two hand-picked lists) and the seven weighting rules (CD, CES, CESw, SG, MinVar, Tangent, and LS) gave 70 cells for each hedge setting, traded over the 420 trading days from 2 January 2025 to 4 September 2026. Each of the 140 cells held 30 runs with seeds 5660–5689, for 4,200 real runs, and the seed set the controller's forecast paths and, for the random and bandit rules, the names picked. A size comparison set global γ against random picks from the whole universe at $K = 5$, 8, 13, 21, 34, 100, and 200 names, again with 30 runs per cell. The synthetic futures were price paths from the JumpHMM generator. We used 600 typical paths drawn at random from 10,000 and two stress sets of 200 paths each, the deepest SPY drawdowns in the trading window among paths with a jump burst (burst stress) and among paths without one (bear stress), and because 22 typical paths were also stress paths, the three sets held 978 distinct paths. On each path the controller recalibrated on the synthetic 2014–2024 history and traded the synthetic 2025–2026 window, and it still forecast with the multi-asset GBM while JumpHMM supplied the prices. Each path carried 532 runs, from 14 selection rules (the ten above, plus global γ and random whole universe at 5 and 34 names), the seven weighting rules, and both hedge settings, with five seeds (5660–5664) for the six random and bandit rules, for 520,296 runs in all, none of which failed. The benchmarks were SPY buy-and-hold and the equal-weight universe, which put $W_0/373$ into each name at the first close and never rebalanced. We measured performance by the terminal wealth ratio W_T/W_0 and by continuously compounded growth rates in yr^{-1}. Every reported median of W_T/W_0 belongs to one cell and is never pooled across selection rules, and on synthetic paths the random and bandit rules first took the median over their five seeds within each path.

We compared the seven weighting rules within each selection rule on real prices with the hedge off, and found that the selection rule moved terminal wealth more than the weighting rule did (Fig. 3; Table 2). Within the γ-per-sector basket, the seven weighting rules ended between 1.086 (MinVar) and 1.234 (SG) times initial wealth, and the four utilities lay between 1.219 and 1.234, whereas across all 70 cells the medians ran from 1.031 (random whole universe, with LS) to 1.697 (hand-picked

13, with Tangent). The spread within a selection rule, the highest minus the lowest of its seven medians, was 0.078 to 0.284, with a median of 0.14, and the spread across the ten selection rules within a weighting rule was 0.357 (MinVar) to 0.591 (LS), with a median of 0.518. Without the two hand-picked lists, the spread across the other eight selection rules was 0.246 to 0.422 (median 0.274), against 0.078 to 0.272 (median 0.135) within them. SPY buy-and-hold ended at 1.317 and the equal-weight universe at 1.296. Of the 70 cells, 29 ended above SPY, all from five selection rules: global γ (six of the seven weighting rules), correlation-aware γ (five), liquid then γ (four), and the two hand-picked lists (seven each). None of the three per-sector rules and neither random 13-name rule ended above SPY with any weighting rule. In the 13-name baskets, global γ ended at 1.248–1.520 against 1.031–1.160 for random names from the whole universe, whereas at three names per sector the γ rule (1.086–1.234) and the bandit (1.140–1.239) ended close to random picks (1.127–1.205). The universe-wide rules were designed after the per-sector results were known, so the 13-name comparison is exploratory. The hand-picked lists ended highest, at 1.443–1.697 for the 13 names and 1.362–1.646 with WMT and PG added. Each is one list chosen by one person, so it shows what that choice earned in this window and not the result of a repeatable rule. MinVar or LS gave the lowest median under every selection rule, MinVar under four rules and LS under six. For the deterministic rules the seed changed only the controller's forecast paths, so their 30 runs clustered tightly (hand-picked 13 with CD: interquartile range 1.546–1.553), whereas the runs of random picks spread over a wider band (random whole universe with CD: 1.048–1.220).

We repeated every real run with the beta hedge on and compared terminal wealth, portfolio beta, and drawdown with the hedge-off runs of the same selection and weighting rules, and found that the hedge lowered market exposure but cost terminal wealth in most cells (Fig. 3; Table S7). At each re-plan with the hedge on, the allocation obeyed a cap on portfolio β that fell as the market state ξ rose, and up to five names whose online β and down-day β were at or below zero could enter as a hedge pool. The median W_T/W_0 was lower with the hedge on in 50 of the 70 cells, and the median of the 70 changes was -0.033. Wealth fell in six or seven of the seven cells for the bandit rule, the three random rules, and liquid then γ, and in five and six for the two hand-picked lists. Correlation-aware γ ended higher in five of its seven cells, and LS weights ended higher under six of the ten selection rules. We measured exposure by the portfolio β, the value-weighted online β of the holdings divided by total wealth, averaged over the days of a run. The hedge lowered this mean portfolio β in all 70 cells (cell medians 0.388–0.903 with the hedge off and 0.314–0.849 with it on), and hedge-pool names held 8.4–39.8% of wealth on an average day (cell medians). The maximum drawdown nonetheless changed little, by at most 0.019 in either direction, with a median change of $+0.001$ across the 70 cells (Table S8). When the hedge-on and hedge-off runs were paired by seed, the hedge-on run had the smaller maximum drawdown in no more than 14 of 30 pairs for any weighting rule under the bandit and random per-sector rules. The cap explained part of the wealth loss. The state ξ was positive on 75 of the 420 trading days (Fig. 1), and the portfolio's mean growth rate on the 75 days that followed a positive ξ fell with the hedge on in 63 of the 70 cells (median change $-0.153\,\mathrm{yr}^{-1}$). On those days SPY's mean daily growth rate was $0.266\,\mathrm{yr}^{-1}$, against $0.144\,\mathrm{yr}^{-1}$ on the other 344 of the 419 growth days, so a tighter cap held less of a market that rose on average. The difference of $0.12\,\mathrm{yr}^{-1}$ was small next to its standard error of $0.59\,\mathrm{yr}^{-1}$, so it describes this window and not a property of the market state. The real window was one historical path, with a largest SPY drawdown of 0.190, so we also tested the hedge on simulated sell-offs.

We followed the controller through one real episode and counted re-plans by trigger across the 4,200 main-comparison runs, and found that the forecast band and the drawdown stop brought forward a share of re-plans that depended on the selection rule (Fig. 2; Table S9). The episode came from the hand-picked 13 with CD weights, the hedge off, and seed 5660, and ran from a re-plan on 10 October 2025 to 22 December 2025. At each re-plan the controller forecast portfolio value over the next 21 trading days with the multi-asset GBM and 1,000 paths, while the portfolio traded real prices. At the horizon re-plan on 10 November the portfolio stood at 1.101 times initial wealth, 0.61 standard deviations (SD) above the mean of the forecast made on 10 October. Eight trading days into the next forecast, on 20 November, the portfolio had fallen to 1.0489 against a forecast mean of 1.1235 with an SD of 0.0355, 2.10 SD below the mean and outside the 1.96 SD band, and the controller re-planned. The third forecast ran to its horizon on 22 December, when the portfolio stood at 1.148, 1.09 SD above its mean. In the cell of this run, the median share of re-plans that came before the horizon was 0.37. Across the 140 cells, the median share brought forward by a band exit or a drawdown stop ranged from 0.20 to 0.61 (Table S9). With the hedge off, the share was lowest for the three random rules, at 0.20 to 0.40, and highest for the four γ rules, whose highest cells ran

from 0.53 to 0.61. The other re-plans came at the horizon or on a basket change, a day on which a timer pick changed the set of names, and no run re-planned on a cash revisit.

We repeated the main comparison with the band exit and the drawdown stop switched off, and found that the early re-plans lowered the worst drawdown more consistently than they changed terminal wealth (Table S10). Each of the 4,200 calendar-only runs kept the selection rule, weighting rule, hedge setting and seed of one triggered run, so the two runs of a pair differed only in the two triggers, and every calendar-only run re-planned at the 21-day horizon alone, 19 times over the window, whereas in every cell the triggered runs re-planned more often in the median. We ran this comparison because band exits measure how often the forecast missed, not whether re-planning early helped the portfolio. With the hedge off, the median maximum drawdown was smaller with the triggers in 55 of the 70 cells, by up to 0.053 (global γ), and at most 0.015 larger in the others. With the hedge on it was smaller in 61 of the 70 cells, by up to 0.054. Median terminal wealth was higher with the triggers in 38 of the 70 cells with the hedge off and in 26 with the hedge on, and the differences ran from -0.093 to 0.191 of initial wealth with the hedge off and from -0.180 to 0.183 with it on. The sign of the wealth difference did not follow the kind of selection rule: with the hedge off, the triggers raised terminal wealth under all seven weighting rules for random liquid 100 and under none for random whole universe or liquid then γ. Every triggered cell paid more trading cost than its calendar-only twin, 3 to 33% more with the hedge off and 4 to 45% more with it on (Table S9), since the triggered runs re-planned more often. On the two stress sets of synthetic futures, the triggers likewise lowered the worst drawdown in most cells without consistently raising terminal wealth (Tables S11 and S12).

We compared the K highest-γ names in the universe with K names drawn at random from the whole universe for K from 5 to 200 on real prices, and found that the advantage of selection by γ shrank as the basket grew (Fig. S2). For each basket size, weighting rule, and hedge setting, we formed all 900 pairs of one of the 30 global γ runs with one of the 30 random runs, regardless of seed, and counted the share of pairs in which global γ ended with the higher W_T/W_0. Like the 13-name comparison, this size comparison is exploratory. Stone–Geary was left out at 100 and 200 names, because its baseline holding of 1% of wealth per name used the whole budget at 100 names and twice the budget at 200, so most of its decisions were infeasible. With the hedge off, global γ ended higher in 91.8–100% of pairs at 5 names, 94.7–100% at 8, and 90.2–100% at 13 across the weighting rules. The share fell to 74.9–96.0% at 21 names and 56.6–96.0% at 34, and to 33.3–83.2% at 100 names and 34.6–82.7% at 200. Random picks ended higher in most pairs for MinVar at 100 names and for CD, CES, and Tangent at 200. With CD weights, the global γ median fell from 1.568 at $K = 5$ to 1.520 at $K = 13$, 1.333 at $K = 34$, 1.244 at $K = 100$, and 1.185 at $K = 200$, while the random median rose from 1.048 to 1.098, 1.182, 1.203, and 1.211 at the same sizes. With the hedge on, global γ ended higher in 85.9–100% of pairs at 5 to 13 names, and in 90.3–100% at 21 and 34 names for every weighting rule except MinVar, whose share fell to 43.2% at 21 names and 26.1% at 34. At 100 and 200 names with the hedge on, the shares were 48.4–85.2% and 30.7–81.6%. Selection by γ thus mattered most for baskets of 5 to 13 names, and at 100 or more of the 373 names the CD medians of the two rules differed by less than 0.05.

We compared daily growth rates of the JumpHMM market generator with those of real SPY over 2014–2024 and checked the composed multi-asset paths against the real universe, and found that the generator reproduced fat tails but not the persistence of volatility (Fig. 4; Table S6). For the comparison with SPY we used the market factor of the first 100 typical paths by path number. Across these paths the median kurtosis of the daily growth rate was 13.7, against 16.8 for SPY and 3 for a normal distribution. In the tails the median generator quantile lay at or beyond SPY's, at -17.2 against $-16.3\,\mathrm{yr}^{-1}$ for the 0.1% quantile, -10.1 against -7.8 for the 1% quantile, 7.3 against 6.4 for the 99% quantile, and 15.1 against 14.8 for the 99.9% quantile. Both series had small autocorrelation of g at lag 1 (-0.04 for the generator median, -0.11 for SPY). The autocorrelation of $|g|$ differed, since for the generator median it fell from 0.22 at lag 1 to 0.08 at lag 2, 0.03 at lag 3, and 0.007 at lag 21, whereas SPY's was 0.36 at lag 1 and still 0.13 at lag 21, about one trading month. Of the 22 gated checks on the composed paths, each with a pass tolerance, 21 passed. The median synthetic SPY growth rate over 2014–2024 was $0.105\,\mathrm{yr}^{-1}$ across the 10,000 paths, against $0.106\,\mathrm{yr}^{-1}$ for real SPY. The median name's variance was 0.994 of the variance of its own JumpHMM draw, and 1.066 of its real variance, a comparison reported without a tolerance. The mean residual correlation within each of the 11 sectors lay within 0.05 of the real value (for example 0.535 against 0.542 in Energy), and the SD of the online β was 0.985 of the real one. The failing check was the cross-name SD of α, the intercept of each name's least-squares regression on the market over 2014–2024, which was

$0.089\,\mathrm{yr}^{-1}$ against $0.076\,\mathrm{yr}^{-1}$ for the real universe, outside the tolerance of $0.01\,\mathrm{yr}^{-1}$. Comparisons without a tolerance showed what the composition left out. The mean residual correlation across sectors was 0.0002 against 0.031, because cross-sectional dependence entered only through the market and sector factors. The median name's autocorrelation of $|g|$ was 0.10 below the real one at lag 1 and 0.17 below at lag 5, but within 0.04 of its own JumpHMM draw, so the shortfall came from the per-name draws and not from the composition. The synthetic paths also offered the hedge fewer candidates, with a median of seven names per path whose online β reached zero or below in the 2025–2026 window, against 120 in the real 2025–2026 window and a median of five in real 2014–2024 windows of the same length (range 0 to 53). The synthetic paths thus resembled a typical stretch of 2014–2024, and the real trading window held more hedge candidates than any earlier stretch of that length.

We ran all selection and weighting rules on the 600 typical futures, compared selection by γ with random picks of the same size, and placed each real result among the typical futures of its cell, and found that selection by γ beat random picks in most typical futures (Tables 3 and S13; Fig. S1). In the typical futures the median synthetic SPY ended at 1.185 and the equal-weight universe at 1.230, and the 70 hedge-off cell medians of the ten main selection rules ran from 1.097 (random whole universe, with LS) to 1.244 (liquid then γ, with SG). Of these 70 cells, 45 ended above the SPY median and seven above the equal-weight median. The equal-weight universe ended above its own path's SPY in 77.5% of typical futures, whereas on real prices it ended below SPY (1.296 against 1.317). With the hedge off, global γ ended above 13 random names from the whole universe in 62.3–71.2% of futures across the weighting rules, with a median difference in W_T/W_0 of 0.042 to 0.101, and γ per sector ended above random per sector in 60.0–73.7%. With the hedge on, the 13-name share was 63.2–71.5%. The hand-picked 13 ended between 1.175 and 1.242, within the 1.142–1.244 range of the four γ rules (γ per sector, global γ, correlation-aware γ, and liquid then γ), whereas on real prices both hand-picked lists ended highest. To place the real results, we used the seeds of the synthetic runs, 5660 for the deterministic rules and the median over 5660–5664 for the random and bandit rules, rather than the 30-seed medians, and took the percentile of each real value among the 600 typical futures of its cell. Real SPY sat at percentile 70 of synthetic SPY and the real equal-weight universe at 60.7, so the real window was a better-than-median period for the market, and part of any real margin over the typical futures reflected that period. The real results of the hand-picked 13 sat at percentiles 79.2 to 90.7 of their cells, those of global γ at 47.5 to 81.7, those of γ per sector at 39.8 to 56.2, and those of 13 random names from the whole universe at 43.3 to 64.3. Of the 70 hedge-off cells, 52 sat above the median future and 19 above SPY's percentile of 70. The hand-picked lists sat at or above percentile 74 in all 14 of their hedge-off cells, above the market's 70, whereas the four γ rules sat at percentiles 39.8 to 81.7 and above 70 in five of their 28 cells (Fig. S1). In the real window the hand-picked lists thus exceeded their typical futures by more than the market did, and the γ rules in most cells did not.

We repeated the comparison on the two stress sets of simulated sell-offs, and found that every selection and weighting median ended above SPY and the equal-weight universe (Fig. 6; Table 3). The median synthetic SPY drawdown in the 2025–2026 window was 0.364 in the burst stress set and 0.406 in the bear stress set, against 0.176 in the typical futures, and SPY's median growth rate was -0.087 and $-0.199\,\mathrm{yr}^{-1}$, against $0.102\,\mathrm{yr}^{-1}$. In both stress sets all 140 cell medians of the ten main selection rules ended above the SPY median and the equal-weight median, and so did all 98 cells per hedge setting of the 14 synthetic selection rules. Among the ten main selection rules, the lowest median in the burst stress set was 0.924 (random per sector, MinVar, hedge off), against 0.865 for SPY and 0.916 for the equal-weight universe, and the highest was 1.031 (hand-picked 13, LS, hedge off). In the bear stress set the lowest median was 0.774 (random whole universe, Tangent, hedge off), against 0.718 for SPY and 0.753 for the equal-weight universe, and the highest was 1.001 (hand-picked 13 with WMT and PG, LS, with either hedge setting). Each cell also ended above its own path's SPY in 69.0–88.5% of burst futures and 72.5–100% of bear futures. Random picks also ended above SPY, so the margin over SPY did not come from selection by γ. The equal-weight universe, which neither caps volatility nor re-plans, also ended above SPY, by 0.051 in the burst stress set and 0.035 in the bear stress set, and the lowest cell medians exceeded it by 0.009 and 0.021, so part of the margin over SPY came from the synthetic market and the rest from parts of the system that every cell shared. With the hedge off, the cell medians of the mean portfolio β were 0.322 to 0.764 in the burst stress set and 0.229 to 0.850 in the bear stress set, all below one (Table S14). In bear futures selection by γ lost its typical-future advantage for five of the seven weighting rules. With the hedge off, global γ ended above 13 random names in 28.0–38.0% of bear futures for CD, CES, CESw, SG, and LS, and

at 5 names the same five weighting rules gave 38.5–45.0%. MinVar and Tangent, whose weights do not use γ, reversed this pattern, with global γ ahead in 70.0% and 71.0% of bear futures. In burst futures global γ ended above 13 random names in 52.5–67.5% of futures with the hedge off. With the hedge on, the median W_T/W_0 ended higher in 65 of the 70 cells of the bear stress set (median of the 70 cell changes $+0.029$), and the lowest cell median rose from 0.774 to 0.814. In the burst stress set it ended higher in 34 of the 70 cells (median change -0.001), and in the typical futures in 11 of the 70 (median change -0.006). The β cap and hedge pool thus raised terminal wealth in prolonged declines and lowered it in typical futures, as they did on real prices.

We repeated the calendar-only comparison on the 400 stress futures, and found that the early re-plans again lowered the worst drawdown more consistently than they changed terminal wealth, by about 0.01 in the median cell as on real prices, although the stress declines were deeper (Tables S11 and S12). The median maximum drawdown was smaller with the triggers in 65 of the 70 cells with the hedge off and in all 70 with it on in the burst stress set, and in 65 and 67 of the 70 in the bear stress set, by up to 0.022, and at most 0.007 larger in the others. Median terminal wealth was higher with the triggers in 33 and 28 of the 70 cells in the burst stress set and in 51 and 56 in the bear stress set, with differences from -0.031 to 0.029 of initial wealth. Every calendar-only run re-planned 19 times, every 21 trading days, while the triggered cells re-planned 25 to 32 times in the median. With the two triggers switched off, every cell median still ended above the SPY median and the equal-weight median in both stress sets, the lowest at 0.930 in the burst stress set and 0.772 in the bear stress set, so the margin over SPY did not depend on the early re-plans.

5 Discussion

The clearest result of the study concerned allocation. On real prices without the market hedge (when on, the hedge capped the portfolio's beta at each re-plan, at a level that was lower when the market was weaker, and let up to five stocks with zero or negative beta enter), the seven weighting rules applied to one selection of 33 stocks (three in each of the 11 sectors) ended within 0.15 of one another, and all seven ended below SPY. The selections themselves ended further apart, from 1.03 to 1.70 times initial wealth across all combinations, and the wide spread came from the 13-stock selections, made across the whole universe of 373 stocks without a sector quota, and from the two lists the author picked in advance (13 stocks, and the same 13 plus two). One explanation, consistent with these runs but not tested by them, is that the shared limits left the weighting rules little room. Every rule, including the three benchmark portfolios (minimum variance, tangency and 130/30 long-short), faced the same cap on each stock's weight and the same cap on portfolio volatility, so the rules differed mainly in how they spread the invested wealth among the selected stocks within those caps. The results do not show that any one selection rule is reliably better on real data. The hand-picked 13-stock list ended highest, but it is one list chosen by one person, and its result cannot be separated from chance. The three score-based 13-stock rules, which took the 13 highest scores (the expected growth rate from a single-index fit divided by a power of the stock's market beta that depended on the market's recent direction) with no screen, a correlation screen or a liquidity screen, ended above the two random 13-stock controls, drawn from the whole universe or from its 100 most liquid stocks, under every weighting rule, while the three 33-stock rules (ranked by score, chosen by a bandit, or drawn at random within each sector) ended in overlapping ranges. The 13-stock rules were added after the backtests of the 33-stock rules on the same trading window had been seen, so their comparison with random selections is exploratory, and one historical path cannot show whether the difference would persist [10, 11].

The rebalancing controller re-planned, re-solving the allocation from the current holdings, on a schedule set partly by its own forecast errors. Across combinations of selection and weighting rule, the median of each combination's 30 runs brought a fifth to three fifths of its re-plans forward of the 21-day horizon, through a band exit (the realized portfolio value leaving the range the forecast considered likely) or a drawdown stop (wealth falling 8% below its peak since the last drawdown stop), and early re-plans were more frequent for stocks selected by score than for random stocks, a difference the study did not examine. The forecast, a geometric Brownian motion, represents neither the heavy tails nor the volatility clustering of daily growth rates [16, 17], so band exits are expected to occur more often than they would if the forecast model were correct, and the band served both as a trigger and as a running check on the forecast. A forecast model that represents volatility clustering [33, 34] would widen the band after volatile days and narrow it after calm ones, so its

effect on the number of early re-plans cannot be signed in advance, and this was not tested. Against re-planning every 21 trading days alone, the early re-plans made the median worst drawdown smaller in most combinations of selection and weighting rule, with the hedge on or off, and raised trading cost in all of them. Median terminal wealth, by contrast, rose with the early re-plans in about half of the combinations with the hedge off and in fewer than half with it on. The same pattern held in the two stress sets of synthetic futures, whose declines were deeper than the real one: the early re-plans made the median worst drawdown smaller in nearly every combination, by about as much as on real prices, and raised median terminal wealth in fewer than half of the combinations in the set with jump bursts and in most of those in the set without (Tables S11 and S12). The triggers therefore limited losses more than they added growth, and the protection they gave did not grow with the depth of the decline.

Results on the simulated markets describe how this system behaved in one generator's markets and are not forecasts. In both stress sets, the futures with the deepest declines of simulated SPY, one drawn from futures whose market path contained a jump episode (a sustained burst of large daily moves) and one from futures without, every combination of selection and weighting rule ended above simulated SPY bought and held in the median, with the hedge on or off. Random selections did so as well, so choosing stocks by score or by bandit was not needed for this result, and every combination also did so with the band exit and the drawdown stop switched off, so the early re-plans were not needed either (Tables S11 and S12). The equal-weight universe, which has neither limits nor triggers, also ended above simulated SPY in both stress sets, so part of the margin over SPY belonged to the simulated market itself, and every combination's median exceeded the equal-weight median by as little as 0.009 in the set with jump episodes and 0.021 in the set without. A likely cause of the margin over SPY is that the portfolios carried less market exposure than the index. With the hedge off, the average market beta of the portfolios was below one in every combination, at 0.32 to 0.76 in the stress set with jump episodes and 0.23 to 0.85 in the set without (medians over futures), as on real prices (0.39 to 0.90), and the market term of a portfolio with a beta below one falls less than the market when the market factor falls, whatever stocks it holds. The volatility cap, set at SPY's historical level, is one reason the beta stays below one, but the study did not run the system without it.

The synthetic market also lacked two features of real sell-offs. Each stock's own simulated residual was generated independently of the others, so co-movement among stocks came only from the simulated SPY and the sector factors. Real residual co-movement tends to rise in sell-offs [58, 59], and its absence may understate how much a portfolio of a few stocks loses relative to the index. The generator's volatility clustering also decayed within about a day, faster than in the data, so turbulent stretches in the synthetic futures were shorter than in real markets, which may affect both the depth of drawdowns and how often the triggers fired. The effect of the market hedge also changed sign across markets. The hedge lowered market exposure in every combination on real prices but reduced terminal wealth in most of them, and it reduced terminal wealth in most combinations of the 600 typical futures, which were drawn at random from $10{,}000$ simulated paths. In the stress set without jump episodes it raised terminal wealth in most combinations, and in the stress set with jump episodes it raised terminal wealth in about half of them. Since the hedge's beta cap tightens after the market has started to fall, and takes effect only at the next re-plan, it pays only when the decline continues. Between re-plans the shares were held, so the portfolio's beta was not controlled, and in a slow decline that stayed inside the forecast band and above the drawdown limit the book could keep the shares of its last re-plan for up to 21 trading days.

The comparison of score-based and random selections on synthetic futures needs a further qualification. The generator gave each synthetic stock a constant intercept, its 2014–2024 single-index intercept shrunk toward the average across stocks. In the synthetic market, a stock's intercept on the calibration years was therefore positively correlated with its intercept on the trading window: across stocks, the two correlated at 0.10, against -0.12 in the real market. The score's trend term is a daily updated estimate of this intercept, so in the synthetic futures it carried information about future growth by construction, and the share of typical futures in which score-based selections beat random ones is partly a consequence of that design. The spread of estimated intercepts across stocks was also wider in the synthetic market than in the real one (0.089 against 0.076 per year), and a wider spread of persistent intercepts gives a score-based selection more to find. The intercept spread was the one check the generator failed among the 22 that had a pass criterion, whereas the intercept correlation and the clustering shortfall were reported without one. On real prices, the score-based 13-stock rules still ended above random selections even though intercepts did not persist from 2014–2024

into 2025–2026. The study does not separate whether that came from the score's market term, from shorter-lived trends picked up by the daily updated estimate, or from chance.

The study has further limits. The trading window covered 420 trading days with one broad sell-off, in the spring of 2025. The universe held only stocks that were members of the S&P 500 index at the date of the constituent snapshot and traded on every day from 2014 to 2026, so it carries survivorship bias [60]. Neither the portfolios nor SPY earned dividends, so every comparison is in prices alone. The simulated fills charged half the bid-ask spread and a linear market-impact term, the turnover limit, a cap on the half-spread cost of the solver's share changes at each re-plan, used the half spread alone, and taxes were not modeled. A paper-trading deployment of the system is running, but its results are not part of this study. Natural extensions are a forecast model with heavy tails and volatility clustering for the controller, runs without the volatility cap to measure its share of the stress-set results, residual co-movement across stocks and less persistent intercepts in the synthetic market, and evaluation of selection rules over many historical windows.

6 Conclusion

We built a system that chose stocks, set their share counts by maximizing a utility under limits on weight, volatility and turnover, and re-planned the holdings with a model-predictive controller, and we tested it on 373 United States equities over 420 trading days of 2025–2026 and on simulated markets. On real prices, the choice of stocks moved terminal wealth more than the choice among seven weighting rules: one 33-stock selection ended between 1.09 and 1.23 times initial wealth across the weighting rules, while different selections ended between 1.03 and 1.70, the highest being a 13-stock list the author picked in advance. The share of the controller's re-plans that came before the 21-day horizon, because the portfolio value left the 95% band of its forecast or fell 8% below its peak since the last drawdown stop, ranged from a fifth to three fifths across combinations of selection and weighting rule, highest for stocks selected by their score, the expected growth rate from a single-index fit divided by a power of the stock's market beta that depended on the market's recent direction. Compared with re-planning every 21 trading days alone, these early re-plans made the median worst drawdown smaller in 55 of the 70 combinations of the ten selection rules and seven weighting rules without the market hedge and in 61 with it, and raised trading cost in all of them, but raised median terminal wealth in 38 and 26 of the 70. Among 10,000 simulated price paths (synthetic futures), we kept two sets of 200 with the deepest drawdowns of simulated SPY, one from futures whose SPY went through a burst of extreme daily moves and one from futures without such a burst. In both sets, the median terminal wealth of every combination ended above those of simulated SPY bought and held and of the equal-weight universe. Random selections did so as well, and so did every combination with the band exit and the drawdown stop switched off. The equal-weight universe also ended above simulated SPY, so part of the margin over SPY came from the simulated market and the rest from parts of the system that every combination shared other than those two triggers, such as the volatility cap, whose separate effect the study did not measure. In these sets the early re-plans again made the median worst drawdown smaller, in 65 to 70 of the 70 combinations for each set and hedge setting, without consistently raising median terminal wealth. In 600 futures drawn at random from the 10,000 (typical futures), the 13 stocks with the highest scores ended above 13 random stocks in 62 to 71% of futures across the weighting rules, but the generator kept each stock's 2014–2024 trend constant, and the score's trend term estimates that trend, whereas in the real market those trends did not carry over into 2025–2026, so part of this advantage came from the generator's design and may not carry over to real prices. A market hedge, which at each re-plan capped the portfolio's beta more tightly when the market was weaker and added stocks with zero or negative beta, reduced terminal wealth on real prices and added to it in the set without bursts, whose declines were the most sustained. Deciding which selection rule is better, and when such a hedge pays, will need more historical windows and generators whose trends persist no more than real ones do.

Acknowledgments

The author thanks Abdulrahman Alswaidan for the JumpHMM generator and its multi-asset composition, which supplied the synthetic markets used here; Aaron Wheeler for helpful discussions; Evan Vu for an early version of the Cobb–Douglas allocation and the bandit selection rule; and Alpaca

Markets. Computations were run on a single desktop computer (Apple Mac Studio, M3 Ultra, 32 cores); no shared-resource allocation was required.

Author contributions. J.D.V. designed the study, wrote the code, ran the experiments, and wrote the manuscript.

Competing interests. The author declares no competing interests.

Code and data availability. The study code, run scripts, result summaries, and the scripts that draw every figure from those summaries are provided in the study repository (https://github.com/varnerlab/INFORMS-Annual-Meeting-2026-Varnerlab); the per-asset generator library is released separately as `JumpHMM.jl` (https://github.com/varnerlab/JumpHMM.jl).

Generative artificial intelligence use. During the preparation of this work, the author used Anthropic Claude Code and OpenAI Codex to assist with the implementation and checking of experiment code, figure preparation, drafting and language editing, consistency checks, and LaTeX formatting. The author reviewed and edited the assisted work and takes full responsibility for the content of the publication.

References

[1] Harry Markowitz. Portfolio selection. *The Journal of Finance*, 7(1):77–91, 1952. doi: 10.1111/j.1540-6261.1952.tb01525.x.

[2] Robert C. Merton. Lifetime portfolio selection under uncertainty: The continuous-time case. *The Review of Economics and Statistics*, 51(3):247–257, 1969. doi: 10.2307/1926560.

[3] Richard O. Michaud. The Markowitz optimization enigma: Is 'optimized' optimal? *Financial Analysts Journal*, 45(1):31–42, 1989. doi: 10.2469/faj.v45.n1.31.

[4] Olivier Ledoit and Michael Wolf. Honey, I shrunk the sample covariance matrix. *The Journal of Portfolio Management*, 30(4):110–119, 2004. doi: 10.3905/jpm.2004.110.

[5] Victor DeMiguel, Lorenzo Garlappi, and Raman Uppal. Optimal versus naive diversification: How inefficient is the 1/N portfolio strategy? *The Review of Financial Studies*, 22(5):1915–1953, 2009. doi: 10.1093/rfs/hhm075.

[6] Michael J. P. Magill and George M. Constantinides. Portfolio selection with transactions costs. *Journal of Economic Theory*, 13(2):245–263, 1976. doi: 10.1016/0022-0531(76)90018-1.

[7] M. H. A. Davis and A. R. Norman. Portfolio selection with transaction costs. *Mathematics of Operations Research*, 15(4):676–713, 1990. doi: 10.1287/moor.15.4.676.

[8] Florian Herzog, Gabriel Dondi, and Hans P. Geering. Stochastic model predictive control and portfolio optimization. *International Journal of Theoretical and Applied Finance*, 10(2): 203–233, 2007. doi: 10.1142/S0219024907004196.

[9] Stephen Boyd, Enzo Busseti, Steven Diamond, Ronald N. Kahn, Kwangmoo Koh, Peter Nystrup, and Jan Speth. Multi-period trading via convex optimization. *Foundations and Trends in Optimization*, 3(1):1–76, 2017. doi: 10.1561/2400000023.

[10] David H. Bailey, Jonathan M. Borwein, Marcos López de Prado, and Qiji Jim Zhu. Pseudo-mathematics and financial charlatanism: The effects of backtest overfitting on out-of-sample performance. *Notices of the American Mathematical Society*, 61(5):458–471, 2014. doi: 10.1090/noti1105.

[11] Campbell R. Harvey and Yan Liu. Backtesting. *The Journal of Portfolio Management*, 42(1): 13–28, 2015. doi: 10.3905/jpm.2015.42.1.013.

[12] Charles W. Cobb and Paul H. Douglas. A theory of production. *The American Economic Review*, 18(1):139–165, 1928.

[13] Kenneth J. Arrow, Hollis B. Chenery, Bagicha S. Minhas, and Robert M. Solow. Capital-labor substitution and economic efficiency. *The Review of Economics and Statistics*, 43(3):225–250, 1961. doi: 10.2307/1927286.

[14] Richard Stone. Linear expenditure systems and demand analysis: An application to the pattern of British demand. *The Economic Journal*, 64(255):511–527, 1954. doi: 10.2307/2227743.

[15] R. C. Geary. A note on "a constant-utility index of the cost of living". *The Review of Economic Studies*, 18(1):65–66, 1950. doi: 10.2307/2296107.

[16] Benoit Mandelbrot. The variation of certain speculative prices. *The Journal of Business*, 36(4): 394–419, 1963.

[17] R. Cont. Empirical properties of asset returns: Stylized facts and statistical issues. *Quantitative Finance*, 1(2):223–236, 2001. doi: 10.1080/713665670.

[18] William F. Sharpe. A simplified model for portfolio analysis. *Management Science*, 9(2): 277–293, 1963. doi: 10.1287/mnsc.9.2.277.

[19] Abdulrahman Alswaidan and Jeffrey D. Varner. Variance-corrected multi-asset equity simulation with hybrid hidden Markov marginals, 2026. arXiv:2603.10202.

[20] Eugene F. Fama and Kenneth R. French. Common risk factors in the returns on stocks and bonds. *Journal of Financial Economics*, 33(1):3–56, 1993. doi: 10.1016/0304-405X(93)90023-5.

[21] William Brock, Josef Lakonishok, and Blake LeBaron. Simple technical trading rules and the stochastic properties of stock returns. *The Journal of Finance*, 47(5):1731–1764, 1992. doi: 10.1111/j.1540-6261.1992.tb04681.x.

[22] Tobias J. Moskowitz, Yao Hua Ooi, and Lasse Heje Pedersen. Time series momentum. *Journal of Financial Economics*, 104(2):228–250, 2012. doi: 10.1016/j.jfineco.2011.11.003.

[23] Andrea Frazzini and Lasse Heje Pedersen. Betting against beta. *Journal of Financial Economics*, 111(1):1–25, 2014. doi: 10.1016/j.jfineco.2013.10.005.

[24] Bruce I. Jacobs, Kenneth N. Levy, and David Starer. On the optimality of long-short strategies. *Financial Analysts Journal*, 54(2):40–51, 1998. doi: 10.2469/faj.v54.n2.2161.

[25] Tor Lattimore and Csaba Szepesvári. *Bandit Algorithms*. Cambridge University Press, Cambridge, 2020. doi: 10.1017/9781108571401.

[26] Nicolae Gârleanu and Lasse Heje Pedersen. Dynamic trading with predictable returns and transaction costs. *The Journal of Finance*, 68(6):2309–2340, 2013. doi: 10.1111/jofi.12080.

[27] D. Q. Mayne, J. B. Rawlings, C. V. Rao, and P. O. M. Scokaert. Constrained model predictive control: Stability and optimality. *Automatica*, 36(6):789–814, 2000. doi: 10.1016/S0005-1098(99)00214-9.

[28] Giuseppe C. Calafiore. Multi-period portfolio optimization with linear control policies. *Automatica*, 44(10):2463–2473, 2008. doi: 10.1016/j.automatica.2008.02.007.

[29] Sanford J. Grossman and Zhongquan Zhou. Optimal investment strategies for controlling drawdowns. *Mathematical Finance*, 3(3):241–276, 1993. doi: 10.1111/j.1467-9965.1993.tb00044.x.

[30] Samuel A. Assefa, Danial Dervovic, Mahmoud Mahfouz, Robert E. Tillman, Prashant Reddy, and Manuela Veloso. Generating synthetic data in finance: Opportunities, challenges and pitfalls. In *Proceedings of the First ACM International Conference on AI in Finance*, pages 1–8, 2020. doi: 10.1145/3383455.3422554.

[31] James Jordon, Lukasz Szpruch, Florimond Houssiau, Mirko Bottarelli, Giovanni Cherubin, Carsten Maple, Samuel N. Cohen, and Adrian Weller. Synthetic data: What, why and how? *arXiv preprint arXiv:2205.03257*, 2022.

[32] Dimitris N. Politis and Joseph P. Romano. The stationary bootstrap. *Journal of the American Statistical Association*, 89(428):1303–1313, 1994. doi: 10.1080/01621459.1994.10476870.

[33] Robert F. Engle. Autoregressive conditional heteroscedasticity with estimates of the variance of United Kingdom inflation. *Econometrica*, 50(4):987–1007, 1982.

[34] Tim Bollerslev. Generalized autoregressive conditional heteroskedasticity. *Journal of Econometrics*, 31(3):307–327, 1986. doi: 10.1016/0304-4076(86)90063-1.

[35] James D. Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2):357–384, 1989. doi: 10.2307/1912559.

[36] Magnus Wiese, Robert Knobloch, Ralf Korn, and Peter Kretschmer. Quant GANs: Deep generation of financial time series. *Quantitative Finance*, 20(9):1419–1440, 2020. doi: 10.1080/14697688.2020.1730426.

[37] Michael Stenger, André Bauer, Thomas Prantl, Robert Leppich, Nathaniel Hudson, Kyle Chard, Ian Foster, and Samuel Kounev. Thinking in categories: A survey on assessing the quality for time series synthesis. *Journal of Data and Information Quality*, 16(2):1–32, 2024. doi: 10.1145/3666006.

[38] Board of Governors of the Federal Reserve System (US). 3-month treasury bill secondary market rate, discount basis [DTB3]. FRED, Federal Reserve Bank of St. Louis, 2026. URL https://fred.stlouisfed.org/series/DTB3. Daily series; observations through 2026-10-16 were retrieved.

[39] Robert Almgren and Neil Chriss. Optimal execution of portfolio transactions. *Journal of Risk*, 3(2):5–39, 2001.

[40] Narasimhan Jegadeesh and Sheridan Titman. Returns to buying winners and selling losers: Implications for stock market efficiency. *The Journal of Finance*, 48(1):65–91, 1993. doi: 10.1111/j.1540-6261.1993.tb04702.x.

[41] Stephen Boyd and Lieven Vandenberghe. *Convex Optimization*. Cambridge University Press, Cambridge, 2004. doi: 10.1017/CBO9780511804441.

[42] Miles Lubin, Oscar Dowson, Joaquim Dias Garcia, Joey Huchette, Benoît Legat, and Juan Pablo Vielma. JuMP 1.0: Recent improvements to a modeling language for mathematical optimization. *Mathematical Programming Computation*, 15:581–589, 2023. doi: 10.1007/s12532-023-00239-3.

[43] Paul J. Goulart and Yuwen Chen. Clarabel: An interior-point solver for conic programs with quadratic objectives, 2024. arXiv:2405.12762.

[44] Brendan O'Donoghue, Eric Chu, Neal Parikh, and Stephen Boyd. Conic optimization via operator splitting and homogeneous self-dual embedding. *Journal of Optimization Theory and Applications*, 169(3):1042–1068, 2016. doi: 10.1007/s10957-016-0892-3.

[45] James Tobin. Liquidity preference as behavior towards risk. *The Review of Economic Studies*, 25(2):65–86, 1958. doi: 10.2307/2296205.

[46] Abraham Charnes and William W. Cooper. Programming with linear fractional functionals. *Naval Research Logistics Quarterly*, 9(3–4):181–186, 1962. doi: 10.1002/nav.3800090303.

[47] Carlos E. García, David M. Prett, and Manfred Morari. Model predictive control: Theory and practice – a survey. *Automatica*, 25(3):335–348, 1989. doi: 10.1016/0005-1098(89)90002-2.

[48] Olivier Ledoit and Michael Wolf. Improved estimation of the covariance matrix of stock returns with an application to portfolio selection. *Journal of Empirical Finance*, 10(5):603–621, 2003. doi: 10.1016/S0927-5398(03)00007-0.

[49] Andrew J. Patton. Volatility forecast comparison using imperfect volatility proxies. *Journal of Econometrics*, 160(1):246–256, 2011. doi: 10.1016/j.jeconom.2010.03.034.

[50] Paul Glasserman. *Monte Carlo Methods in Financial Engineering*. Springer, 2003. doi: 10.1007/978-0-387-21617-1.

[51] William F. Sharpe. Mutual fund performance. *Journal of Business*, 39(1):119–138, 1966.

[52] L. Rabiner and B. Juang. An introduction to hidden Markov models. *IEEE ASSP Magazine*, 3 (1):4–16, 1986. doi: 10.1109/MASSP.1986.1165342.

[53] Jan Bulla and Ingo Bulla. Stylized facts of financial time series and hidden semi-Markov models. *Computational Statistics & Data Analysis*, 51(4):2192–2209, 2006. doi: 10.1016/j.csda.2006.07.021.

[54] Robert C. Merton. Option pricing when underlying stock returns are discontinuous. *Journal of Financial Economics*, 3(1-2):125–144, 1976. doi: 10.1016/0304-405X(76)90022-2.

[55] S. G. Kou. A jump-diffusion model for option pricing. *Management Science*, 48(8):1086–1101, 2002. doi: 10.1287/mnsc.48.8.1086.166.

[56] Andrew J. Patton. Modelling asymmetric exchange rate dependence. *International Economic Review*, 47(2):527–556, 2006. doi: 10.1111/j.1468-2354.2006.00387.x.

[57] Umberto Cherubini, Elisa Luciano, and Walter Vecchiato. *Copula Methods in Finance*. John Wiley & Sons, 2004. ISBN 9780470863459.

[58] François Longin and Bruno Solnik. Extreme correlation of international equity markets. *The Journal of Finance*, 56(2):649–676, 2001.

[59] Andrew Ang and Joseph Chen. Asymmetric correlations of equity portfolios. *Journal of Financial Economics*, 63(3):443–494, 2002.

[60] Stephen J. Brown, William Goetzmann, Roger G. Ibbotson, and Stephen A. Ross. Survivorship bias in performance studies. *The Review of Financial Studies*, 5(4):553–580, 1992.

Tables

Table 1: **Weighting rules built on the score.** Each utility sums over the names with $\gamma_i > 0$ in the decision's name set $\mathcal{N}$; n_i are shares, S_i prices (USD/share), W wealth (USD) and $w_i = n_i S_i / W$ weights. The optimum under the budget alone pays the floors first and splits the rest among the names with $\gamma_i > 0$: the proportionality is normalized so that those weights sum to $\hat{w} = 1 - \sum_{\gamma_j \le 0} \varepsilon S_j / W$ for CD, CES and CESw, whose names with $\gamma_j \le 0$ sit at $\varepsilon = 10^{-3}$ shares, and to $1 - 0.01\,|\mathcal{N}|$ for SG, whose baseline is $b_i = 0.01\,W/S_i$ shares in every name. A hedge-pool name with $\gamma_i \le 0$ is instead free above its floor, with no utility term. Both CES rules use $\eta = 0.5$ and $\rho = 1 - 1/\eta = -1$; at $\eta = 1$ they reduce to CD. With the study's limits (Eq. 7) no closed form exists, and each utility is maximized numerically as a convex cone program over the cone in the last column, with the volatility cap as a second-order cone and the other limits as linear constraints. Derivations are in Appendix S1.

Rule	Utility $U(n;\gamma)$	Optimum under the budget alone	Names with $\gamma_i \le 0$	Cone
CD	$\sum \gamma_i \ln n_i$	$w_i \propto \gamma_i$	ε shares	exponential
CES	$\left(\sum \gamma_i n_i^{\rho}\right)^{1/\rho}$	$n_i \propto (\gamma_i / S_i)^{\eta}$, so $w_i \propto \gamma_i^{\eta} S_i^{1-\eta}$	ε shares	power
CESw	$\left(\sum \gamma_i w_i^{\rho}\right)^{1/\rho}$	$w_i \propto \gamma_i^{\eta}$	ε shares	power
SG	$\sum \gamma_i \ln(n_i - b_i)$	$w_i - 0.01 \propto \gamma_i$	b_i shares	exponential

Table 2: **Real-data backtest, hedge off.** Terminal wealth ratio W_T/W_0 from 2 January 2025 to 4 September 2026, the median of 30 runs (seeds 5660–5689) for each selection rule (row) and weighting rule (column). Range is the highest minus the lowest median in a row; the last row gives the same quantity within each column, across the ten selection rules. SPY buy-and-hold and the equal-weight universe (each of the 373 names bought at the first close and held) do not depend on the weighting rule. CD, Cobb–Douglas; CES, constant elasticity of substitution (CES) utility on share counts; CESw, CES utility on portfolio weights; SG, Stone–Geary; MinVar, minimum variance; Tangent, tangency portfolio; LS, 130/30 long-short Cobb–Douglas. With the market hedge off, no cap on portfolio β applied and no hedge names were held. Hedge-on values are in Table S7.

Selection rule	CD	CES	CESw	SG	MinVar	Tangent	LS	Range
33 names, 3 per sector								
γ per sector	1.230	1.219	1.221	1.234	1.086	1.227	1.110	0.148
Bandit per sector	1.226	1.234	1.223	1.239	1.162	1.229	1.140	0.099
Random per sector	1.205	1.183	1.201	1.183	1.128	1.158	1.127	0.078
13 names								
Global γ	1.520	1.384	1.460	1.441	1.355	1.341	1.248	0.272
Correlation-aware γ	1.352	1.386	1.421	1.380	1.284	1.404	1.219	0.202
Liquid then γ	1.333	1.310	1.379	1.350	1.238	1.361	1.305	0.141
Random liquid 100	1.216	1.173	1.211	1.212	1.158	1.274	1.152	0.122
Random whole universe	1.098	1.132	1.138	1.112	1.104	1.160	1.031	0.129
Hand-picked								
Hand-picked 13	1.548	1.571	1.614	1.631	1.443	1.697	1.549	0.254
Hand-picked 13 + WMT, PG	1.616	1.646	1.637	1.639	1.362	1.642	1.622	0.284
SPY buy-and-hold				1.317				
Equal-weight universe				1.296				
Range across selection rules	0.518	0.514	0.499	0.527	0.357	0.539	0.591	

Table 3: **Synthetic futures, by selection rule.** Lowest and highest of the seven weighting-rule medians of the terminal wealth ratio W_T/W_0 over the synthetic 2 January 2025 to 4 September 2026 window, for each selection rule, in the 600 typical futures and the 200 futures of each stress set. Typical futures were drawn at random from 10,000 synthetic markets; the burst and bear stress sets held the 200 markets with the deepest SPY drawdown in the window among those with a jump burst and among those without one. Each median is taken across futures; for random and bandit rules, each future first contributes the median of five seeds. SPY buy-and-hold and the equal-weight universe give the median across futures of each future's own SPY and equal-weight universe. With the market hedge on, each re-plan capped portfolio β at a level that was lower when the market was weaker, and up to five names with zero or negative β could be held. Every selection and weighting cell is shown in Fig. 6.

Selection rule	Typical	Burst stress	Bear stress
Hedge off			
33 names, 3 per sector			
γ per sector	1.142–1.227	0.967–1.001	0.813–0.874
Bandit per sector	1.135–1.227	0.946–0.988	0.803–0.844
Random per sector	1.100–1.180	0.924–0.951	0.777–0.876
13 names			
Global γ	1.200–1.240	0.980–1.010	0.837–0.858
Correlation-aware γ	1.206–1.230	0.978–1.016	0.837–0.868
Liquid then γ	1.208–1.244	0.972–1.021	0.842–0.861
Random liquid 100	1.138–1.191	0.944–0.997	0.774–0.969
Random whole universe	1.097–1.166	0.929–0.990	0.774–0.982
Hand-picked			
Hand-picked 13	1.175–1.242	0.951–1.031	0.785–0.998
Hand-picked 13 + WMT, PG	1.164–1.225	0.951–1.029	0.782–1.001
Hedge on			
33 names, 3 per sector			
γ per sector	1.135–1.223	0.961–0.991	0.846–0.892
Bandit per sector	1.122–1.223	0.956–0.980	0.838–0.872
Random per sector	1.090–1.176	0.942–0.964	0.814–0.880
13 names			
Global γ	1.189–1.239	0.982–1.012	0.863–0.889
Correlation-aware γ	1.197–1.229	0.976–1.007	0.872–0.896
Liquid then γ	1.186–1.236	0.986–1.006	0.871–0.896
Random liquid 100	1.137–1.178	0.945–0.990	0.827–0.978
Random whole universe	1.096–1.147	0.936–0.995	0.828–0.990
Hand-picked			
Hand-picked 13	1.164–1.217	0.946–1.029	0.832–0.999
Hand-picked 13 + WMT, PG	1.161–1.223	0.946–1.023	0.835–1.001
SPY buy-and-hold	1.185	0.865	0.718
Equal-weight universe	1.230	0.916	0.753

Figures

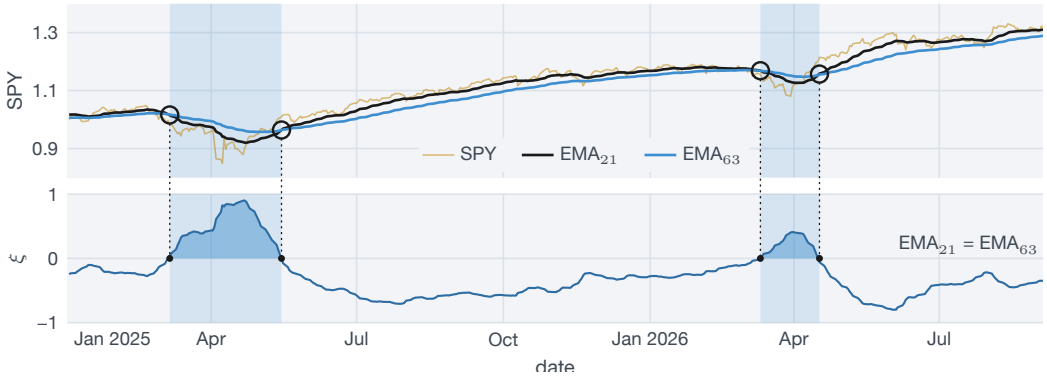


Figure 1: **SPY and the market state over the trading window, 2 January 2025 to 4 September 2026.** Top: SPY's daily close (light gold) with its 21-day (black) and 63-day (blue) exponential moving averages (EMAs), each divided by SPY's first close in the window. Bottom: the market state $\xi = -G(\mathrm{EMA}_{21}/\mathrm{EMA}_{63} - 1)$, where the gain $G = 21.23$ is the reciprocal of the 99th percentile of $|\mathrm{EMA}_{21}/\mathrm{EMA}_{63} - 1|$ over 2014–2024. Its positive part is filled. Shaded days have $\xi > 0$, when the 21-day EMA lies below the 63-day EMA; circles mark the crossings of the two EMAs, and dotted lines carry them down to the zeros of ξ.

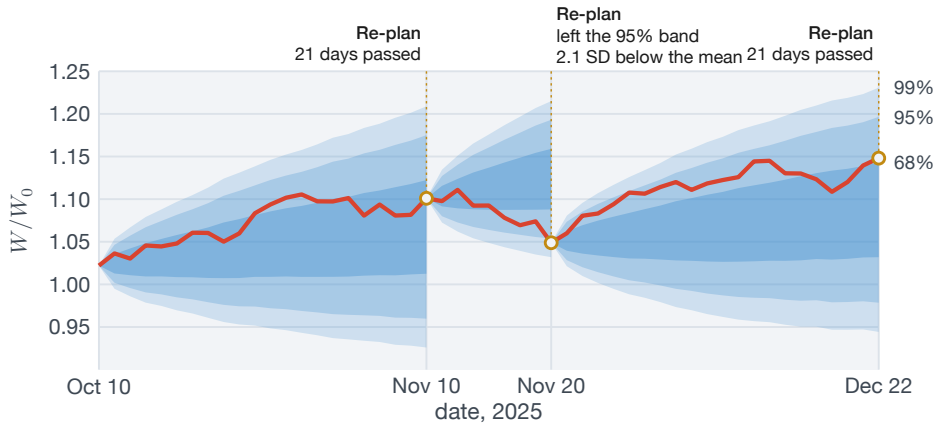


Figure 2: **One episode of the rebalancing controller on real prices, 10 October to 22 December 2025.** The red line is the wealth W of one backtest (the hand-picked 13 names, Cobb–Douglas weighting, market hedge off, seed 5660) divided by its initial wealth W_0. At each re-plan the controller projected 1,000 paths of a multi-asset geometric Brownian motion 21 trading days ahead; the nested cones are the 68, 95 and 99% bands of the projected portfolio value. The first and third re-plans came when 21 trading days had passed. The second came on 20 November, when wealth left the 95% band, 2.1 standard deviations below the projected mean.



Figure 3: **Real-data backtest: terminal wealth by selection rule and weighting rule, 2 January 2025 to 4 September 2026.** Each row is one selection rule, grouped by basket size (33 names, three per sector; 13 names; and the hand-picked lists of 13 and 15 names). Each dot is one of the seven weighting rules (Cobb–Douglas, CES on shares, CES on weights, Stone–Geary, minimum variance, tangency, and $130/30$ long-short), placed at the median over 30 runs (seeds 5660–5689) of terminal wealth divided by initial wealth. Filled dots have the market hedge off and open dots have it on; with the hedge on, each re-plan capped portfolio beta at a level that was lower when the market was weaker, and up to five stocks with zero or negative beta could be held. The pale bar spans each row's lowest to highest dot. The dashed gold line is SPY buy-and-hold (1.32) and the dotted line is the equal-weight universe (1.30).

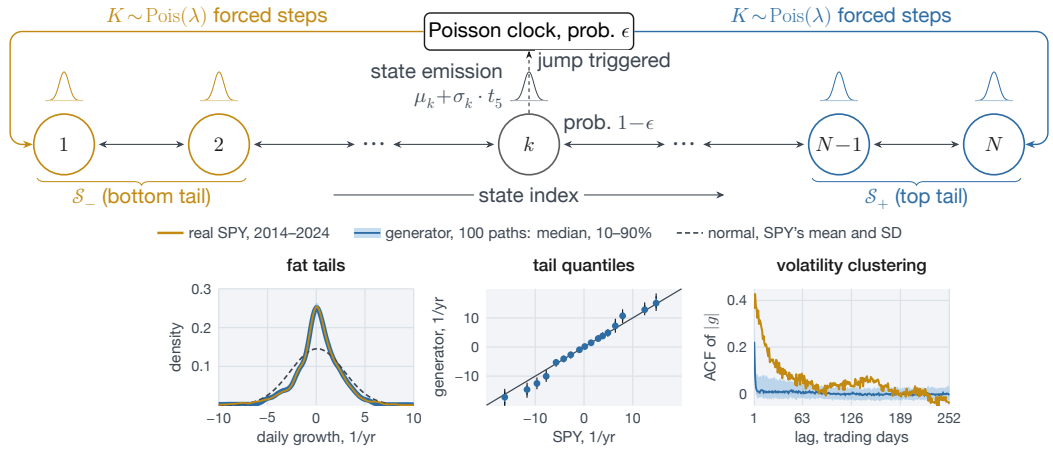


Figure 4: **The per-asset hidden Markov generator with jumps (JumpHMM) and its stylized facts for SPY.** Top: the generator's hidden states $1, \ldots, N$ are ordered from the bottom to the top tail of the daily growth rate, and state k emits $\mu_k + \sigma_k t_5$, a scaled Student-t draw with five degrees of freedom. With probability $1 - \epsilon$ the chain moves by its fitted transition matrix; with probability ϵ a jump episode starts and lasts $K \sim \mathrm{Pois}(\lambda)$ forced steps, each drawn from the bottom tail set $\mathcal{S}_-$ or the top tail set $\mathcal{S}_+$ [19]. Bottom: SPY's daily growth rates for 2014–2024 (gold) against 100 generator paths (median line and 10–90% band, blue) and a normal distribution with SPY's mean and standard deviation (dashed): the density, the generator's quantiles against SPY's, and the autocorrelation of the absolute growth rate $|g|$ out to 252 trading days.

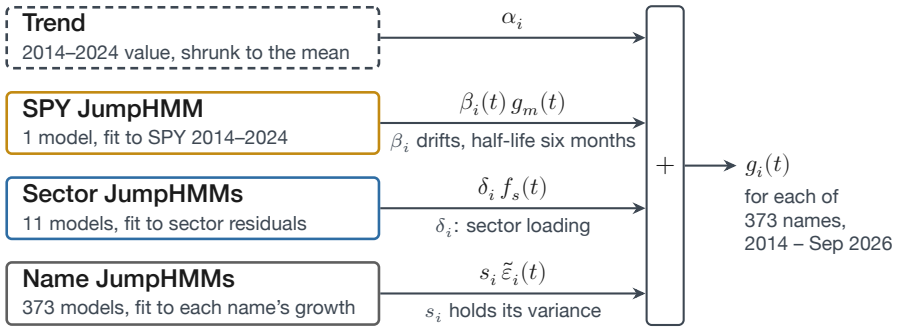


Figure 5: **Composition of one name's daily growth rate in the synthetic market.** Each name's growth rate $g_i(t)$ is the sum of a trend α_i (its 2014–2024 value, shrunk toward the cross-sectional mean), a market term $\beta_i(t)\, g_m(t)$ from one JumpHMM (hidden Markov generator with jumps) fit to SPY, a sector term $\delta_i f_s(t)$ from one of 11 sector JumpHMMs fit to sector residuals, and a name term $s_i\tilde{\varepsilon}_i(t)$ from the name's own JumpHMM, scaled so that the sum holds the name's variance. The loading $\beta_i(t)$ drifts with a half-life of 126 trading days. The name terms are drawn independently across names, so dependence across names enters only through the market and sector terms.

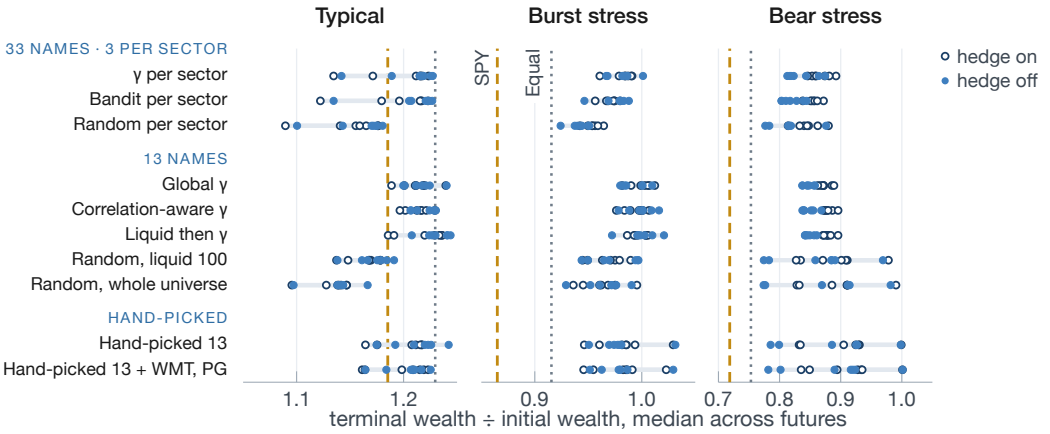


Figure 6: **Terminal wealth on synthetic futures of the trading window, by selection rule and weighting rule.** The three panels are 600 typical futures and two stress sets drawn from 10,000 futures: the 200 with the deepest SPY drawdowns that contain a jump burst (burst stress), and the 200 deepest without one (bear stress). Rows, dots, bars and lines are as in Fig. 3; each dot is the median across the set's futures (after a median over seeds within a future) of terminal wealth divided by initial wealth, and each panel has its own horizontal axis. SPY buy-and-hold ended at 1.19, 0.86 and 0.72, and the equal-weight universe at 1.23, 0.92 and 0.75, in the typical, burst-stress and bear-stress sets.

Online Appendix

S1 Allocation details

Closed forms under the budget alone. We derive the optima of Table 1 for a decision without hedge names. Let $\mathcal{P}$ be the names with $\gamma_i > 0$ and $F = W - \sum_{j \notin \mathcal{P}} \varepsilon S_j$ (USD) the wealth left after the floors of the other names. Each utility increases in every n_i with $i \in \mathcal{P}$, so the budget $\sum_{i \in \mathcal{P}} S_i n_i \le F$ binds. For CD, stationarity of the Lagrangian gives $\gamma_i / n_i = \nu S_i$ with multiplier $\nu > 0$ (USD^{-1}), so $S_i n_i = \gamma_i / \nu$, and the budget sets $\nu = \sum_{j \in \mathcal{P}} \gamma_j / F$, hence $w_i = (F/W)\, \gamma_i / \sum_{j \in \mathcal{P}} \gamma_j$. For CES with $0 < \rho < 1$ ($\eta > 1$) the rule maximizes the concave sum $\sum \gamma_i n_i^\rho$; with $\rho < 0$ ($0 < \eta < 1$) the outer power $1/\rho$ is negative, so the rule minimizes the convex sum $\sum \gamma_i n_i^\rho$. In both cases stationarity gives $\rho\, \gamma_i\, n_i^{\rho - 1} = \nu S_i$ with ν of the sign of ρ, and because $1/(\rho - 1) = -\eta$,

$$n_i = C \left(\frac{\gamma_i}{S_i} \right)^\eta, \qquad C = \frac{F}{\sum_{j \in \mathcal{P}} S_j (\gamma_j / S_j)^\eta}, \qquad w_i = \frac{F}{W} \frac{\gamma_i^\eta S_i^{1-\eta}}{\sum_{j \in \mathcal{P}} \gamma_j^\eta S_j^{1-\eta}}. \tag{S1}$$

As $\eta \to 0$ every name in $\mathcal{P}$ receives the same share count, a price-weighted book; as $\eta \to \infty$ the budget goes to the name with the largest γ_i / S_i. The same steps with w_i in place of n_i and the budget $\sum_{i \in \mathcal{P}} w_i \le F/W$ give CESw, $w_i = (F/W)\, \gamma_i^\eta / \sum_{j \in \mathcal{P}} \gamma_j^\eta$, which tends to equal weights as $\eta \to 0$ and to the name with the largest γ_i as $\eta \to \infty$. For SG every name of $\mathcal{N}$ holds $n_i \ge b_i$ and the names outside $\mathcal{P}$ hold exactly b_i. Writing $e_i = n_i - b_i$ for the excess shares, SG maximizes $\sum_{i \in \mathcal{P}} \gamma_i \ln e_i$ subject to $\sum_{i \in \mathcal{P}} S_i e_i \le W - \sum_{j \in \mathcal{N}} S_j b_j = W(1 - 0.01\,|\mathcal{N}|)$, which is the CD problem on the excess shares, so

$$w_i = 0.01 + (1 - 0.01\,|\mathcal{N}|)\, \frac{\gamma_i}{\sum_{j \in \mathcal{P}} \gamma_j}, \qquad i \in \mathcal{P}. \tag{S2}$$

A feasible SG book therefore needs $|\mathcal{N}| \le 100$, and a positive excess needs $|\mathcal{N}| < 100$. The unit tests checked Eq. S2 against the solver on a three-name problem to a relative tolerance of 10^{-4}, and on the hand-picked list the CD, CES and CESw solver weights stayed within 1.0×10^{-5} of their closed forms.

Cone programs. We use the cone conventions of MathOptInterface. The exponential cone is the closure of $\{(x, y, z) : y > 0,\ y\, e^{x/y} \le z\}$, so, with an auxiliary variable t_i per name, $(t_i, 1, n_i)$ in the cone is $t_i \le \ln n_i$. CD maximizes $\sum_{i \in \mathcal{P}} \gamma_i t_i$ with $n_i \ge 10^{-8}$, and SG puts $(t_i, 1, n_i - b_i)$ in the cone. The power cone with exponent $a \in (0, 1)$ is $\{(x, y, z) : x^a y^{1-a} \ge |z|,\ x, y \ge 0\}$. The CES rules act on a scaled variable $u_i = n_i \sum_{j \in \mathcal{P}} S_j / W$ (CES), which equals one at equal share counts that spend W, or $u_i = |\mathcal{P}|\, w_i$ (CESw), which equals one at equal weights; a common positive scale multiplies $\sum \gamma_i u_i^\rho$ by a constant and leaves the optimum unchanged, and it keeps u_i near one at any wealth and basket size. For $\eta > 1$ the program puts $(u_i, 1, t_i)$ in the power cone with exponent ρ, so $t_i \le u_i^\rho$, and maximizes $\sum \gamma_i t_i$. For $\eta < 1$ it puts $(t_i, u_i, 1)$ in the power cone with exponent η, so $t_i^\eta u_i^{1-\eta} \ge 1$, which is $t_i \ge u_i^\rho$, and minimizes $\sum \gamma_i t_i$. At $\eta = 1$ both CES rules call CD. The volatility cap puts $(\sigma_{\mathrm{SPY}}, L^\top w)$ in the second-order cone. The turnover budget uses slack variables $v_i \ge |n_i - n_i^{\mathrm{prev}}|$ with $\bar{c}\, \bar{S} \sum_i v_i \le B$, added only when the largest possible share change could break the budget. A name held at its floor enters the limits as a constant, and a limit row made only of constants is checked directly instead of being sent to the solver, because all-zero rows degraded Clarabel's accuracy (a share error of 1.3×10^{-5} with them against 2×10^{-6} without, on a CD test). LS uses two exponential-cone blocks, $(t_i^+, 1, n_i^+)$ for the long names with $\gamma_i > 0$ and $(t_i^-, 1, n_i^-)$ for the short names, with $n_i^\pm \ge 10^{-8}$, and applies the limits to the signed weights $w_i = (n_i^+ - n_i^-) S_i / W$. We set Clarabel's duality-gap tolerances to 10^{-12} because the utility is flat at its optimum, so the share error scales with the square root of the gap: on a three-name CD test the solver's shares missed the closed form by 2.5×10^{-5} at a gap tolerance of 10^{-9} and by 6×10^{-9} at 10^{-12}. An exception thrown inside a solver, such as the domain error that Clarabel 0.11.1 raised in its power-cone barrier for some CES problems at $\eta \le 0.25$ and $\eta \ge 4$, counted as that solver not accepting the problem, and the decision moved on to SCS.

The hedge stage. Let $\mathcal{H}_f$ be the free hedge names, the pool names with $\gamma_i \le 0$, and $n^{(1)}$ the first-stage shares. The second stage holds every name with a utility term at $n_i^{(1)}$ as a fixed variable, so

its share change still counts toward the turnover budget, holds every other name outside $\mathcal{H}_f$ at its floor, and solves

$$\min_{n_i \geq f_i,\ i \in \mathcal{H}_f} \sum_{i \in \mathcal{H}_f} w_i \quad \text{subject to the limits of Eq. 7 (and the 130/30 rows for LS),} \tag{S3}$$

with f_i the floor of hedge name i (ε for CD, CES and CESw, b_i for SG and 10^{-8} for LS). The first stage already subtracts $10^{-6} \sum_{i \in \mathcal{H}_f} w_i$ from the utility, but that cost is below the solver's dual accuracy and alone would leave hedge holdings of several percent of wealth along the flat direction; Eq. S3 removes them. For SG with no positive score and a finite β cap, the second stage runs with no fixed names and finds the least hedge holding that lets the baselines meet every limit.

Minimum variance and tangency. Both benchmarks use long-only shares with no floors and the expected growth rates $\mu_{g,i} = \alpha_i + \beta_i \tilde{g}_M$. They first solve $s = \min\left(1, \max \sum_i w_i\right)$ over the budget, single-name, β and turnover limits without the volatility cap. MinVar then minimizes $w^\top \Sigma_g w$ over the same limits with $\sum_i w_i \geq s - 10^{-7}$. Tangent, with excess growth $e = \mu_g - g_f \mathbf{1}$, first finds the best feasible excess growth $\bar{e} = \max e^\top w$ over that set; when $\bar{e} \leq 10^{-6}\, s \max_i |e_i|$ no feasible mix has positive excess growth and the rule falls back to the MinVar portfolio. Otherwise it applies the Charnes–Cooper change of variables [46] $y = \kappa n$, $v = \kappa w$ with $\kappa \geq 0$, and solves

$$\min_{y \geq 0,\ \kappa \geq 0} v^\top \Sigma_g\, v \quad \text{subject to} \quad e^\top v = \bar{e}, \qquad \textstyle\sum_i v_i \geq (s - 10^{-7})\, \kappa, \tag{S4}$$

with the right-hand side of every limit multiplied by κ, and sets $w_{\text{mix}} = v/\kappa$. Fixing $e^\top v$ at $\bar{e}$ instead of one leaves the optimum unchanged, because the Sharpe ratio does not depend on the scale of w, and it keeps $\kappa = \bar{e}/e^\top w_{\text{mix}} \geq 1$, so the division does not amplify solver error. Each mix is then scaled by $c = \min\left(1, \sigma_{\text{SPY}} / \sqrt{w_{\text{mix}}^\top \Sigma_g\, w_{\text{mix}}}\right)$ with the rest in cash. Scaling down keeps the budget, the single-name caps and the β cap, whose bound is never negative; the turnover budget is checked again, and a scaled book that breaks it keeps the previous holdings. An earlier version forced $\sum_i w_i$ to the largest investment the volatility cap allowed, which left a single feasible portfolio whenever the cap bound, so MinVar and Tangent coincided; the two-fund form, a risky mix scaled against cash [45], avoids this.

The bandit selection rule. One sector's bandit at a pick on day t runs as in Algorithm S1. The window starts d run over $\mathcal{D}_t = \{t - 272, \ldots, t - 21\}$, the 252 starts whose 21-day windows end on days $t - 251$ through t. The random stream of the sector with rank r (sectors in alphabetical order) is seeded with $(\text{seed} \cdot 1{,}000{,}003 + \text{date}) \cdot 64 + r$, where date is the day count of the pick date, so a pick depends only on the run seed, the date and the sector.

Parameters. The settings of the allocation, the score and the selection rules, with the calibrated values from 2014–2024, are in Table S1.

Algorithm S1 Bandit selection rule for one sector at a pick on day t.

Require: The sector's names $j = 1, \ldots, A$; excess growth rates $x_j(d)$ (yr^{-1}) for window starts $d \in \mathcal{D}_t$; a random stream seeded by the run seed, the date and the sector.
Ensure: The 3 names with the highest mean reward.

if $A \leq 3$ **then**
  return all A names
end if
$M \leftarrow \min(\max(50A, 500), 5000)$; $\hat{\mu}_j \leftarrow 0$ and $c_j \leftarrow 0$ for every j; draw a random permutation π of $1, \ldots, A$ as the tie-break rank
for $k = 1, \ldots, M$ **do**
  $\epsilon_k \leftarrow 1$ if $k = 1$, else $\max\{0.05,\ \min[1,\ k^{-1/3}(A \ln k)^{1/3}]\}$
  if $\mathrm{Uniform}(0, 1) < \epsilon_k$ or no arm has been pulled **then**
    $a \leftarrow$ an arm drawn uniformly from $1, \ldots, A$ {explore}
  else
    $a \leftarrow \arg\max_{j : c_j > 0} \hat{\mu}_j$, ties to the smallest π_j {exploit}
  end if
  $d \leftarrow$ a start drawn uniformly from $\mathcal{D}_t$; $c_a \leftarrow c_a + 1$; $\hat{\mu}_a \leftarrow \hat{\mu}_a + (x_a(d) - \hat{\mu}_a)/c_a$
end for
return the first 3 names ordered by: pulled before unpulled, then $\hat{\mu}_j$ descending, then π_j ascending

Table S1: **Settings of the allocation, the score and the selection rules.** Windows and spans are in trading days. Calibrated values were fitted once on 2014-01-03 to 2024-12-31 and then held fixed.

Setting	Symbol	Value
Score		
Half-life of the weighted least-squares fit	h ($\delta = 2^{-1/h}$)	63 days
Prior weight of the 2014–2024 fit	N_0	63 observations
SPY growth mean; variance (prior, calibrated)		0.106 yr^{-1}; 7.51 yr^{-2}
market growth average, EMA span	$\tilde{g}_M$	252 days
market state, EMA spans	ξ	21 and 63 days
market state gain (calibrated, 99% rule)	G	21.23
floor on $\lvert\beta_i\rvert$ in the score		0.1
Limits		
single-name cap: picks; other names	$\bar{w}_i$	$\max(0.10, 1.5/K)$; 0.10
volatility cap (calibrated)	σ_{SPY}	2.741 yr^{-1} (17.3% annualized)
residual covariance window; shrinkage		126 days; 0
half spread; turnover budget	$\bar{c}$; B	5 bps; $0.10\,W$
sector cap		none
Hedge		
β cap	$\bar{\beta}$	$\max(0, 1 - \xi_t)$
pool size; down-day beta lookback	H	5 names; 252 days
Weighting		
CES elasticity	η; ρ	0.5; -1
Stone–Geary baseline	$b_i S_i$	$0.01\,W$
floor of names with $\gamma_i \leq 0$ (CD, CES, CESw)	ε	10^{-3} shares
smallest traded target		$0.001\,W$
hedge tie-break cost		10^{-6} per unit weight
130/30 long; short; net weight		1.30; 0.30; 1.00
Selection		
names per sector (per-sector rules)		3 (33 in all)
names, universe-wide rules (main comparison)	K	13
liquid shortlist; dollar-volume window		100 names; 63 days
correlation limit; window		0.7; 252 days
bandit lookback; reward window		252 days; 21 days
bandit exploration floor; pulls	ϵ	0.05; $50A$ within $[500, 5000]$
pick interval; pick at band exit		21 days; yes
seeds		5660–5689
Solvers		
modeling package; solvers		JuMP 1.32.1; Clarabel 0.11.1, SCS 2.7.0
Clarabel gap; feasibility tolerance; iterations		10^{-12}; 10^{-9}; 500

S2 Controller and estimation details

This appendix gives the preparation of the price data, the estimators behind the controller's forecast and the details of the simulated fills. The settings of every real-data run are collected in Table S2 and the experiment grids in Table S3.

S2.1 Universe and price files

The sector snapshot lists 503 tickers. Of the 418 tickers that appear in all three price files with a sector, 373 have a positive open and close on all 3,187 trading days from 2014-01-03 to 2026-09-04: 15 in Communication Services, 39 in Consumer Discretionary, 31 in Consumer Staples, 16 in Energy, 53 in Financials, 44 in Health Care, 60 in Industrials, 50 in Information Technology, 17 in Materials, 22 in Real Estate and 26 in Utilities. Each price file is split-adjusted by the vendor as of its own creation date, so a split between the creation dates of two files shows as a one-day jump at their boundary. At each of the two boundaries we flagged a name when its overnight log price move exceeded 0.4 in magnitude, the move matched a standard split ratio to within a factor of $e^{0.1}$, its volume moved in the matching direction by a factor within 3 of that ratio, and the overnight log move of SPY was

within 0.05 in magnitude. We then divided the earlier prices of a flagged name by the split ratio and multiplied its earlier volumes by it, which leaves the growth rates within each file and the dollar volumes unchanged. This changed seven names, ORLY (ratio 15), NFLX (10), NOW (5) and FAST (2) at 2025-01-02, and KLAC (10), MNST (2) and APH (2) at 2026-01-02. Two one-day close moves larger than 0.9 in log magnitude remain, IR on 2020-03-02 and NI on 2015-07-02, both inside the calibration window, and we left them as given.

S2.2 Online single-index fit

For each name i the controller updates the intercept α_i (yr^{-1}) and loading β_i (dimensionless) of the single-index model $g_i(t) = \alpha_i + \beta_i\, g_M(t) + \varepsilon_i(t)$ once per trading day by exponentially weighted least squares. With half-life h (trading days) the decay factor is $\delta = 2^{-1/h}$, and the weighted sums are updated with the growth rates into day t as

$$\begin{aligned} &S_w \leftarrow \delta S_w + 1, && S_x \leftarrow \delta S_x + g_M(t), && S_y \leftarrow \delta S_y + g_i(t), \\ &S_{xx} \leftarrow \delta S_{xx} + g_M(t)^2, && S_{xy} \leftarrow \delta S_{xy} + g_i(t)\, g_M(t), \end{aligned} \tag{S5}$$

after which

$$\beta_i(t) \;=\; \frac{S_w S_{xy} - S_x S_y}{S_w S_{xx} - S_x^2}, \qquad \alpha_i(t) \;=\; \frac{S_y - \beta_i(t)\, S_x}{S_w}, \tag{S6}$$

with the previous estimates kept on a day when the denominator is at most 10^{-12} in magnitude. The sums start from a prior worth N_0 pseudo-observations. Let α_i^{cal} and β_i^{cal} be the ordinary least-squares fit of name i on SPY over the 2,766 daily growth rates of the calibration window, and let m_M (yr^{-1}) and v_M (yr^{-2}) be the mean and population variance of SPY's daily growth rate over the same days. The prior is

$$\begin{aligned} &S_w = N_0, \qquad S_x = N_0\, m_M, \qquad S_y = N_0\big(\alpha_i^{\mathrm{cal}} + \beta_i^{\mathrm{cal}} m_M\big), \\ &S_{xx} = N_0\big(v_M + m_M^2\big), \qquad S_{xy} = N_0\big(\alpha_i^{\mathrm{cal}} m_M + \beta_i^{\mathrm{cal}}(v_M + m_M^2)\big), \end{aligned} \tag{S7}$$

the sums that N_0 observations would give if SPY's growth had the calibration mean and variance and name i followed its calibration fit in the mean. Before the first update Eq. (S6) gives α_i^{cal} and β_i^{cal} exactly, and the weight of the prior halves every h days. We used $h = 63$ and $N_0 = 63$, with $m_M = 0.1061\ \mathrm{yr}^{-1}$ and $v_M = 7.511\ \mathrm{yr}^{-2}$, and the first update used the growth into 2025-01-02. SPY's standard deviation $\sigma_M(t)$ follows the same decay,

$$\begin{aligned} &Q_w \leftarrow \delta Q_w + 1, \qquad Q_x \leftarrow \delta Q_x + g_M(t), \qquad Q_{xx} \leftarrow \delta Q_{xx} + g_M(t)^2, \\ &\sigma_M^2(t) \;=\; \frac{Q_{xx}}{Q_w} - \left(\frac{Q_x}{Q_w}\right)^2, \end{aligned} \tag{S8}$$

from the prior $Q_w = N_0$, $Q_x = 0$ and $Q_{xx} = N_0\, \sigma_{\mathrm{SPY}}^2$, where $\sigma_{\mathrm{SPY}} = 2.741\ \mathrm{yr}^{-1}$ is the sample standard deviation of SPY's daily growth rate over the calibration window (0.1727 when multiplied by $\sqrt{\Delta t}$ to an annual volatility, the value that also sets the volatility cap).

S2.3 Residual covariance and the choice of the shrinkage weight

At a re-plan on day t the residual matrix holds $\varepsilon_i(s) = g_i(s) - \alpha_i(t) - \beta_i(t)\, g_M(s)$ for the names of the allocation problem over the L_Ω daily growth rates that end on day t. We write the sample covariance of these residuals as $\hat{\Omega}_\varepsilon = DRD$, with D the diagonal matrix of the sample standard deviations of the residual columns and R their sample correlation matrix, and Eq. (9) replaces R by $(1-\lambda)R + \lambda I$. A residual column with zero variance carries no correlation information and gets zero correlation with the others. We chose λ once, on the calibration window, by the accuracy of the residual part of Eq. (9) as a forecast of portfolio variance. The anchors were the last trading days of each month with at least 252 trading days of the window behind them and 21 ahead. At each anchor we fitted α_i and β_i by ordinary least squares on the trailing 252 daily growth rates, because the online fit has no warm-up period inside the calibration window, and formed $\hat{\Omega}_\varepsilon$ over the trailing $L_\Omega = 126$ growth rates. We then drew 200 portfolios per anchor, each of up to 33 names taken in turn from the 11 sectors with weights w drawn from a flat Dirichlet distribution. For each portfolio the realized

variance r is the sample variance of $w^{\top}\varepsilon(s)$ over the 21 days after the anchor and the prediction is $p(\lambda) = w^{\top}\big[(1-\lambda)\hat{\Omega}_{\varepsilon} + \lambda\,\mathrm{diag}(\hat{\Omega}_{\varepsilon})\big]w$. The loss of λ is the mean over portfolios and anchors of

$$\mathcal{L}(\lambda) \;=\; \frac{r}{p(\lambda)} - \ln\frac{r}{p(\lambda)} - 1, \tag{S9}$$

the QLIKE loss, which ranks variance forecasts consistently when the realized variance is a noisy proxy for the true one [49]. Over the grid $\lambda = 0, 0.1, \ldots, 1$ we took the λ with the smallest mean loss, breaking ties toward the larger λ. The mean loss was 0.353 at $\lambda = 0$ and rose steadily to 0.619 at $\lambda = 1$, so the study used $\lambda = 0$, the full sample residual covariance.

S2.4 Forecast simulation and fills

The simulation of Eq. (10) draws the daily log-price increments as $\Delta t\,\mu_g + \sqrt{\Delta t}\,L_{\Delta}z$, with z a vector of independent standard normal draws and L_{Δ} the lower Cholesky factor of $\Delta t\,\Sigma_g + 10^{-12}I$. The diagonal term keeps the factorization defined when Σ_g is singular and has no other effect at the scale of the forecast. The value of a 130/30 book on each path counts its short positions with negative shares. The forecast standard deviation uses the $N_p - 1$ denominator. The paths of a re-plan come from a Mersenne Twister generator seeded with $1{,}000{,}003$ times the run seed plus the serial day number of the decision date. The picks of a random rule use the same integer for the pick date multiplied by 64, and each sector's bandit adds the rank of its sector to that product, so every run is reproducible and does not depend on the thread that runs it. At a fill, the common factor that scales the buys is the largest value in $[0, 1]$ that keeps cash nonnegative after the buys and their costs, found by 60 bisection steps, and the average daily volume A_i in Eq. (3) is floored at one share. Names held before a re-plan that are not part of the new allocation problem are sold at the next open.

Table S2: **Controller and backtest settings.** Settings of every real-data run, from the study configuration, with the values fitted on the calibration window. Windows, horizons and intervals are in trading days, and growth rates are in yr^{-1}.

Setting	Symbol	Value
Windows and wealth		
Calibration window		2014-01-03 to 2024-12-31 (2,767 days)
Trading window		2025-01-02 to 2026-09-04 (420 days)
Length of a trading day	Δt	$1/252$ yr
Initial wealth	W_0	100,000 USD
Estimation		
Half-life of the online single-index fit	h	63 days
Prior weight of the online fit	N_0	63 pseudo-observations
SPY growth mean, 2014–2024	m_M	$0.1061\ \mathrm{yr}^{-1}$
SPY growth variance, 2014–2024	v_M	$7.511\ \mathrm{yr}^{-2}$
SPY growth SD, 2014–2024	σ_{SPY}	$2.741\ \mathrm{yr}^{-1}$ (0.1727 annualized)
Span of SPY's average growth rate		252 days
Residual covariance window	L_Ω	126 days
Residual shrinkage weight	λ	0 (chosen from $0, 0.1, \ldots, 1$)
Forecast, triggers and picks		
Forecast horizon	T	21 days
Forecast paths	N_p	1,000
Band half-width	z	1.96 standard deviations
Drawdown limit	$D_{\max}$	0.08, peak reset at each drawdown
Cash revisit		21 days
Pick interval		21 days, and at band exits
Smallest position traded		$0.001\,W$
Costs and cash		
Half spread	$\bar{c}$	5 bps
Impact coefficient	κ	0.001
Window of the average daily volume		63 days
Borrow fee on short positions	b	$0.0075\ \mathrm{yr}^{-1}$
Commission		0
Cash growth	g_f	DTB3, Eq. (2)
Experiment grids		
Seeds		5660 to 5689 (30 per cell)
Names per sector (per-sector rules)		3 (33 names)
Basket sizes (universe-wide rules)	K	$5, 8, 13, 21, 34$, and $100, 200$

Table S3: **Experiment grids of the real-data backtest.** Every grid runs the full trading window (2025-01-02 to 2026-09-04) with the seven weighting rules (CD, CES, CESw, SG, MinVar, Tangent and LS) unless the row says otherwise. Seeds run from 5660, and a deterministic selection rule ran with seed 5660 alone except in the seed rerun. The seed rerun repeats 140 runs of the Tier 1, universe-wide and hand-picked grids, and each repeated run counts once in the total.

Grid	**Selection rules and weighting rules**	**Hedge**	**Seeds**	**Runs**
Tier 1	γ, bandit and random per sector (33 names)	on, off	30 (γ: 1)	854
Fixed schedule	as Tier 1, picks on the 21-day timer only	on	30 (γ: 1)	427
Universe-wide	global γ, liquid then γ, correlation-aware γ, random from the 100 most liquid, random from the universe, at $K = 5, 8, 13, 21, 34$	on, off	30 (γ rules: 1)	4,410
Hand-picked	the 13-name list and its 15-name extension	on, off	1	28
Seed rerun	γ per sector; global γ at $K = 5, 8, 13, 21, 34$; correlation-aware γ and liquid then γ at $K = 13$; both hand-picked lists	on, off	30	4,200
Large book	global γ and random from the universe at $K = 100, 200$	on, off	30	1,680
Calendar only	the ten selection rules of the main comparison, with the band exit and the drawdown stop off	on, off	30	4,200
CES elasticity	as Tier 1; CES and CESw at $\eta = 0.1, 0.25, 0.5, 2, 4, 10$ with the single-name cap on and off, CD with the cap off (CES at $\eta = 0.1$ and the cap-on points at $\eta = 0.5$ left out)	on, off	30 (γ: 1)	2,562
Budget only	hand-picked 13; CD, and CES and CESw at $\eta = 0.05, 0.1, 0.25, 0.5, 2, 4, 10, 20$, with no single-name or volatility cap	off	1	17
Total	distinct runs			18,238

S3 Synthetic market details

The settings and fitted values of the generator are collected in Table S4, and those of the 11 sector models in Table S5. This appendix gives the calculations behind the sector loading scale, the variance of a composed name, the fits of the loading drift and the intercept pull, the choice of centering, and the frequency of jump bursts in the trading window.

For a sector s with n_s names, the real sector factor is $f_s(t) = n_s^{-1} \sum_{i \in s} e_i(t)$, with variance σ_s^2, so the sum of the real residual covariances over the ordered pairs $i \neq j$ of the sector is $n_s^2 \sigma_s^2 - \sum_{i \in s} \mathrm{Var}(e_i)$. In the generator, the shared part of two names in the sector is $\delta_i f_s(t)$ and $\delta_j f_s(t)$, whose covariance is $\delta_i \delta_j \sigma_s^2$ when the sector draw has the real factor variance, and the sum of $\delta_i \delta_j$ over the ordered pairs is $(\sum_i \delta_i)^2 - \sum_i \delta_i^2$. With δ_i^{loo} the slope of e_i on the mean residual of the other $n_s - 1$ names of the sector, the scale κ_s that makes the generator's mean pairwise covariance equal the real one is

$$\kappa_s^2 = \frac{n_s^2 \sigma_s^2 - \sum_{i \in s} \mathrm{Var}(e_i)}{\left[\left(\sum_{i \in s} \delta_i^{\mathrm{loo}} \right)^2 - \sum_{i \in s} \left(\delta_i^{\mathrm{loo}} \right)^2 \right] \sigma_s^2}, \tag{S10}$$

and $\delta_i = \kappa_s \delta_i^{\mathrm{loo}}$. The leave-one-out mean contains the common factor plus the average noise of the other $n_s - 1$ names, so the slope on it is attenuated as in an errors-in-variables regression. For names with equal loadings and a common residual correlation r_s, the slope is attenuated by the factor $r_s / (r_s + (1 - r_s)/(n_s - 1))$, and to first order

$$\kappa_s^2 \approx 1 + \frac{1 - r_s}{(n_s - 1)\, r_s}, \tag{S11}$$

so the scale is largest in small sectors with weak residual co-movement. With 15 names and $r_s = 0.094$, Eq. (S11) gives $\kappa_s \approx 1.30$ for Communication Services against a fitted 1.34, and with 17 names and $r_s = 0.166$ it gives 1.15 for Materials, as fitted. In Financials (53 names, $r_s = 0.239$), Utilities (26, 0.549) and Energy (16, 0.542), the fitted scales were 1.02. With the scaled loadings, every sector's mean residual correlation on the check paths came within 0.01 of its real value (Table S6).

The draws g_m, f_s and ε_i come from separate models with separate seeds, and the drift shocks $z_i(t)$ are independent of all three, so the variance of Eq. (12) on a day is approximately

$$\mathrm{Var}\left(g_i(t)\right) \approx \theta_i^2 \left(\mathrm{E}\left[\beta_i(t)^2 \right] v_m + \delta_i^2 v_s \right) + s_i^2 v_i, \tag{S12}$$

where v_m, v_s and v_i (yr^{-2}) are the variances of the market, sector and name draws. The omitted term, the variance of the drift times the squared mean of g_m, is at most $\bar{g}_m^2 / v_m$ times the first term, about $0.105^2 / 7.51 \approx 0.0015$. The stationary variance of the drift in Eq. (13) is $a^2 k_i^2 / (1 - (1 - a)^2) = k_i^2 a / (2 - a)$, so $\mathrm{E}[\beta_i(t)^2] = \hat{\beta}_i^2 + k_i^2 a / (2 - a)$, the term in Eq. (14). If each draw variance were exactly its real counterpart, with $v_m = \sigma_m^2$, $v_s = \sigma_s^2$ and $v_i = \sigma_i^2$, the rule $s_i^2 = 1 - \bar{\rho}_i$ would give $\mathrm{Var}(g_i) = \sigma_i^2$. Each JumpHMM instead drew a few percent more variance than its training series (the median standard deviation ratio of the name models was 1.037), and when the market, sector and name models inflate variance by similar ratios, Eq. (S12) lands near v_i, the variance of the name's own draw. The checks found median ratios of 0.994 to v_i and 1.066 to σ_i^2.

The loading drift was fitted on a bank of 20 synthetic paths, drawn once without drift and reused at every candidate value. On each bank path, a name was composed with a trial (h_β, k_i), the online estimator was started from the OLS fit of that path's 2014–2024 and the market moments of that path, and the standard deviation and lag-252 autocorrelation of $\hat{\beta}_i(t)$ over 2014–2024 were averaged over the bank. For each half-life, each k_i came from 30 bisection steps on the stationary standard deviation $k_i \sqrt{a/(2-a)}$ between 0 and 1.5, with $k_i = 0$ reserved for a name whose estimation noise alone reached its target, which happened for no name. The median lag-252 autocorrelation then rose with the half-life, from 0.059 at 63 trading days to 0.178 at 126, 0.307 at 252, 0.397 at 504, 0.447 at 1,008 and 0.481 at 2,016, and 126 came closest to the real 0.206. The median drift scale at the chosen half-life was 4.97, and the real median standard deviation of $\hat{\beta}_i(t)$ that the fit reproduced was 0.208.

The intercept pull used the same bank. OLS is linear in the growth path and α_i does not enter $\bar{\rho}_i$, so on each bank path and window the synthetic OLS intercept of name i is $\bar{\alpha} + \phi(\hat{\alpha}_i - \bar{\alpha}) + r_i$, where r_i is the

intercept of the same name composed with a zero intercept. We computed r_i once for every name, path and window (2014–2024, 2014–2019 and 2020–2024) and searched ϕ from 0 to 1.5 in steps of 0.01, jointly with an intercept drift whose stationary standard deviation ran from 0 to 0.15 yr^{-1} in steps of 0.0025 yr^{-1}. With spread the cross-name standard deviation of the 2014–2024 OLS intercepts and corr the cross-name correlation of the 2014–2019 and 2020–2024 intercepts, each computed on a path and averaged over the bank, the loss was $((\mathrm{spread} - 0.0764)/0.01)^2 + ((\mathrm{corr} - 0.175)/0.05)^2$. The joint minimum lay at $\phi = 0.6$ with no drift, where the bank gave a spread of 0.088 yr^{-1} and a correlation of 0.150 (loss 1.61). The real pooled statistic, the mean over names of the squared change in OLS intercept between 2014–2019 and 2020–2024 minus the mean over names of the sum of its two squared standard errors, was -0.0043 yr^{-2}, so the real changes were no larger than their noise. The synthetic spread contains the estimation noise of each name's intercept, whose real root-mean-square standard error is 0.075 yr^{-1}, so the spread could not fall below about 0.075 yr^{-1} for any ϕ, close to the real spread of 0.076 yr^{-1}, while the half-window correlation fell toward zero as ϕ shrank. A value of ϕ low enough to bring the spread within its tolerance would therefore have pushed the correlation further below its target, and the fitted ϕ left the spread 0.013 yr^{-1} above the real value on the check paths, the one failed check.

The composition centers each name's draw on the model's long-run mean $\bar{\varepsilon}_i$, not on the draw's own mean over the path. The full calendar holds 2,766 growth days in the calibration window, one growth day from 2024-12-31 to 2025-01-02, and 419 growth days in the trading window. Centering on the path mean forces $2{,}766\, m_{\mathcal{C}} + \tilde{\varepsilon}_i(2{,}767) + 419\, m_{\mathcal{W}} = 0$, where $m_{\mathcal{C}}$ and $m_{\mathcal{W}}$ are the mean draws in the two windows, so $m_{\mathcal{W}} \approx -6.6\, m_{\mathcal{C}}$. A name whose own draw raised its calibration intercept by chance would then have its trading-window intercept lowered by about 6.6 times that amount. On the bank, the noise parts of the two intercepts correlated -0.48 across names under path centering (from -0.71 to -0.16 across paths) and 0.003 when centered on the model's mean. Leaving the draws uncentered is not an alternative, since the long-run mean of a name model is the name's own mean growth (median 0.087 yr^{-1}), which would act as a second, persistent intercept and in the same diagnostic made the noise parts of the intercepts of 2014–2019 and 2020–2024 correlate 0.20.

A jump episode of the market model overlaps a trading window of $L = 419$ growth days if it starts inside the window or starts before it and is still under way on its first day. With ϵ small, the expected number of such episodes is ϵL for the first kind plus $\epsilon\, \mathrm{E}[K] = \epsilon\lambda$ for the second, where K is the episode length, so the probability that a window contains a burst is approximately

$$1 - \exp\left(-\epsilon(L + \lambda)\right) = 1 - \exp\left(-10^{-4}(419 + 118.9)\right) = 0.052, \tag{S13}$$

against 523 of 10,000 windows in the simulation. Over the whole calendar of 3,186 growth days, with no episode under way before the first day, $1 - \exp(-3{,}186\,\epsilon)$ gives a probability of 0.27 that a path contains at least one episode, most of which fall in the calibration window.

Every random stream is seeded from the path number p. The market draw uses seed p, sector s uses $10^{12} + 10^4 p + s$, name j uses $2 \times 10^{12} + 10^4 p + j$, and the drift shocks of all names come from one stream seeded with $3 \times 10^{12} + p$. The long-run means use fixed seeds above 4×10^{12}, and the tuning of name j and sector s used seeds $5660 + j$ and $7000 + s$. JumpHMM.jl reseeds the generator at each simulation call, so the draws do not depend on the number of threads, and building one path twice gave identical market data (Table S6).

Table S4: **Synthetic generator settings and fitted values.** JumpHMM settings shared by every model, the tuning grids, the tuned jump parameters, and the composition parameters of Eq. (12). Ranges across names are given as median (interquartile range) or as the smallest and largest value. The fit targets are real 2014–2024 statistics, and the fitted values are averages over the bank of 20 synthetic paths used for fitting. Growth rates and intercepts are in yr^{-1}, variances in yr^{-2}, and half-lives and episode lengths in trading days.

Quantity	**Symbol**	**Value**
JumpHMM settings, every model		
Number of states		50
Student-t degrees of freedom	ν	5
Probability that a forced step targets the negative tail	p_{neg}	0.52
States per tail set	N_{tail}	5
Risk-free rate; time step	r_f; Δt	0; $1/252$ yr
Tuning objective		ACF lags 1 to 25, kurtosis weight 0.20
Tuning grids		
Sector and name models, jump probability	ϵ	10 values from 10^{-4} to 2.5×10^{-2}
Sector and name models, mean episode length	λ	10 values from 10 to 120
Sector and name models, paths per grid point		100
Market model (fitted in earlier work)	ϵ, λ	15×15 grid, ϵ to 0.03, λ from 5 to 150, 200 paths
Tuned jump parameters		
Market model	ϵ, λ	10^{-4}, 118.9
Sector models (11)	ϵ, λ	10^{-4} for all; 10 to 120 (Table S5)
Name models (373), jump probability	ϵ	10^{-4} for 370; 2.87×10^{-3} for 2; 1.12×10^{-2} for 1
Name models, mean episode length	λ	58.9 (34.4 to 83.3); 10 to 120
Name models, draw SD divided by real SD (gate 0.5 to 2)		1.037; 1.006 to 1.251
Name models, long-run mean of the draw	$\bar{\varepsilon}_i$	median 0.087
Composition		
Real SPY growth variance, 2014–2024	σ_m^2	7.51 (SD 2.74)
Intercept pull toward the universe mean	ϕ	0.6
Universe mean of the real intercepts	$\bar{\alpha}$	-0.014
Cross-name SD of the synthetic intercepts	α_i	0.046
Intercept drift		none (fitted size zero)
Half-life of the loading drift	h_β	126
Drift scale	k_i	4.97 (3.74 to 6.45)
Stationary SD of the loading drift	$k_i\sqrt{a/(2-a)}$	0.26 (0.20 to 0.34); 0.03 to 0.77
Real market loading, 2014–2024	$\hat{\beta}_i$	0.98; 0.33 to 1.74
Sector loading	δ_i	0.94; -0.34 to 2.70
Sector scale of the leave-one-out loadings	κ_s	1.02 to 1.34
Floor on the own draw's variance share	$s^2_{\min}$	0.10
Real R^2 above which R^2 is kept		0.80 (largest real R^2: 0.645)
Factor share of variance	$\bar{\rho}_i$	0.51; 0.097 to 0.918
Names scaled by $\theta_i < 1$		2 (0.54% of name-days)
Fit targets (real) and fitted values (bank)		
Cross-name SD of OLS intercepts, 2014–2024		0.0764; 0.0881
Correlation of OLS intercepts, 2014–2019 with 2020–2024		0.175; 0.150
Median lag-252 autocorrelation of $\hat{\beta}_i(t)$		0.206; 0.178
Median SD of $\hat{\beta}_i(t)$		0.208; matched per name
Pooled change in intercept between halves, net of noise		-0.0043; -0.0002 yr^{-2}

Table S5: **Sector factors and their JumpHMMs.** For each sector, the number of names, the standard deviation of the real sector factor f_s over 2014–2024 (σ_s, yr^{-1}), the scale κ_s applied to the leave-one-out loadings, the median sector loading δ_i after scaling, the tuned mean episode length λ of the sector JumpHMM in trading days (every sector model has $\epsilon = 10^{-4}$), and the long-run mean of its draw (yr^{-1}), which enters the composition uncentered.

Sector	**Names**	σ_s	κ_s	**Median** δ_i	λ	**Mean**
Communication Services	15	1.500	1.337	0.784	71.1	−0.0007
Consumer Discretionary	39	1.719	1.084	0.834	95.6	0.0010
Consumer Staples	31	1.576	1.072	0.941	95.6	0.0013
Energy	16	4.215	1.024	0.968	10.0	0.0048
Financials	53	1.724	1.020	1.047	107.8	−0.0067
Health Care	44	1.400	1.085	0.911	46.7	0.0003
Industrials	60	1.350	1.056	1.002	107.8	−0.0010
Information Technology	50	1.471	1.069	0.813	71.1	−0.0005
Materials	17	2.004	1.151	0.770	95.6	−0.0027
Real Estate	22	2.293	1.040	0.952	71.1	−0.0028
Utilities	26	2.460	1.021	1.014	120.0	−0.0014

Table S6: **Generator checks against the real market.** A gated check passes when the synthetic value lies within the tolerance of the reference. References are real values, except for V3 and V4, which compare each composed name with its own JumpHMM draw, and V7, which builds one path twice. Each of the first 50 typical paths was recalibrated on its synthetic 2014–2024 data. A statistic with one value per path is the median over these paths, and one with a value per name is the median over names of each name's path median. V1 uses all 10,000 market paths. Statistics cover 2014–2024 unless a row names the trading window (2025-01-02 to 2026-09-04). Twenty-one of the 22 gated checks passed, and the test suite (V8) passed separately. Growth rates and intercepts are in yr^{-1}. ACF: autocorrelation of absolute growth rates. Residual correlation: mean pairwise correlation of single-index-model residuals. $\hat{\beta}_i$: the method's online estimate of the market loading.

Check	Statistic	Synthetic	Reference	Tolerance	Pass
Gated checks					
V1	SPY growth rate	0.105	0.106	± 0.02	yes
V2	Equal-weight growth rate, daily rebalancing minus buy and hold	-0.023	-0.020	± 0.02	yes
V3	Composed variance divided by the variance of the own draw	0.994	1	$[0.95, 1.05]$	yes
V3	Absolute gap between the OLS β and the mean of $\beta_i(t)$	0.021	0	≤ 0.05	yes
V4	ACF at lag 1, composed minus own draw	-0.038	0	± 0.05	yes
V4	ACF at lag 5, composed minus own draw	0.004	0	± 0.05	yes
V4	ACF at lag 20, composed minus own draw	0.006	0	± 0.05	yes
V5	Residual correlation, Communication Services	0.091	0.094	± 0.05	yes
V5	Residual correlation, Consumer Discretionary	0.108	0.112	± 0.05	yes
V5	Residual correlation, Consumer Staples	0.201	0.204	± 0.05	yes
V5	Residual correlation, Energy	0.535	0.542	± 0.05	yes
V5	Residual correlation, Financials	0.234	0.239	± 0.05	yes
V5	Residual correlation, Health Care	0.121	0.128	± 0.05	yes
V5	Residual correlation, Industrials	0.134	0.135	± 0.05	yes
V5	Residual correlation, Information Technology	0.098	0.102	± 0.05	yes
V5	Residual correlation, Materials	0.159	0.166	± 0.05	yes
V5	Residual correlation, Real Estate	0.361	0.366	± 0.05	yes
V5	Residual correlation, Utilities	0.540	0.549	± 0.05	yes
V6	Standard deviation of $\hat{\beta}_i$ divided by the real one	0.985	1	$[0.9, 1.1]$	yes
V7	Market-data fields that differ when one path is built twice	0	0	$= 0$	yes
V9	Cross-name standard deviation of the OLS intercepts	0.089	0.076	± 0.01	no
V9	Correlation of OLS intercepts, 2014–2019 with 2020–2024	0.134	0.175	± 0.05	yes
Reported, not gated					
V3	Composed variance divided by the real variance	1.066	1		
V4	ACF at lag 1, synthetic minus real	-0.097	0		
V4	ACF at lag 5, synthetic minus real	-0.171	0		
V4	ACF at lag 20, synthetic minus real	-0.097	0		
V5	Residual correlation across sectors	0.0002	0.031		
V6	Lag-252 autocorrelation of $\hat{\beta}_i$	0.188	0.206		
V6	Share of paths with a name whose $\hat{\beta}_i \leq 0$ on some day of the trading window (reference: real 2025–2026, 1 = yes)	1	1		
V6	Names with $\hat{\beta}_i \leq 0$ on some day of the trading window (synthetic range 2 to 11 over paths; reference: real 2025–2026)	7	120		
V6	As above, reference: median over 112 real 2014–2024 windows of the same length, ending every 21 trading days up to 2024-12-31 (range 0 to 53)	7	5		
V6	As above, reference: the real window of the same length ending 2024-12-31	7	53		
V9	Correlation of OLS intercepts, 2014–2024 with the trading window (reference: real 2025–2026)	0.100	-0.121		

S4 Additional results

The real-data backtest with the market hedge on, the market exposure of the portfolios on real prices and on the stress sets, the share of re-plans that the triggers brought forward and the trading cost they added, the comparison of the re-plan triggers with calendar-only re-planning on real prices and on the two stress sets, the paired comparisons of selection rules on synthetic futures, the percentile of each real result among the typical synthetic futures, and the comparison of global γ with random picks at basket sizes from 5 to 200 names are reported here (Tables S7, S8, S9, S10, S11, S12, S13 and S14; Figs. S1 and S2).

Table S7: **Real-data backtest, hedge on.** Terminal wealth ratio W_T/W_0 from 2 January 2025 to 4 September 2026, the median of 30 runs (seeds 5660–5689) for each selection rule (row) and weighting rule (column). Range is the highest minus the lowest median in a row; the last row gives the same quantity within each column, across the ten selection rules. SPY buy-and-hold and the equal-weight universe (each of the 373 names bought at the first close and held) do not depend on the weighting rule. CD, Cobb–Douglas; CES, constant elasticity of substitution (CES) utility on share counts; CESw, CES utility on portfolio weights; SG, Stone–Geary; MinVar, minimum variance; Tangent, tangency portfolio; LS, 130/30 long-short Cobb–Douglas. With the hedge on, every allocation also obeyed the β cap and could hold names from the hedge pool. Hedge-off values are in Table 2.

Selection rule	CD	CES	CESw	SG	MinVar	Tangent	LS	Range
33 names, 3 per sector								
γ per sector	1.212	1.232	1.223	1.218	1.073	1.285	1.130	0.212
Bandit per sector	1.125	1.140	1.167	1.194	1.166	1.214	1.120	0.094
Random per sector	1.108	1.087	1.114	1.155	1.155	1.092	1.092	0.068
13 names								
Global γ	1.503	1.305	1.464	1.461	1.203	1.589	1.283	0.386
Correlation-aware γ	1.409	1.429	1.479	1.447	1.213	1.366	1.292	0.266
Liquid then γ	1.322	1.285	1.347	1.331	1.184	1.317	1.336	0.163
Random liquid 100	1.167	1.156	1.174	1.178	1.123	1.155	1.112	0.066
Random whole universe	1.069	1.106	1.093	1.115	1.079	1.062	0.994	0.121
Hand-picked								
Hand-picked 13	1.554	1.517	1.517	1.511	1.233	1.582	1.586	0.353
Hand-picked 13 + WMT, PG	1.564	1.571	1.543	1.542	1.159	1.383	1.630	0.471
SPY buy-and-hold				1.317				
Equal-weight universe				1.296				
Range across selection rules	0.495	0.484	0.450	0.427	0.160	0.527	0.636	

Table S8: **Market exposure on real prices.** For each selection rule (row) and weighting rule (column), the median over the 30 runs (seeds 5660–5689) of the mean portfolio β with the hedge off and on, of the mean weight of hedge-pool names with the hedge on, and of the maximum drawdown, as the median with the hedge on minus the median with it off. The portfolio β of a day is the value-weighted online β of the holdings divided by total wealth, and its mean is taken over the days of a run. The first three rows of each block hold 33 names (3 per sector), the next five hold 13, and the last two are the hand-picked lists. With the market hedge on, each re-plan capped portfolio β at a level that was lower when the market was weaker, and up to five names with zero or negative β could be held.

Selection rule	CD	CES	CESw	SG	MinVar	Tangent	LS
Hedge off, mean portfolio β							
γ per sector	0.798	0.817	0.808	0.810	0.496	0.638	0.710
Bandit per sector	0.821	0.837	0.836	0.842	0.437	0.621	0.717
Random per sector	0.730	0.723	0.717	0.737	0.388	0.582	0.638
Global γ	0.761	0.762	0.749	0.748	0.760	0.742	0.605
Correlation-aware γ	0.774	0.787	0.783	0.779	0.761	0.758	0.639
Liquid then γ	0.784	0.784	0.787	0.785	0.782	0.759	0.621
Random liquid 100	0.764	0.767	0.774	0.785	0.626	0.787	0.626
Random whole universe	0.627	0.631	0.629	0.656	0.555	0.667	0.467
Hand-picked 13	0.868	0.882	0.897	0.903	0.863	0.891	0.771
Hand-picked 13 + WMT, PG	0.883	0.894	0.887	0.899	0.784	0.899	0.765
Hedge on, mean portfolio β							
γ per sector	0.791	0.803	0.793	0.795	0.340	0.544	0.705
Bandit per sector	0.793	0.800	0.799	0.813	0.338	0.536	0.695
Random per sector	0.681	0.663	0.661	0.682	0.314	0.478	0.592
Global γ	0.728	0.733	0.733	0.733	0.575	0.679	0.590
Correlation-aware γ	0.762	0.767	0.760	0.762	0.519	0.663	0.627
Liquid then γ	0.763	0.752	0.764	0.763	0.560	0.660	0.609
Random liquid 100	0.702	0.696	0.706	0.715	0.383	0.563	0.559
Random whole universe	0.564	0.549	0.557	0.578	0.342	0.470	0.379
Hand-picked 13	0.838	0.846	0.848	0.849	0.449	0.692	0.746
Hand-picked 13 + WMT, PG	0.831	0.839	0.836	0.843	0.417	0.643	0.730
Hedge on, hedge-pool weight (% of wealth)							
γ per sector	8.4	11.8	8.6	9.2	34.1	19.8	13.4
Bandit per sector	10.1	11.5	9.5	8.4	30.2	19.2	14.6
Random per sector	11.7	12.1	11.6	10.1	25.9	20.8	14.7
Global γ	27.9	28.9	29.8	28.5	39.8	30.9	33.5
Correlation-aware γ	25.5	26.1	25.8	26.0	39.3	28.4	32.1
Liquid then γ	29.2	29.7	29.7	29.5	39.2	29.9	33.7
Random liquid 100	21.8	21.7	22.0	21.4	34.0	29.7	26.9
Random whole universe	21.1	21.4	21.4	21.0	33.0	27.6	26.0
Hand-picked 13	23.2	22.9	22.4	22.6	37.6	31.4	29.6
Hand-picked 13 + WMT, PG	21.3	22.2	20.6	20.5	35.4	32.3	29.0
Maximum drawdown, hedge on minus hedge off							
γ per sector	0.000	−0.006	−0.001	0.000	0.000	−0.009	0.000
Bandit per sector	0.010	0.007	0.002	0.001	0.011	0.007	0.008
Random per sector	0.002	0.001	−0.001	−0.003	0.000	0.008	0.002
Global γ	0.000	0.008	0.007	0.000	0.001	−0.009	0.003
Correlation-aware γ	0.000	0.000	0.000	0.000	0.004	−0.001	0.005
Liquid then γ	−0.002	0.008	0.001	0.000	−0.009	0.005	0.005
Random liquid 100	0.003	0.005	0.005	−0.001	0.009	0.019	0.015
Random whole universe	0.002	0.005	0.002	0.000	−0.002	−0.018	0.003
Hand-picked 13	0.002	0.000	−0.007	0.003	0.005	0.003	−0.009
Hand-picked 13 + WMT, PG	−0.003	−0.003	−0.003	0.004	−0.006	0.000	0.006

Table S9: **Early re-plans and trading cost on real prices.** For each selection rule (row) and weighting rule (column), the median over the 30 runs (seeds 5660–5689) of the share of re-plans after the first allocation that came before the 21-day horizon, because the portfolio value left the 95% band of the forecast or fell 8% below its peak, and the percentage by which the median trading cost of these runs exceeded that of their 30 calendar-only twins (Table S10). The other re-plans came at the horizon or on a basket change. The first three rows of each block hold 33 names (3 per sector), the next five hold 13, and the last two are the hand-picked lists. With the market hedge on, each re-plan capped portfolio β at a level that was lower when the market was weaker, and up to five names with zero or negative β could be held.

Selection rule	CD	CES	CESw	SG	MinVar	Tangent	LS
Hedge off, share of re-plans brought forward by the band or the drawdown stop							
γ per sector	0.42	0.32	0.42	0.42	0.40	0.56	0.50
Bandit per sector	0.33	0.36	0.31	0.31	0.29	0.37	0.46
Random per sector	0.25	0.24	0.20	0.24	0.30	0.25	0.31
Global γ	0.48	0.36	0.41	0.43	0.41	0.45	0.55
Correlation-aware γ	0.50	0.38	0.41	0.47	0.50	0.52	0.61
Liquid then γ	0.29	0.29	0.28	0.29	0.38	0.53	0.52
Random liquid 100	0.31	0.33	0.29	0.33	0.27	0.32	0.40
Random whole universe	0.26	0.26	0.25	0.26	0.28	0.29	0.37
Hand-picked 13	0.37	0.40	0.40	0.43	0.35	0.44	0.39
Hand-picked 13 + WMT, PG	0.41	0.37	0.37	0.37	0.37	0.40	0.39
Hedge on, share of re-plans brought forward by the band or the drawdown stop							
γ per sector	0.42	0.35	0.42	0.42	0.39	0.52	0.46
Bandit per sector	0.37	0.33	0.35	0.38	0.36	0.44	0.51
Random per sector	0.27	0.29	0.25	0.25	0.39	0.35	0.35
Global γ	0.47	0.41	0.50	0.48	0.45	0.47	0.56
Correlation-aware γ	0.45	0.43	0.47	0.45	0.45	0.50	0.61
Liquid then γ	0.37	0.38	0.37	0.37	0.41	0.47	0.53
Random liquid 100	0.33	0.33	0.34	0.31	0.33	0.36	0.45
Random whole universe	0.33	0.31	0.31	0.31	0.37	0.35	0.42
Hand-picked 13	0.43	0.38	0.38	0.42	0.28	0.36	0.41
Hand-picked 13 + WMT, PG	0.44	0.37	0.37	0.39	0.33	0.48	0.40
Hedge off, trading cost, % above the calendar-only twin							
γ per sector	11	7	9	10	23	21	22
Bandit per sector	26	25	22	24	27	29	31
Random per sector	13	15	12	14	31	20	21
Global γ	28	19	18	19	18	24	29
Correlation-aware γ	30	17	15	20	28	22	32
Liquid then γ	5	3	6	5	10	20	17
Random liquid 100	20	22	20	19	17	21	32
Random whole universe	18	16	18	16	20	15	33
Hand-picked 13	8	3	6	10	5	9	17
Hand-picked 13 + WMT, PG	10	10	10	10	8	12	22
Hedge on, trading cost, % above the calendar-only twin							
γ per sector	11	6	12	12	26	16	18
Bandit per sector	23	18	21	23	34	29	35
Random per sector	15	13	13	9	45	24	21
Global γ	21	24	26	25	21	17	31
Correlation-aware γ	21	17	24	22	25	20	35
Liquid then γ	24	15	24	25	20	25	22
Random liquid 100	21	22	19	18	28	22	27
Random whole universe	23	22	21	22	35	23	30
Hand-picked 13	20	4	4	5	19	8	15
Hand-picked 13 + WMT, PG	12	10	10	6	19	7	18

Table S10: **Re-plan triggers against calendar-only re-planning on real prices.** Each entry is the median of the 30 runs (seeds 5660–5689) with the band exit and the drawdown stop, minus the median of the 30 runs with both switched off, for each selection rule (row) and weighting rule (column). Without the two triggers a run re-planned only at the 21-day horizon, 19 times over the window. A positive terminal wealth entry means the triggered runs ended higher, and a negative maximum drawdown entry means their worst drawdown was smaller. The first three rows of each block hold 33 names (3 per sector), the next five hold 13, and the last two are the hand-picked lists. With the market hedge on, each re-plan capped portfolio β at a level that was lower when the market was weaker, and up to five names with zero or negative β could be held. The triggered medians are those of Tables 2 and S7.

Selection rule	CD	CES	CESw	SG	MinVar	Tangent	LS
Hedge off, terminal wealth ratio							
γ per sector	0.003	−0.016	−0.002	0.007	−0.019	−0.075	−0.062
Bandit per sector	0.013	0.031	0.004	0.010	−0.016	0.003	0.007
Random per sector	0.000	−0.006	0.005	−0.014	−0.015	−0.067	−0.058
Global γ	0.191	0.084	0.104	0.099	0.120	0.002	−0.017
Correlation-aware γ	−0.006	0.071	0.043	0.013	0.069	0.055	−0.093
Liquid then γ	−0.063	−0.057	−0.063	−0.064	−0.046	−0.005	−0.022
Random liquid 100	0.013	0.007	0.017	0.008	0.026	0.021	0.004
Random whole universe	−0.037	−0.008	−0.006	−0.039	−0.014	−0.054	−0.050
Hand-picked 13	−0.001	0.004	0.034	0.037	0.003	0.009	−0.019
Hand-picked 13 + WMT, PG	0.111	0.122	0.112	0.093	0.012	−0.010	0.093
Hedge off, maximum drawdown							
γ per sector	0.005	−0.015	0.006	0.006	−0.015	−0.007	−0.025
Bandit per sector	−0.007	−0.009	−0.005	−0.002	−0.002	0.000	−0.009
Random per sector	−0.005	−0.007	−0.008	−0.001	−0.006	0.009	0.002
Global γ	−0.053	−0.024	−0.052	−0.053	−0.014	−0.021	−0.039
Correlation-aware γ	−0.013	−0.035	−0.013	−0.013	−0.022	−0.012	−0.031
Liquid then γ	0.015	0.001	0.010	0.013	−0.026	−0.012	−0.007
Random liquid 100	−0.011	−0.019	−0.023	−0.010	0.000	−0.004	−0.019
Random whole universe	−0.010	−0.008	−0.010	−0.009	0.004	0.012	−0.006
Hand-picked 13	−0.020	−0.022	−0.001	−0.008	−0.004	−0.002	0.002
Hand-picked 13 + WMT, PG	−0.036	−0.043	−0.038	−0.029	−0.002	0.000	−0.038
Hedge on, terminal wealth ratio							
γ per sector	−0.006	0.010	0.010	0.002	−0.070	−0.007	−0.045
Bandit per sector	−0.042	−0.022	0.002	0.019	−0.046	−0.033	−0.028
Random per sector	−0.049	−0.068	−0.047	0.004	−0.016	−0.107	−0.085
Global γ	0.081	−0.085	0.011	0.025	0.045	0.183	−0.077
Correlation-aware γ	−0.011	0.051	0.034	0.015	0.060	−0.035	−0.113
Liquid then γ	−0.165	−0.152	−0.180	−0.174	0.032	−0.069	−0.086
Random liquid 100	−0.006	0.013	0.007	0.022	−0.007	−0.015	−0.042
Random whole universe	−0.080	−0.026	−0.049	−0.046	−0.041	−0.055	−0.090
Hand-picked 13	0.036	−0.022	−0.007	−0.040	−0.045	0.025	−0.010
Hand-picked 13 + WMT, PG	0.087	0.069	0.063	0.034	−0.058	−0.076	0.064
Hedge on, maximum drawdown							
γ per sector	0.005	−0.021	0.005	0.005	−0.018	−0.016	−0.025
Bandit per sector	−0.003	−0.007	−0.008	−0.008	0.003	−0.002	−0.011
Random per sector	−0.009	−0.011	−0.014	−0.009	−0.010	0.011	0.001
Global γ	−0.053	−0.015	−0.045	−0.054	−0.018	−0.021	−0.036
Correlation-aware γ	−0.013	−0.035	−0.013	−0.013	−0.025	−0.013	−0.027
Liquid then γ	0.003	−0.004	0.001	0.003	−0.046	−0.010	−0.014
Random liquid 100	−0.010	−0.015	−0.018	−0.013	−0.002	−0.008	−0.030
Random whole universe	−0.021	−0.012	−0.021	−0.015	−0.010	−0.013	−0.027
Hand-picked 13	−0.018	−0.022	−0.008	−0.005	−0.001	−0.001	−0.016
Hand-picked 13 + WMT, PG	−0.039	−0.046	−0.040	−0.026	−0.014	−0.003	−0.045

Table S11: **Re-plan triggers against calendar-only re-planning in the burst stress set.** Each entry is the median across the 200 synthetic futures of the set (the deepest SPY drawdowns among futures with a jump burst) of the runs with the band exit and the drawdown stop, minus the same median of the runs with both switched off, for each selection rule (row) and weighting rule (column). Within a future, the bandit and random rules contribute the median over seeds 5660–5664, and the other rules ran once. Without the two triggers a run re-planned every 21 trading days, 19 times over the window, at the horizon or, when its last re-plan had gone to cash, on the cash revisit. A positive terminal wealth entry means the triggered runs ended higher, and a negative maximum drawdown entry means their worst drawdown was smaller. Rows are grouped as in Table S10, and the triggered medians are those of Fig. 6.

Selection rule	CD	CES	CESw	SG	MinVar	Tangent	LS
Hedge off, terminal wealth ratio							
γ per sector	−0.009	0.008	−0.002	−0.002	0.002	−0.016	0.006
Bandit per sector	0.008	−0.007	−0.002	−0.012	0.004	0.015	−0.010
Random per sector	0.000	0.003	0.003	0.007	−0.008	0.003	0.002
Global γ	−0.027	0.001	−0.006	−0.031	−0.022	−0.009	−0.026
Correlation-aware γ	−0.013	−0.008	−0.013	−0.015	−0.025	−0.013	−0.002
Liquid then γ	0.007	−0.009	0.002	0.014	0.003	0.005	0.014
Random liquid 100	0.022	−0.009	0.021	−0.006	0.009	0.013	0.016
Random whole universe	0.007	0.002	0.004	−0.001	−0.007	−0.011	−0.001
Hand-picked 13	−0.004	−0.007	−0.002	0.000	0.004	0.011	0.024
Hand-picked 13 + WMT, PG	0.003	−0.008	−0.001	0.013	0.005	−0.004	0.021
Hedge off, maximum drawdown							
γ per sector	−0.006	−0.003	−0.007	−0.005	−0.014	−0.001	−0.006
Bandit per sector	0.000	−0.001	−0.003	−0.002	−0.002	0.002	−0.008
Random per sector	−0.013	−0.013	−0.016	−0.008	−0.007	−0.001	−0.017
Global γ	0.000	−0.004	−0.009	−0.004	−0.002	−0.005	−0.002
Correlation-aware γ	−0.002	−0.005	−0.006	0.002	0.001	−0.007	0.001
Liquid then γ	−0.009	−0.008	−0.009	−0.010	−0.009	−0.003	0.000
Random liquid 100	−0.010	−0.010	−0.013	−0.011	−0.004	−0.003	−0.004
Random whole universe	−0.011	−0.010	−0.009	−0.004	−0.002	−0.003	−0.012
Hand-picked 13	−0.022	−0.020	−0.020	−0.012	−0.009	−0.006	−0.022
Hand-picked 13 + WMT, PG	−0.020	−0.020	−0.021	−0.015	−0.008	−0.005	−0.019
Hedge on, terminal wealth ratio							
γ per sector	0.009	0.002	0.002	−0.003	−0.001	0.017	−0.010
Bandit per sector	−0.013	−0.019	−0.013	−0.005	−0.002	0.003	−0.007
Random per sector	0.012	0.019	0.014	−0.006	0.001	0.000	0.015
Global γ	−0.003	−0.002	−0.007	0.003	−0.003	−0.010	−0.018
Correlation-aware γ	−0.012	−0.002	−0.023	−0.016	−0.014	−0.015	−0.026
Liquid then γ	−0.004	0.004	−0.016	0.002	0.015	0.003	−0.014
Random liquid 100	0.000	−0.010	0.007	0.009	0.015	0.006	−0.002
Random whole universe	−0.007	0.001	−0.002	0.013	−0.013	−0.011	0.005
Hand-picked 13	−0.007	−0.017	−0.004	−0.011	0.000	−0.010	0.021
Hand-picked 13 + WMT, PG	−0.010	−0.008	−0.007	0.002	−0.007	−0.010	0.010
Hedge on, maximum drawdown							
γ per sector	−0.007	−0.009	−0.008	−0.010	−0.007	−0.004	−0.011
Bandit per sector	−0.005	−0.013	−0.008	−0.013	−0.004	−0.002	−0.011
Random per sector	−0.018	−0.015	−0.017	−0.016	−0.009	−0.010	−0.012
Global γ	−0.012	−0.010	−0.014	−0.012	−0.014	−0.009	−0.006
Correlation-aware γ	−0.011	−0.005	−0.012	−0.014	−0.002	−0.005	−0.004
Liquid then γ	−0.009	−0.009	−0.008	−0.012	−0.014	−0.007	−0.009
Random liquid 100	−0.013	−0.013	−0.014	−0.014	−0.007	−0.011	−0.003
Random whole universe	−0.014	−0.017	−0.009	−0.009	−0.011	−0.003	−0.017
Hand-picked 13	−0.018	−0.015	−0.021	−0.020	−0.011	−0.007	−0.022
Hand-picked 13 + WMT, PG	−0.021	−0.016	−0.020	−0.015	−0.009	0.000	−0.020

Table S12: **Re-plan triggers against calendar-only re-planning in the bear stress set.** Entries as in Table S11, across the 200 synthetic futures of the bear stress set (the deepest SPY drawdowns among futures without a jump burst).

Selection rule	CD	CES	CESw	SG	MinVar	Tangent	LS
Hedge off, terminal wealth ratio							
γ per sector	0.000	0.001	0.005	0.006	0.003	0.007	0.008
Bandit per sector	0.001	0.004	−0.005	0.000	−0.010	−0.001	0.004
Random per sector	0.004	0.008	0.010	0.000	−0.004	−0.002	0.013
Global γ	−0.004	0.011	0.004	−0.005	0.010	−0.004	−0.019
Correlation-aware γ	−0.004	0.004	0.000	−0.007	0.008	−0.003	0.000
Liquid then γ	0.013	0.018	0.012	0.013	0.012	0.006	−0.009
Random liquid 100	0.011	0.004	0.002	0.005	0.002	0.010	0.017
Random whole universe	0.014	0.019	0.017	−0.003	−0.001	0.001	0.008
Hand-picked 13	0.016	0.019	0.020	0.004	0.006	0.002	0.029
Hand-picked 13 + WMT, PG	0.004	0.013	0.006	0.015	0.004	0.004	0.024
Hedge off, maximum drawdown							
γ per sector	−0.010	−0.004	−0.011	−0.006	−0.012	−0.007	−0.009
Bandit per sector	−0.006	−0.002	−0.006	−0.005	−0.001	−0.005	−0.011
Random per sector	−0.013	−0.009	−0.011	−0.005	0.000	−0.001	−0.014
Global γ	−0.012	−0.001	−0.008	−0.011	−0.009	0.002	−0.009
Correlation-aware γ	−0.013	−0.004	−0.015	−0.014	−0.009	0.007	−0.008
Liquid then γ	−0.009	−0.018	−0.014	−0.010	−0.007	−0.011	−0.007
Random liquid 100	−0.010	−0.014	−0.010	−0.009	−0.004	−0.007	−0.006
Random whole universe	−0.013	−0.012	−0.015	−0.007	0.001	−0.003	−0.009
Hand-picked 13	−0.008	−0.009	−0.011	−0.007	−0.006	0.001	−0.013
Hand-picked 13 + WMT, PG	−0.011	−0.011	−0.010	−0.008	−0.003	−0.008	−0.014
Hedge on, terminal wealth ratio							
γ per sector	0.012	0.003	0.014	0.017	0.004	0.006	0.011
Bandit per sector	0.012	−0.002	0.010	0.009	−0.005	−0.003	0.002
Random per sector	−0.003	−0.001	−0.001	−0.007	−0.001	0.002	0.002
Global γ	−0.003	0.012	0.002	0.007	0.014	−0.003	0.005
Correlation-aware γ	−0.001	0.016	0.008	0.001	0.023	0.012	0.014
Liquid then γ	0.005	0.011	0.006	0.005	0.007	0.002	0.001
Random liquid 100	0.021	0.022	0.014	0.008	−0.002	−0.015	0.020
Random whole universe	0.011	0.009	0.012	0.004	0.003	0.009	0.019
Hand-picked 13	−0.002	0.008	0.008	0.009	0.000	0.001	0.022
Hand-picked 13 + WMT, PG	0.015	0.011	0.016	0.009	0.003	0.017	0.017
Hedge on, maximum drawdown							
γ per sector	−0.010	−0.003	−0.008	−0.009	−0.003	−0.014	−0.013
Bandit per sector	−0.004	−0.001	−0.008	−0.007	−0.004	0.002	−0.004
Random per sector	−0.016	−0.012	−0.012	−0.008	−0.004	−0.006	−0.007
Global γ	−0.013	−0.013	−0.005	−0.011	−0.012	0.002	−0.016
Correlation-aware γ	−0.018	−0.014	−0.009	−0.017	−0.006	−0.005	−0.015
Liquid then γ	−0.007	−0.012	−0.006	−0.007	−0.012	−0.007	−0.012
Random liquid 100	−0.016	−0.015	−0.018	−0.011	−0.001	−0.001	−0.011
Random whole universe	−0.014	−0.011	−0.014	−0.010	−0.004	0.004	−0.009
Hand-picked 13	−0.013	−0.013	−0.017	−0.017	0.000	−0.001	−0.015
Hand-picked 13 + WMT, PG	−0.015	−0.015	−0.017	−0.013	−0.004	−0.011	−0.015

Table S13: **Selection by γ against random picks on synthetic futures.** Share of futures (%) in which the rule that selected by γ ended with a higher terminal wealth ratio W_T/W_0 than random picks of the same size, by weighting rule (column), stress set and hedge setting. With the market hedge on, each re-plan capped portfolio β at a level that was lower when the market was weaker, and up to five names with zero or negative β could be held. Global γ held the K highest-γ names in the universe and was compared with K names drawn at random from the whole universe; γ per sector was compared with three random names per sector. Each random rule contributed the median of five seeds in each future. A value of 50 means that each rule ended higher in half of the futures.

Comparison	CD	CES	CESw	SG	MinVar	Tangent	LS
Typical futures (600), hedge off							
Global γ vs. random, 5 names	62.3	62.3	64.8	64.5	62.5	61.3	64.2
Global γ vs. random, 13 names	70.7	71.2	70.5	69.7	64.5	62.3	70.0
Global γ vs. random, 34 names	70.2	68.2	71.0	70.8	67.2	58.0	69.0
γ per sector vs. random per sector	73.2	72.7	73.2	73.3	64.8	60.0	73.7
Typical futures (600), hedge on							
Global γ vs. random, 5 names	63.2	62.7	64.2	64.5	62.0	60.2	63.2
Global γ vs. random, 13 names	68.8	69.8	68.7	68.8	65.7	63.2	71.5
Global γ vs. random, 34 names	67.7	68.3	70.3	71.2	65.3	60.7	68.0
γ per sector vs. random per sector	72.3	71.7	73.7	71.5	66.7	63.5	73.2
Burst stress (200), hedge off							
Global γ vs. random, 5 names	53.0	53.0	52.5	53.5	66.5	62.5	57.5
Global γ vs. random, 13 names	52.5	53.0	55.0	57.5	67.0	67.5	53.5
Global γ vs. random, 34 names	58.5	63.5	61.0	65.0	60.5	55.0	56.5
γ per sector vs. random per sector	66.0	63.0	64.5	70.0	72.0	64.5	63.0
Burst stress (200), hedge on							
Global γ vs. random, 5 names	54.0	54.5	58.0	58.5	60.0	64.0	54.0
Global γ vs. random, 13 names	56.0	59.0	59.0	61.0	64.5	64.0	48.5
Global γ vs. random, 34 names	53.0	60.0	66.0	69.0	61.0	54.0	57.5
γ per sector vs. random per sector	61.5	60.0	60.0	70.0	60.5	60.0	61.0
Bear stress (200), hedge off							
Global γ vs. random, 5 names	40.0	38.5	41.5	45.0	72.0	63.0	43.0
Global γ vs. random, 13 names	29.0	28.5	28.0	38.0	70.0	71.0	29.5
Global γ vs. random, 34 names	53.5	46.0	47.0	65.5	69.0	72.0	35.5
γ per sector vs. random per sector	54.0	44.0	50.5	67.0	68.5	74.5	42.5
Bear stress (200), hedge on							
Global γ vs. random, 5 names	43.0	41.5	44.0	46.5	63.0	60.0	37.5
Global γ vs. random, 13 names	33.5	34.0	34.5	45.0	66.5	61.0	30.5
Global γ vs. random, 34 names	48.5	50.0	49.0	61.0	59.5	63.0	40.5
γ per sector vs. random per sector	53.5	53.5	53.0	63.5	61.0	63.0	40.0

Table S14: **Market exposure on the stress sets of synthetic futures, hedge off.** For each selection rule (row) and weighting rule (column), the mean portfolio β (Table S8) as the median across the 200 futures of each stress set, after a median over seeds within a future (seeds 5660–5664 for the random and bandit rules, 5660 for the rest). The burst stress set holds the 200 futures with the deepest SPY drawdowns among those with a jump burst, and the bear stress set the 200 deepest among those without one. The first three rows of each block hold 33 names (3 per sector), the next five hold 13, and the last two are the hand-picked lists.

Selection rule	CD	CES	CESw	SG	MinVar	Tangent	LS
Burst stress, mean portfolio β							
γ per sector	0.739	0.747	0.749	0.751	0.542	0.637	0.690
Bandit per sector	0.732	0.733	0.739	0.755	0.532	0.675	0.677
Random per sector	0.659	0.667	0.663	0.728	0.557	0.695	0.552
Global γ	0.690	0.695	0.693	0.692	0.652	0.662	0.625
Correlation-aware γ	0.695	0.698	0.701	0.697	0.647	0.662	0.631
Liquid then γ	0.721	0.714	0.721	0.725	0.707	0.710	0.656
Random liquid 100	0.571	0.576	0.572	0.627	0.725	0.739	0.395
Random whole universe	0.528	0.526	0.525	0.582	0.705	0.734	0.322
Hand-picked 13	0.565	0.565	0.567	0.631	0.764	0.761	0.384
Hand-picked 13 + WMT, PG	0.552	0.555	0.555	0.633	0.760	0.761	0.369
Bear stress, mean portfolio β							
γ per sector	0.768	0.788	0.787	0.791	0.510	0.629	0.709
Bandit per sector	0.779	0.788	0.788	0.804	0.521	0.699	0.710
Random per sector	0.649	0.651	0.649	0.765	0.571	0.729	0.483
Global γ	0.715	0.717	0.722	0.717	0.643	0.669	0.631
Correlation-aware γ	0.724	0.725	0.729	0.726	0.647	0.670	0.642
Liquid then γ	0.776	0.768	0.778	0.782	0.739	0.750	0.693
Random liquid 100	0.512	0.514	0.515	0.586	0.798	0.819	0.293
Random whole universe	0.460	0.458	0.460	0.536	0.769	0.798	0.229
Hand-picked 13	0.489	0.486	0.489	0.578	0.850	0.850	0.284
Hand-picked 13 + WMT, PG	0.484	0.487	0.485	0.584	0.847	0.850	0.250

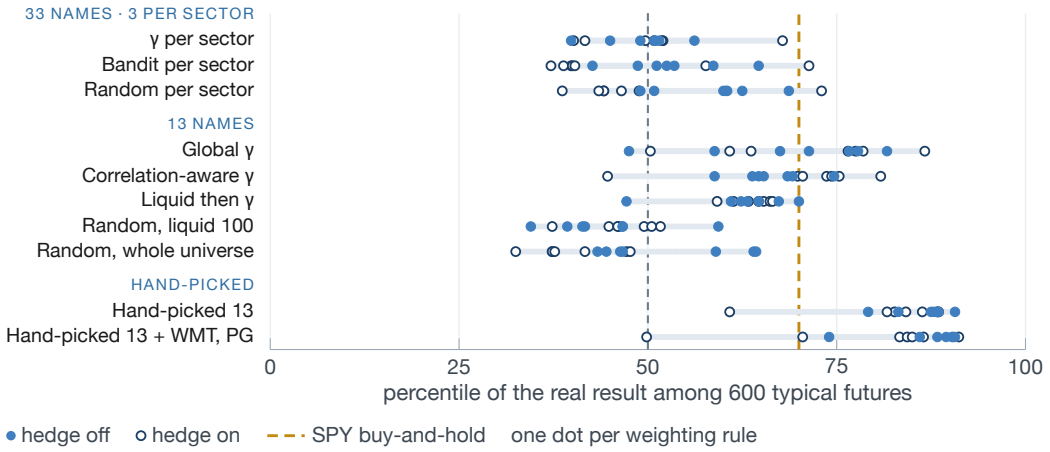


Figure S1: **Real-data results as percentiles of the typical synthetic futures.** For each selection rule and weighting rule, the percentile of the real terminal wealth ratio among the same cell's 600 typical futures Filled dots have the market hedge off and open dots have it on (a β cap, set at each re-plan, that was lower when the market was weaker, plus up to five names with $\beta \leq 0$); the pale bar spans each row's lowest to highest dot. At 50 (dashed) a real result equals the median future; the gold line marks real SPY, at the 70th percentile of synthetic SPY. Real results use the seeds of the synthetic runs: the median over seeds 5660–5664 for the random and bandit rules, and seed 5660 for the rest.

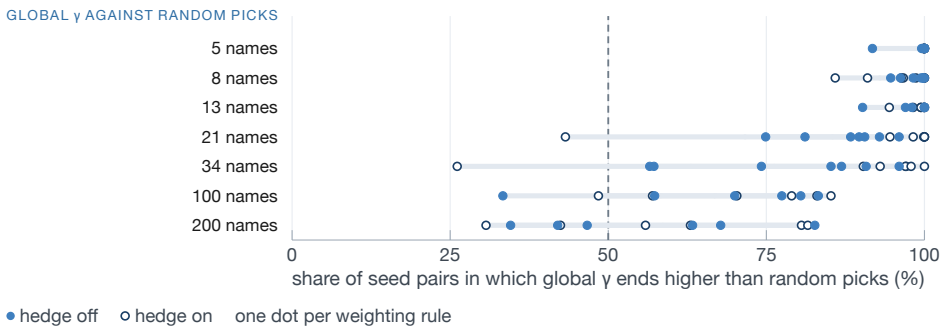


Figure S2: **Global γ against random picks from the whole universe, from 5 to 200 names, on real prices.** For each basket size and weighting rule, the share of the 900 pairs of one global γ run and one random run (30 runs each) in which global γ ended with the higher terminal wealth ratio (market hedge off filled, on open). At 50% (dashed) the two rules end higher equally often. Stone–Geary is left out at 100 and 200 names, where its baseline holding of 1% of wealth per name uses the whole budget or twice the budget.